

DHBW Informatik Semester 4
Wahlpflicht
Cloud Computing
Chapter 3 PaaS towards Serverless
Or what comes next on top of Kubernetes

Juergen Schneider



Agenda



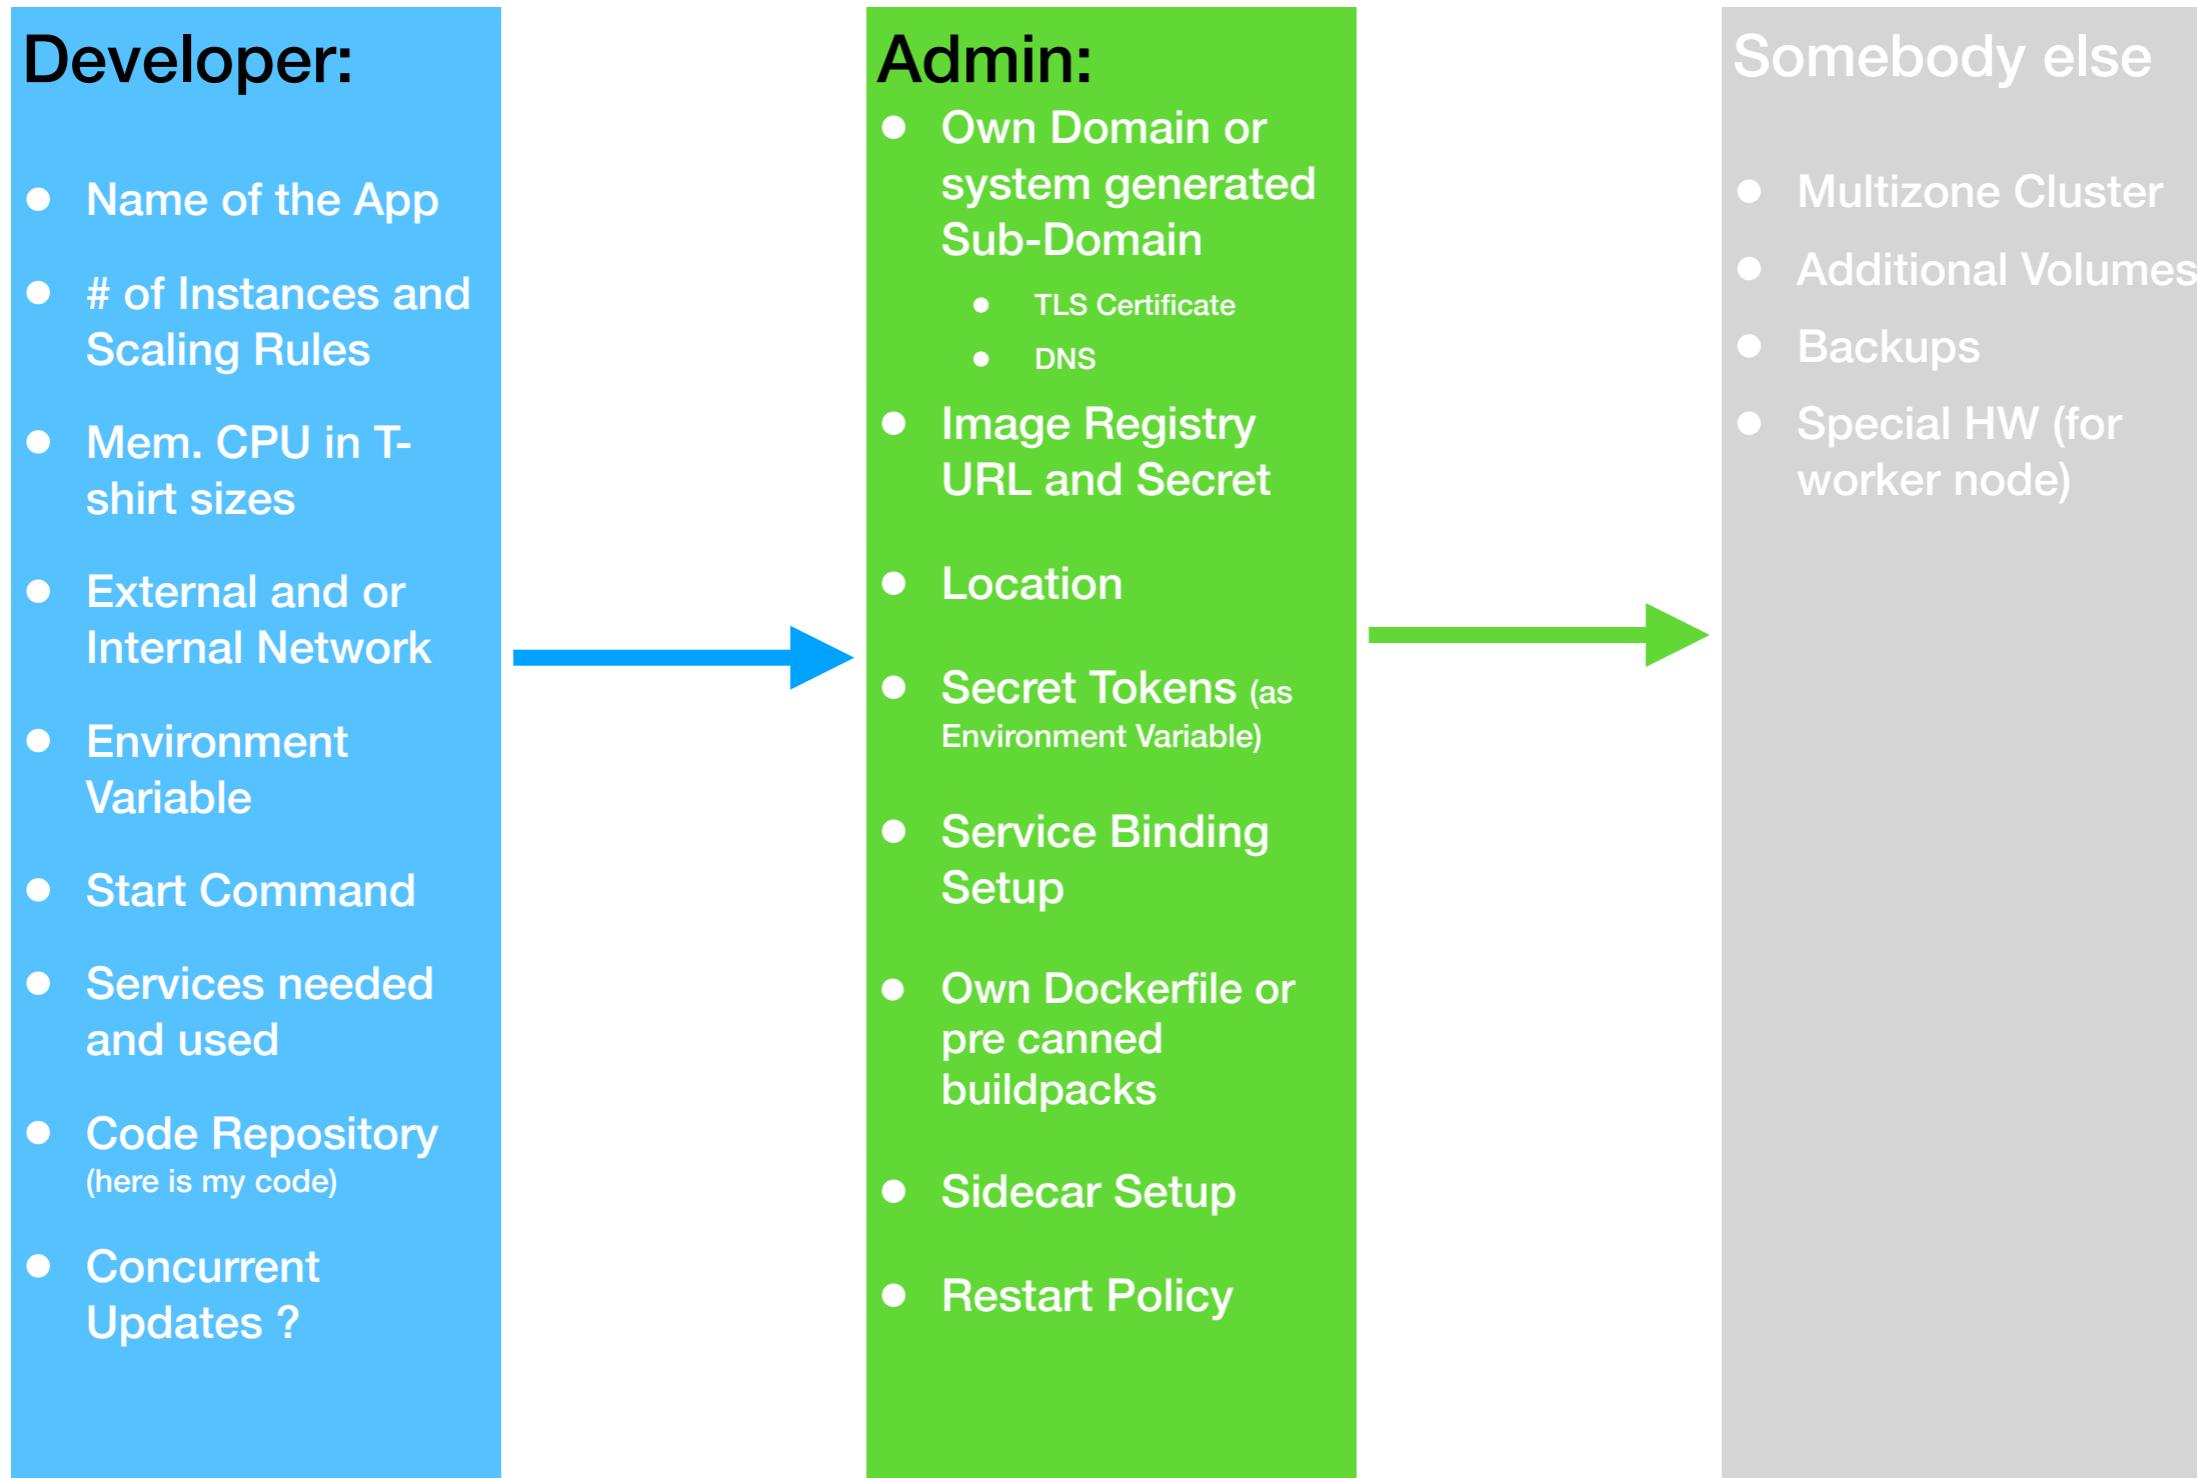
- What is next on top of Kubernetes
- The Principles of Cloud Foundry and Serverless.

Why should I care about containers ?

- All I need is a means to give somebody my code and have it run in the cloud.... ([a Developer view](#))
- Containers ([PaaS but with the need of infrastructure skill](#))
 - More control of the deployment, the resource usage (and limits) the network traffic and security
 - But complicated, needs skill in infrastructure
- PaaS ([the ‘real’ PaaS](#))
 - ‘Explain the Code’ and a **bit of infrastructure** and **accept** the rest to be done by the cloud (platform layer) and admin
 - Service Access is HTTPS or another IP Layer 4 Protocol with a generated subdomain.. (TLS Certificates available or configurable)
 - No control of what is virtualized or not
 - Data (only) via Data (Database) Services (typically no local disk recovery)
 - Programming model driven by the 12 factor directives

What does a Developer and his/her admin really need

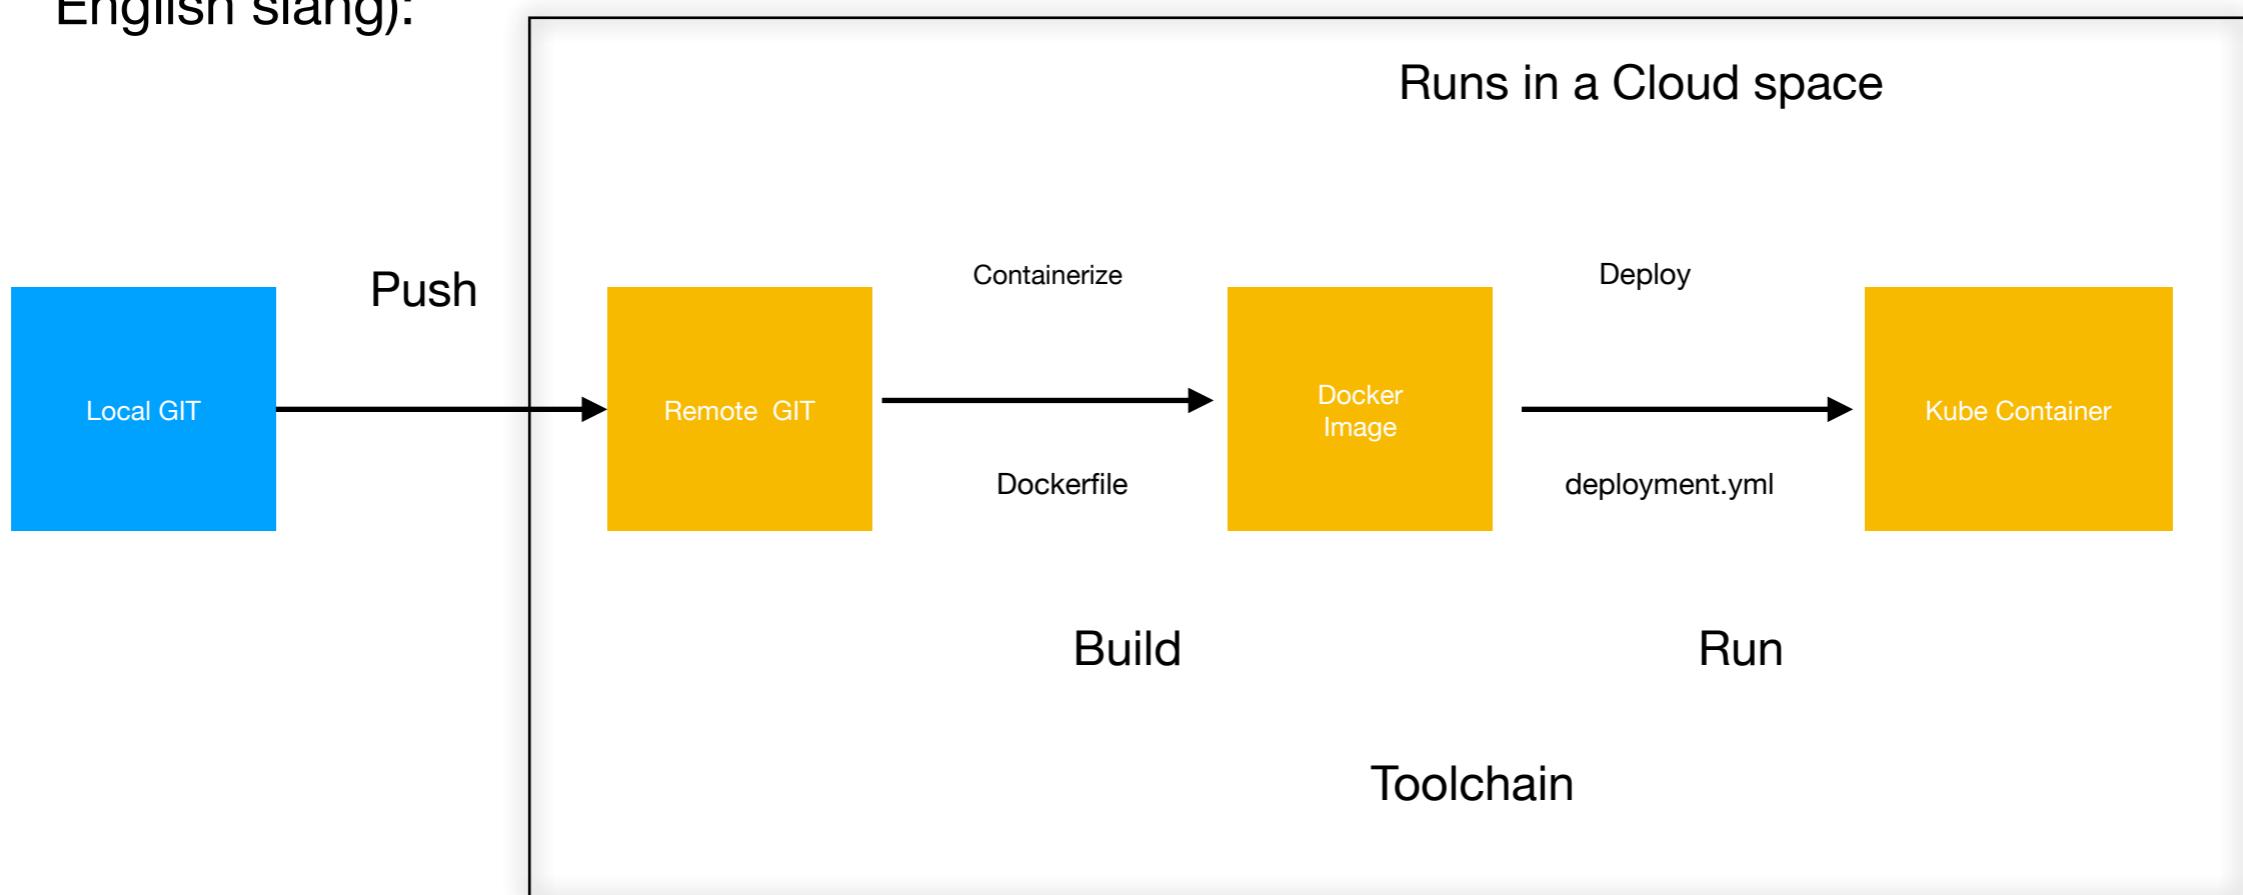
Stateless cloud native App



Enhance Kubernetes to allow a developer Usage (Concept Example)

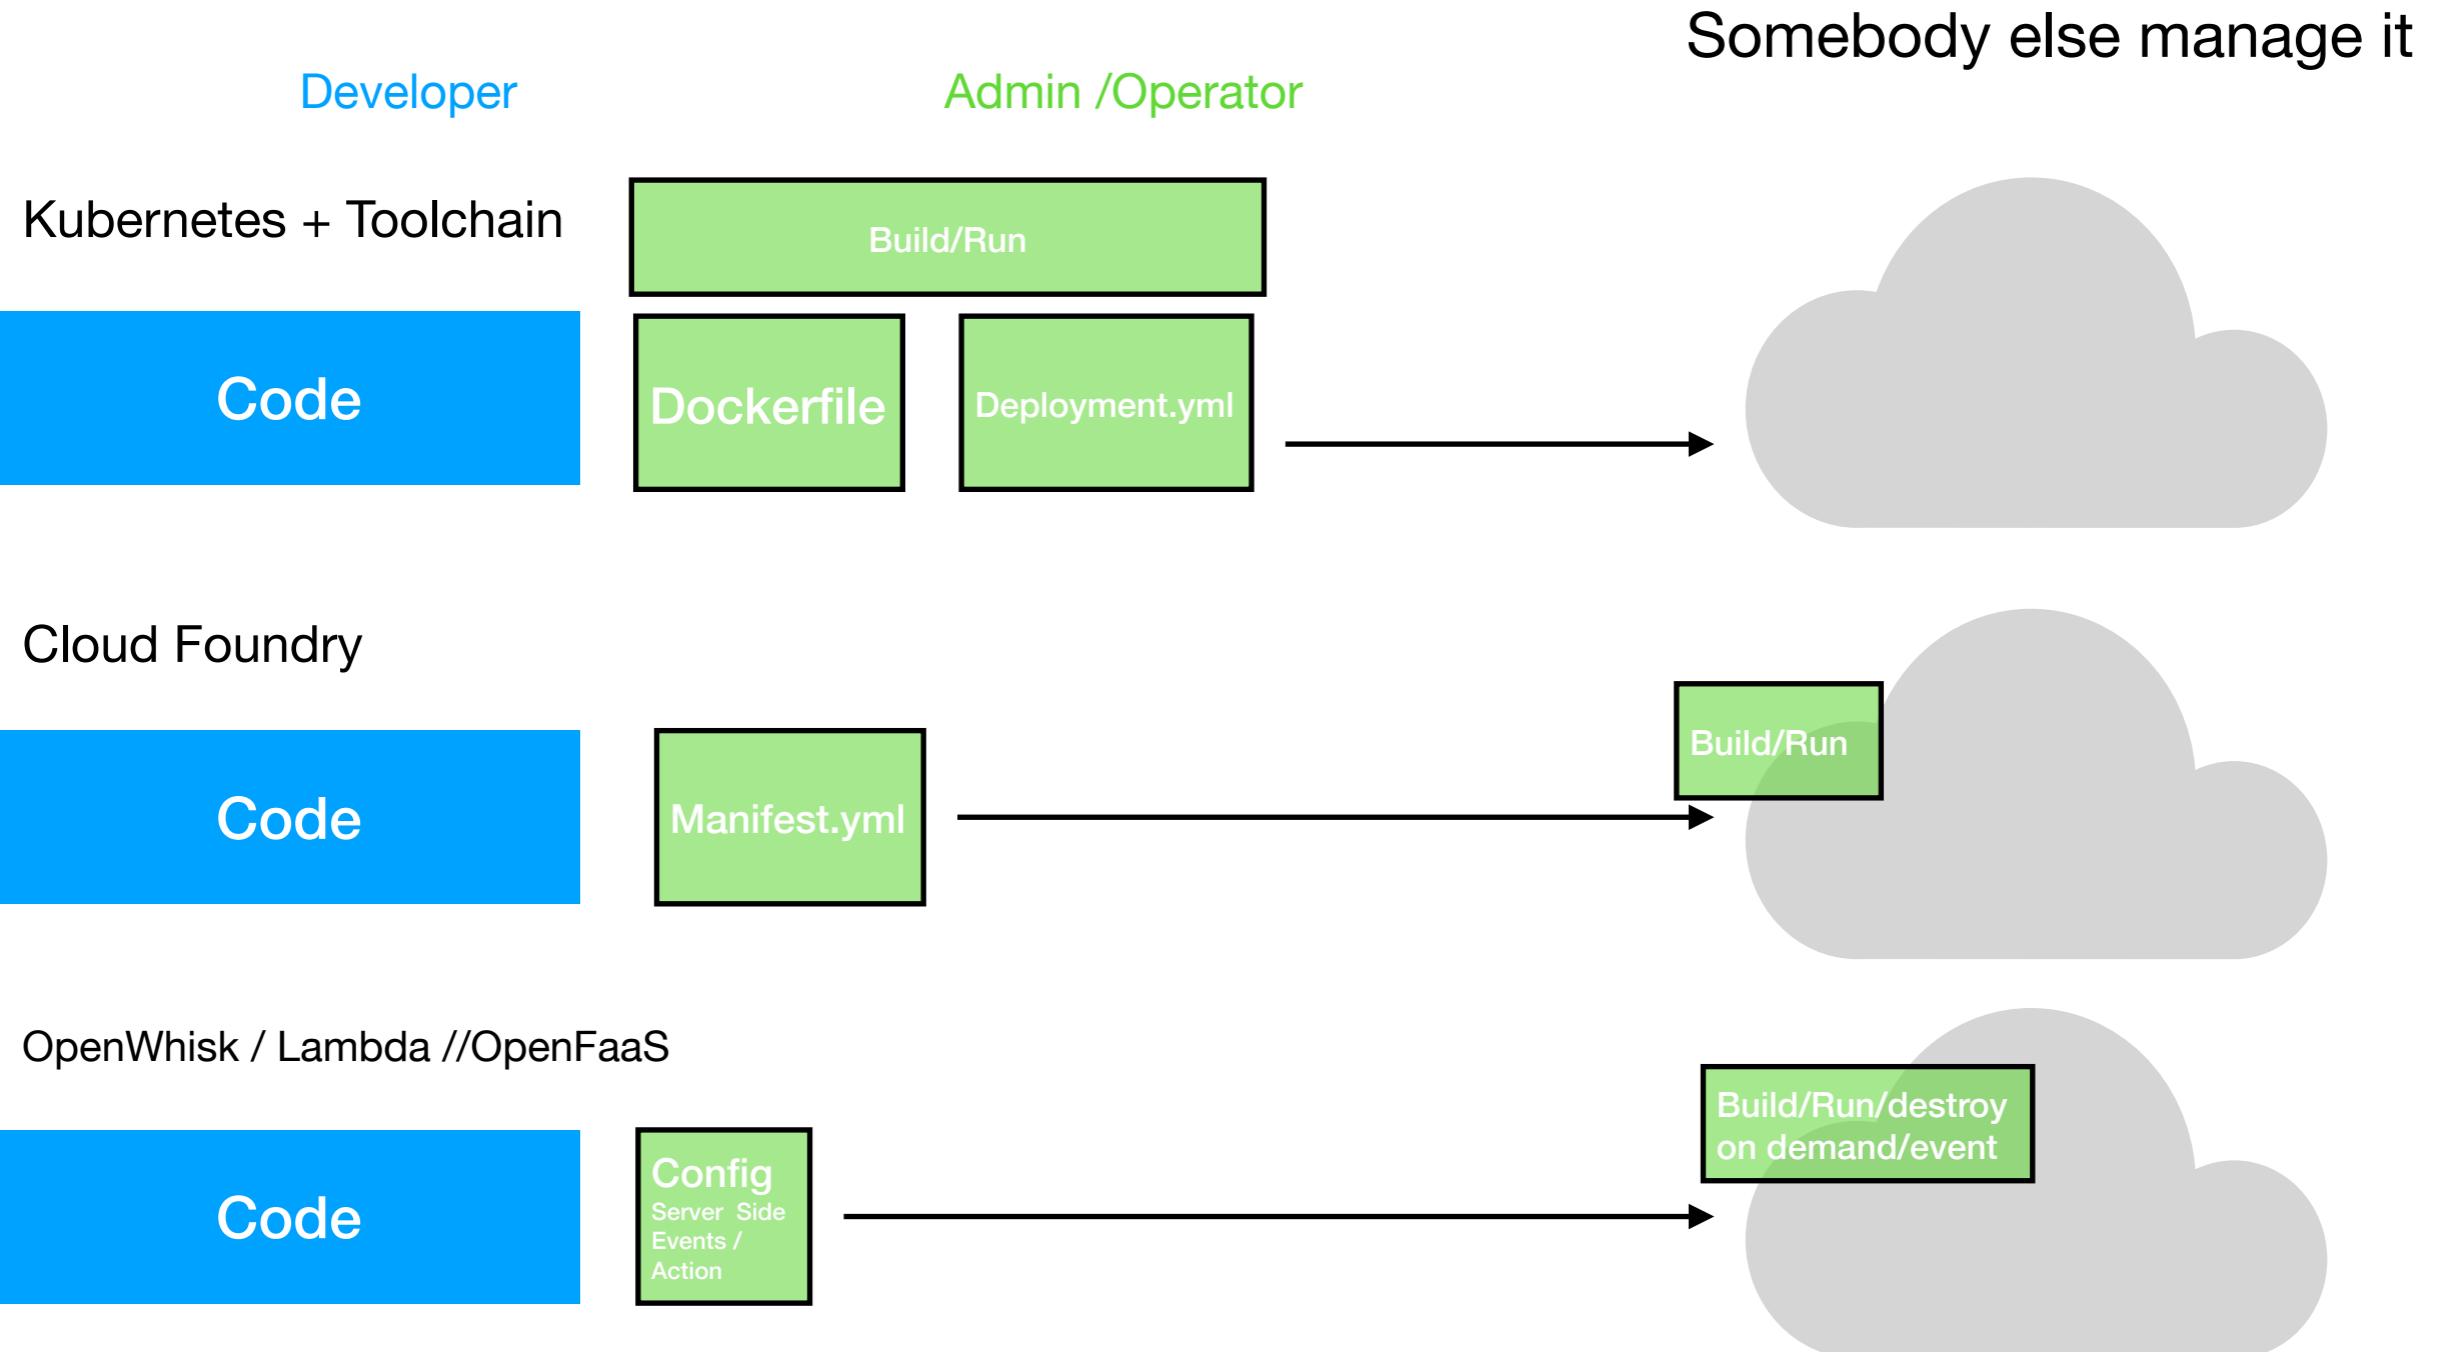
Add development tool chain integration on top of Kubernetes

- Let a admin person deal with all the YAMLs,
 - the proper network setup, the container, the images, the side-car etc.
- Let the developer just use a Code Repository to push code into the cloud.
 - The cloud configuration is in files of the **remote** repository but not in my code (12 factor app)
 - Torvalds sarcastically joked about the name *git* (which means "unpleasant person" in British English slang):

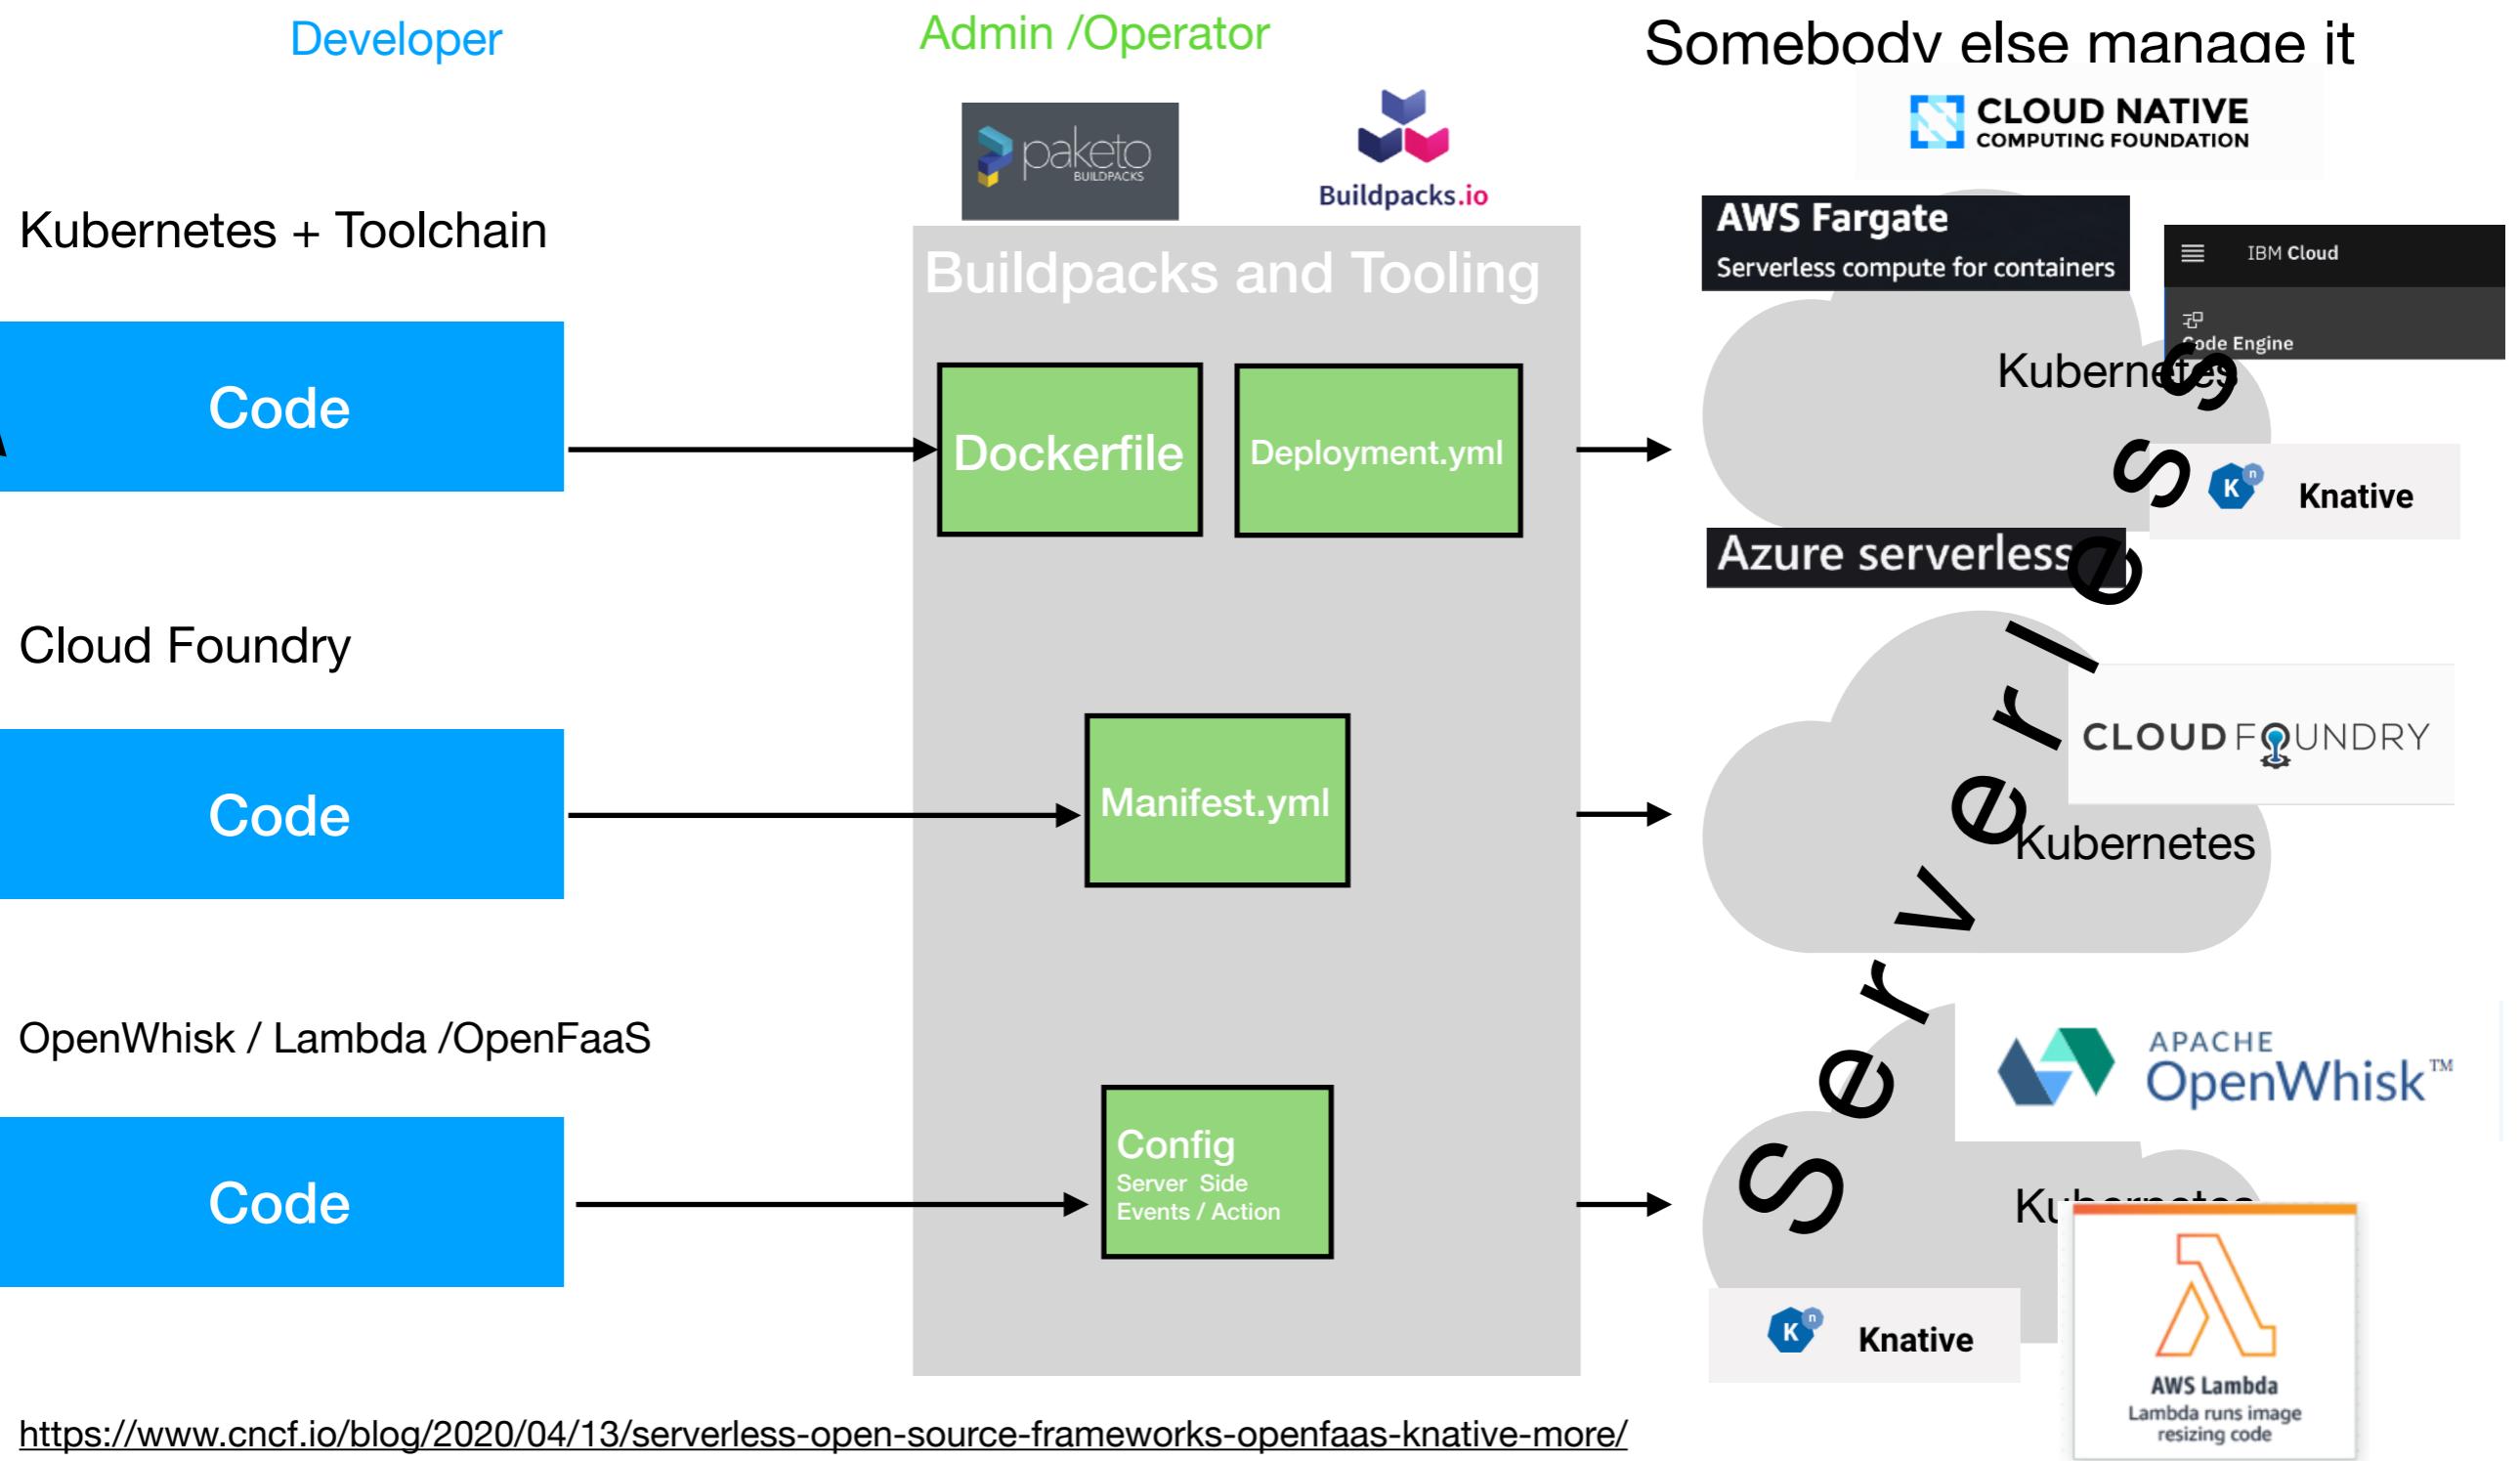


PaaS PaaS++(FaaS)

Dev / Ops (User) provides it



The current Market Place Battle (biased)

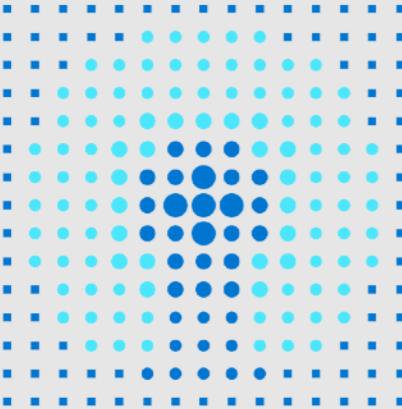


Serverless ? - Cloud Provider View

<https://azure.microsoft.com/en-us/resources/cloud-computing-dictionary/what-is-serverless-computing/>

Serverless application patterns

Developers build serverless applications using a variety of application patterns—many of which align with approaches that are already familiar—to meet specific requirements and business needs.



Serverless functions

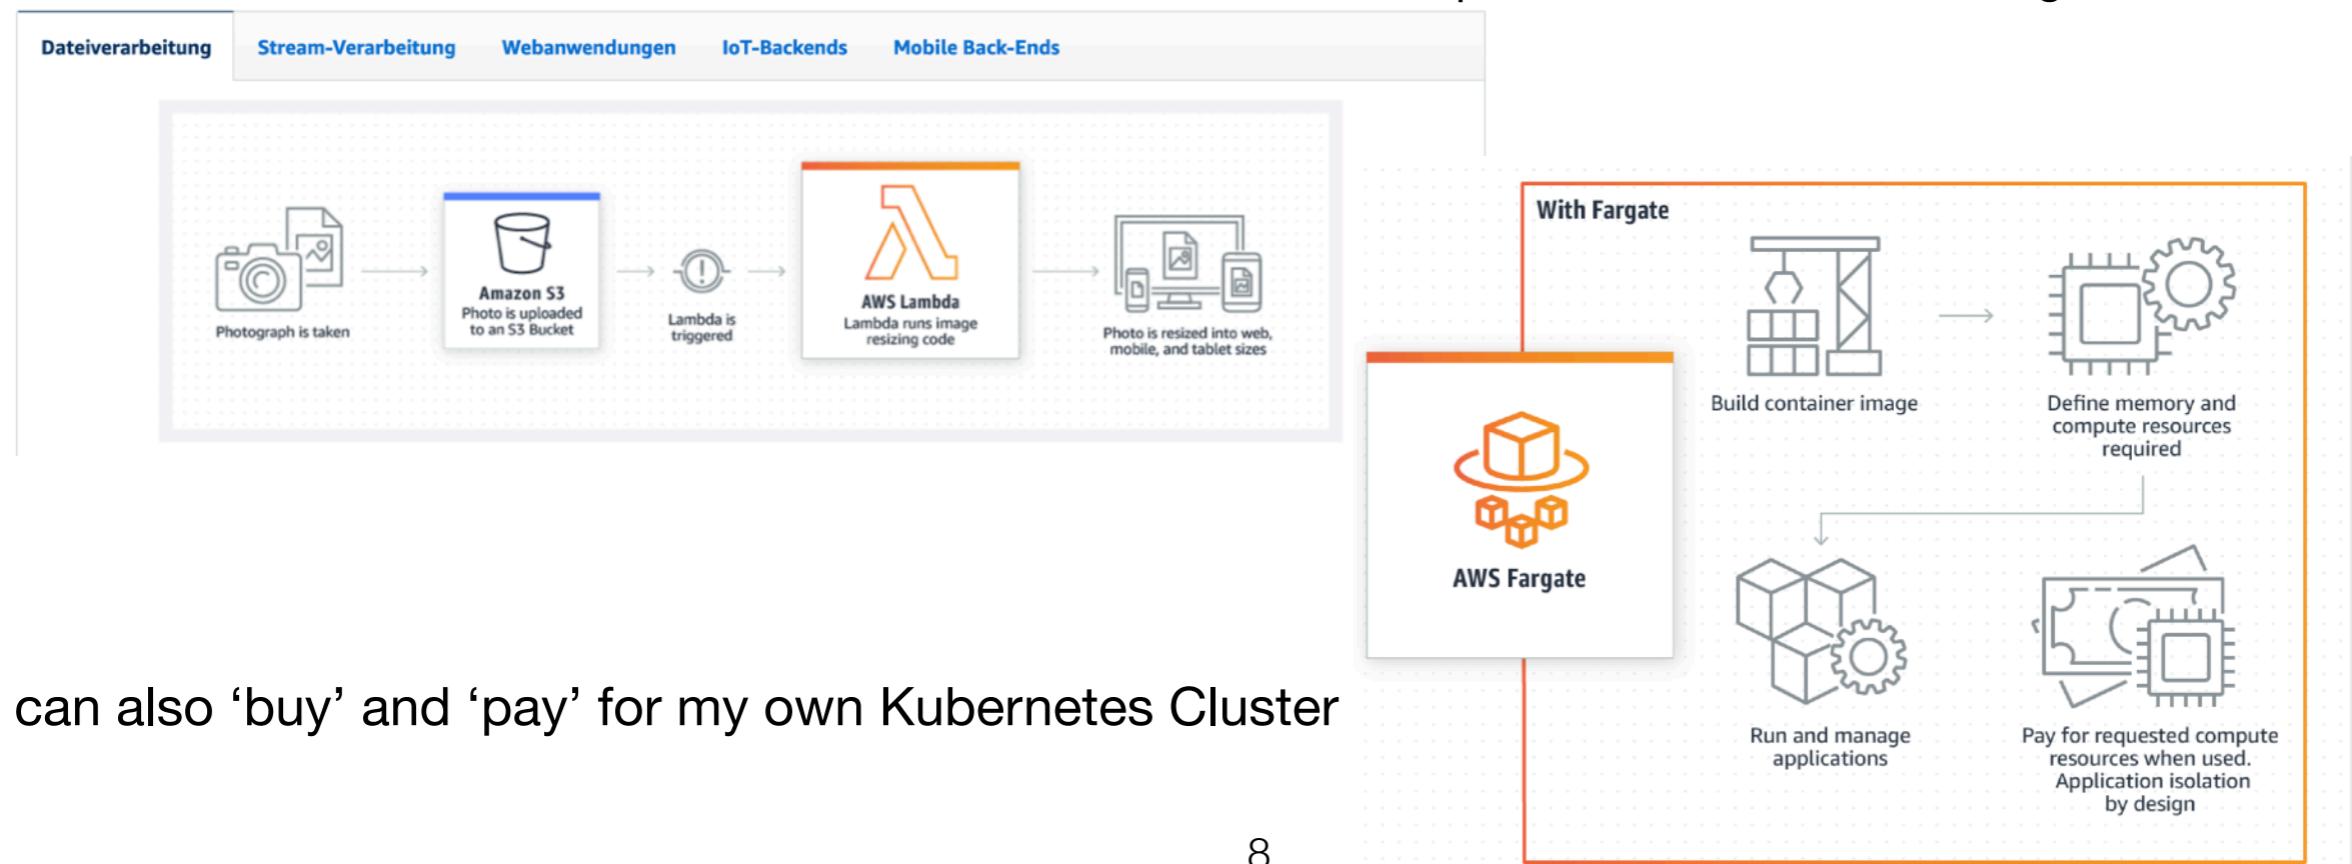
Serverless functions accelerate development by using an event-driven model, with triggers that automatically execute code to respond to events and bindings to seamlessly integrate additional services. A pay-per-execution model with sub-second billing charges only for the time and resources it takes to execute the code.

Serverless Kubernetes

Developers bring their own containers to fully managed, Kubernetes-orchestrated clusters that can automatically scale up and down with sudden changes in traffic on spiky workloads.

<https://aws.amazon.com/de/lambda/>

<https://aws.amazon.com/de/fargate/?c=ser&sec=srv>



Event Driven Architecture

Why not ask ChatGPT, here is the result:

Events as First-Class Citizens

- An **event** is a significant change in system state (e.g., "user placed an order").
- Events are published and consumed as the primary mechanism for communication.

Loose Coupling

- Producers (senders) of events do not know who will consume them.
- Consumers (listeners/subscribers) process events independently.
- This promotes flexibility and scalability.

Asynchronous Communication

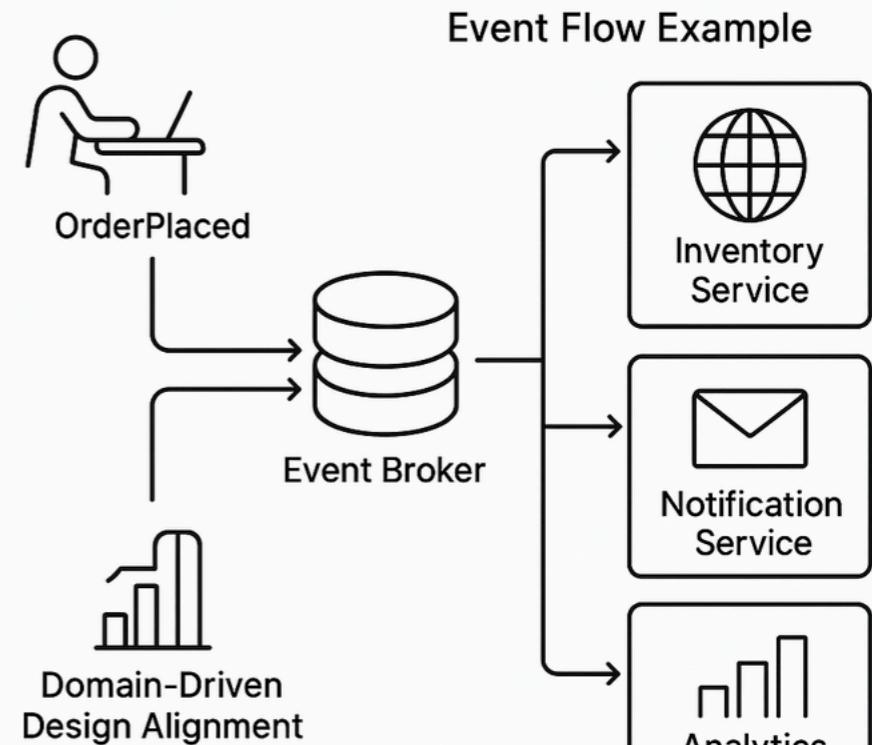
- Events are typically processed asynchronously.
- Producers do not wait for consumers to finish processing.
- Improves system responsiveness and decouples processing speed.

Event Brokers or Event Buses

- Middleware like Kafka, RabbitMQ, HTTP, or AWS EventBridge routes events from producers to consumers.
- Helps manage scalability, filtering, and delivery guarantees.

Principles of Event-Driven Architecture

-  Events as First-Class Citizens
-  Loose Coupling
-  Asynchronous Communication
-  Event Brokers or Event Buses
-  Event Consumers and Replayability
-  Scalability and Resilience
-  Domain-Driven Design Alignment



Event Persistence and Replayability (optional)

- Events can be stored for auditing, debugging, or replaying system state.
- Event Sourcing is a related concept where application state is derived from a sequence of events.

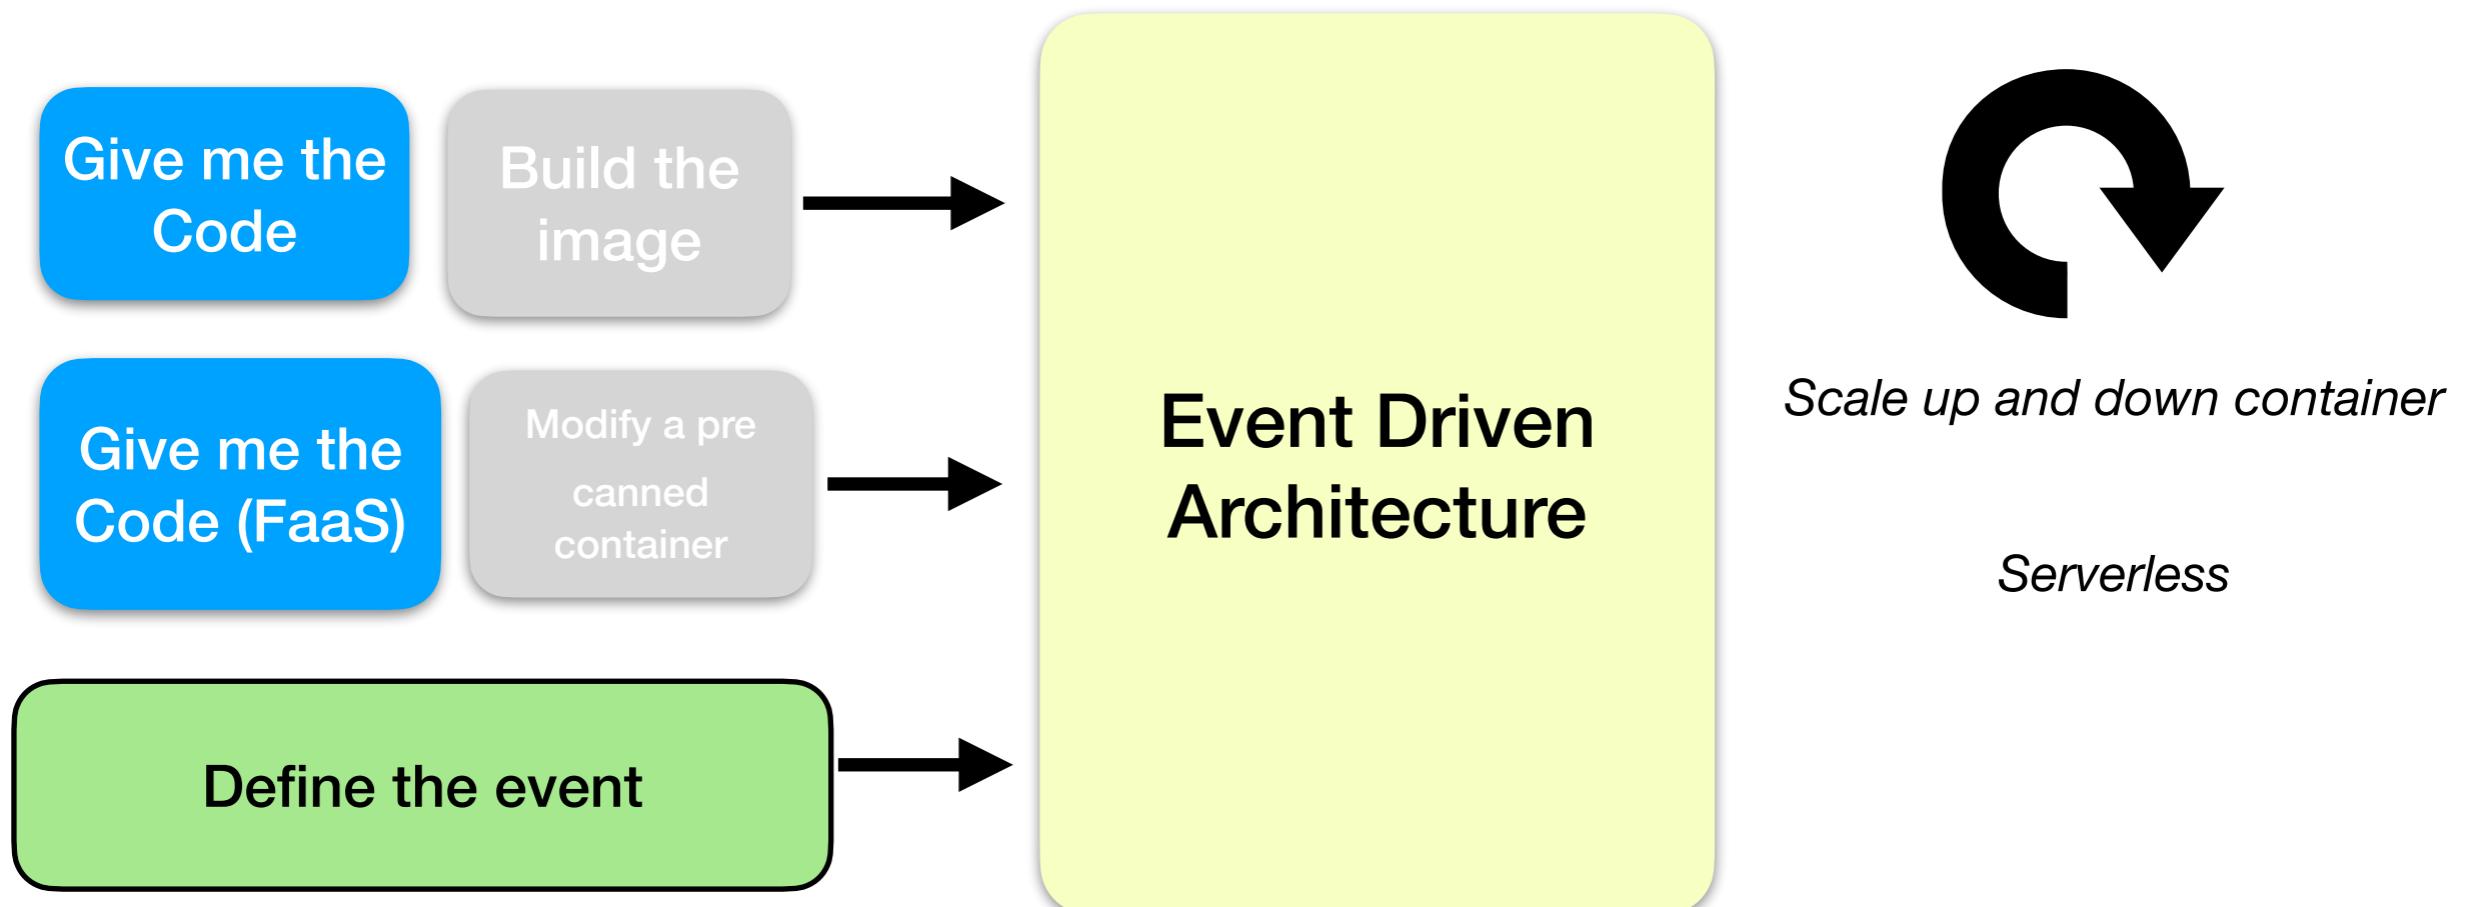
Scalability and Resilience

- **The architecture naturally supports horizontal scaling.**
- Components can fail independently, improving fault tolerance.

Domain-Driven Design Alignment

- Events often map closely to business domain events, aligning the system behaviour with business operations.

Event Driven Architecture → Serverless



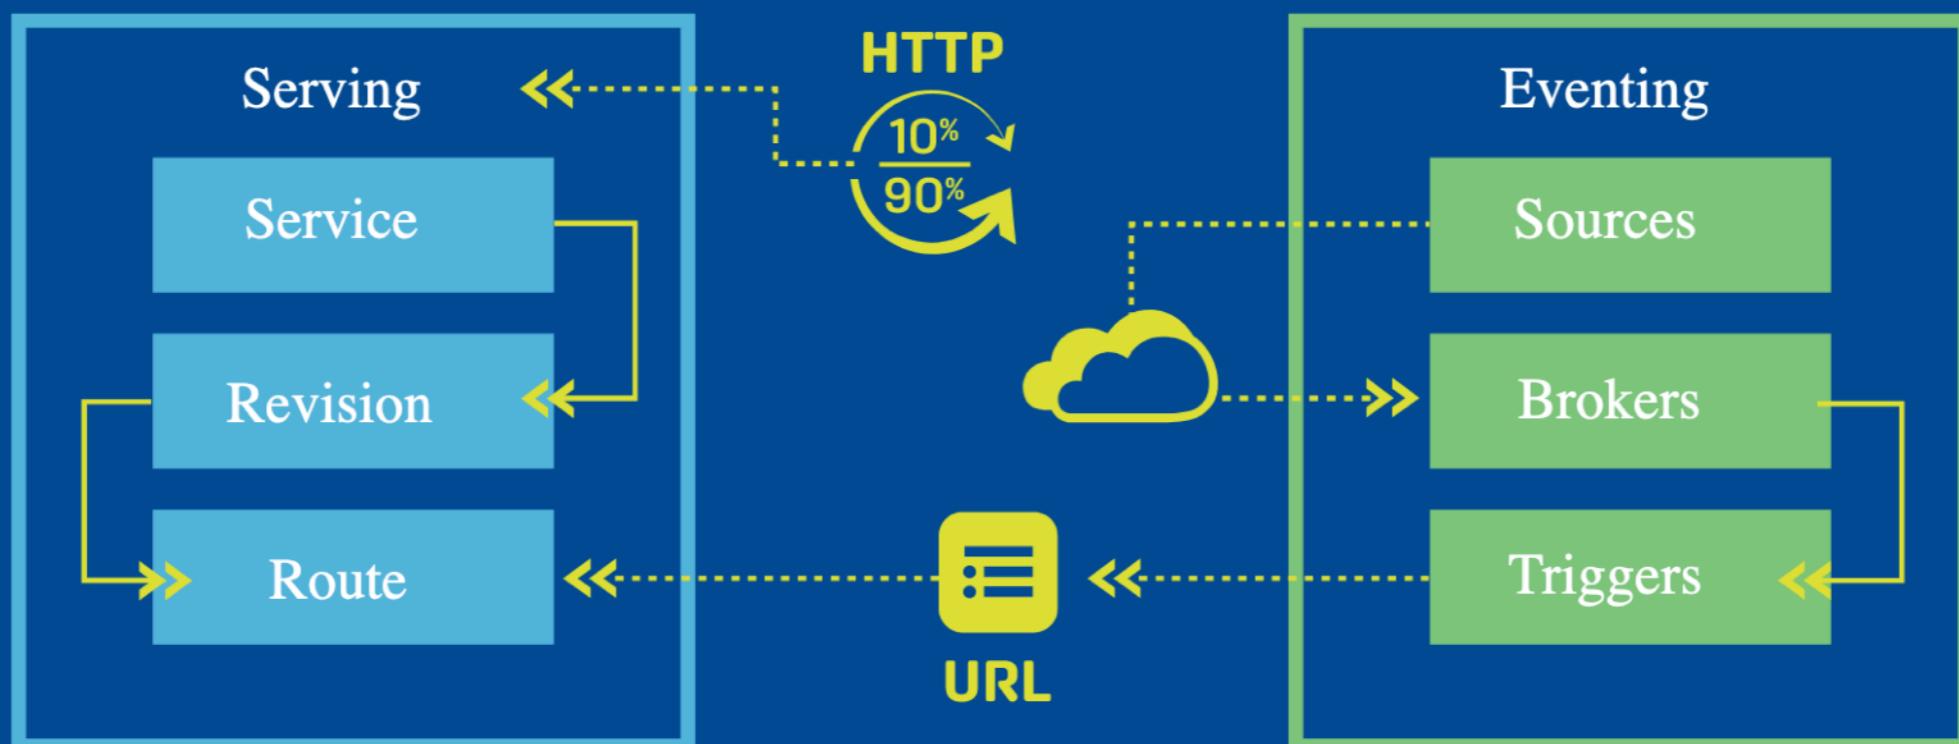
Open Source View

<https://www.cncf.io/>

<https://knative.dev/docs/>

Knative Components

Knative has two main components that empower teams working with Kubernetes. Serving and Eventing work together to automate and manage tasks and applications.



Run serverless containers in Kubernetes with ease. Knative takes care of the details of networking, autoscaling (even to zero), and revision tracking. Teams can focus on core logic using any programming language.

Universal subscription, delivery and management of events. Build modern apps by attaching compute to a data stream with declarative event connectivity and developer friendly object models.

IBM Code Engine - a Knative Implementation

Workload : see next charts

- Applications, Jobs, Functions

Event Sources

- Cron, IBM Object Storage (COS), Kafka, Webhooks (Github) .. more to come

Event Brokering

- related to the Event Source
 - Kafka -> Topics,
 - COS -> Updates for Bucket entries (create, read, update, delete)
 - HTTP (Url)
- Trigger event data follows the cloud events specification: <https://cloudevents.io/>

Revisions:

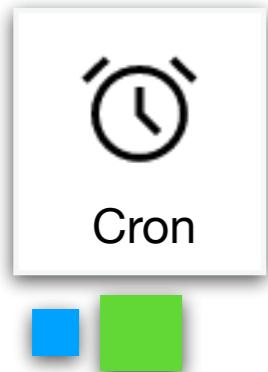
- Revisions represents snapshot-in-time of application code and configuration. Revisions enable progressive rollout and rollback of application changes by changing the HTTP routing between the Revision instances.

Applications - Jobs - Functions

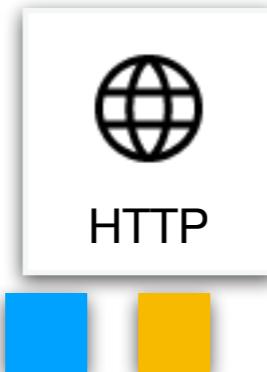
Triggers



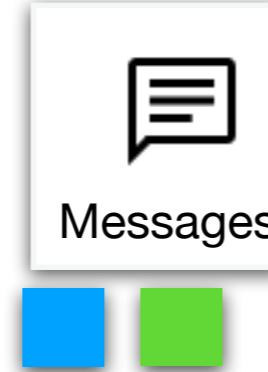
Integration with Git Repository
For Image Build



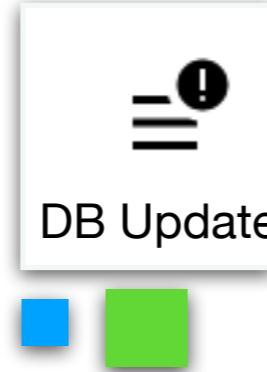
Cron



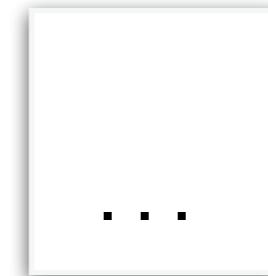
HTTP



Messages



DB Update



Application

Long running
Port listening
Multiple request/threads

Scaling based #
request (min/max)

Jobs

Very long running
No Port listening

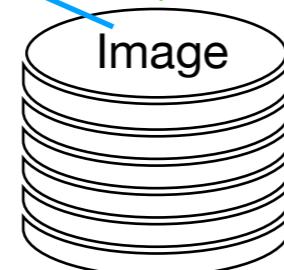
Scaling based
specific workload

Functions

short running
No Port listening
Single request/thread

Scaling based on
request (0 to current)

Source



Bring your image

- build by Dockerfile or buildpacks

```
main.js
1 function main(args) {
2   const body = {
3     args,
4     env: process.env,
5   };
6   console.log(`Return body: ${JSON.stringify(body, null, 2)}`);
7   return {
8     statusCode: 200,
9     headers: {
10       'Content-Type': 'application/json',
11     },
12     body,
13   };
14 }
15 module.exports.main = main;
```

Bring your code - inline or Github

Capacity and Scaling Profiles for Application

The screenshot shows the AWS Lambda function configuration interface. The left sidebar has tabs: 'Code', 'Resources & scaling' (which is selected and highlighted in blue), 'Environment variables', and 'Image start options'. The main content area is divided into sections:

- Instance resources**: Define the amount of CPU and memory resources for each instance. It shows 1 vCPU / 4 GB for CPU and memory, and 0.4 GB for ephemeral storage.
- Autoscaling - instance scaling range**: Specify the range within which Code Engine should be auto-scaling the number of running instances. It shows a minimum of 0 and a maximum of 10 instances. A note says "For unlimited scaling, set to 0".
- Autoscaling - request concurrency and timing settings**: Maximum number of concurrent requests per instance and timeout for requests. It shows target concurrency at 59, max concurrency at 100, request timeout at 300 seconds, and scale-down delay at 0 seconds.

- **Scaling Range** : min/max number of instances
- **Target concurrency** : average concurrent requests to start scale up to avoid max concurrency
- **Max concurrency** : Hard limit to start to scale up
- **Request timeout** : a time in which the app should respond, before the requestors may fail. As consequence the request will be terminated (HTTP Status 503). And it may also be an indication that something is wrong
- **Scale down delay**: Number of seconds to wait (delay) before scale down (as the number of concurrent requests proposes a scale down)

Capacity and Scaling Profiles for Jobs

The screenshot shows the 'Resources & scaling' section of the AWS Lambda configuration interface. It includes:

- Job instances:** Set to "Array size".
 - Array size:** Value 1.
 - Array indexes:** Value 0.
- Instance resources:** Define CPU and memory resources.
 - CPU and memory:** Value 1 vCPU / 4 GB.
 - Ephemeral storage (GB):** Value 0.4.
- Resiliency settings:** Specify retry and timeout policies.
 - Mode:** Set to "Task".
 - Number of job retries:** Value 3.
 - Job timeout (seconds):** Value 7200.

- **Job Instances :** you can split the Batch/Job workload in more than one instances, if that workload allows it and each instance could have a unique index as input (Job Arrays). Each instance will run once.
- **Task or Daemon :**
 - In Daemon mode: the job runs without a maximum execution time, and failed indexes are restarted indefinitely.
 - In Task mode: **Number of retries** a single Job (Task) are restarted before considered as failed and **Job Timeout** reflects the time a single Job can run (maximum execution time)

Capacity and Scaling Profiles for Functions

Configuration Domain mappings Service access

The screenshot shows the 'Configuration' tab selected in the top navigation bar. On the left, there's a sidebar with 'Code', 'Resources & scaling' (which is highlighted with a blue border), and 'Environment variables'. The main area is titled 'Resources & scaling' and contains instructions to 'Specify instance resources and scaling behavior.' It includes three input fields: 'CPU and memory' set to '0.5 vCPU / 2 GB', 'Timeout (seconds)' set to '60', and 'Scale-down delay (seconds)' set to '0'.

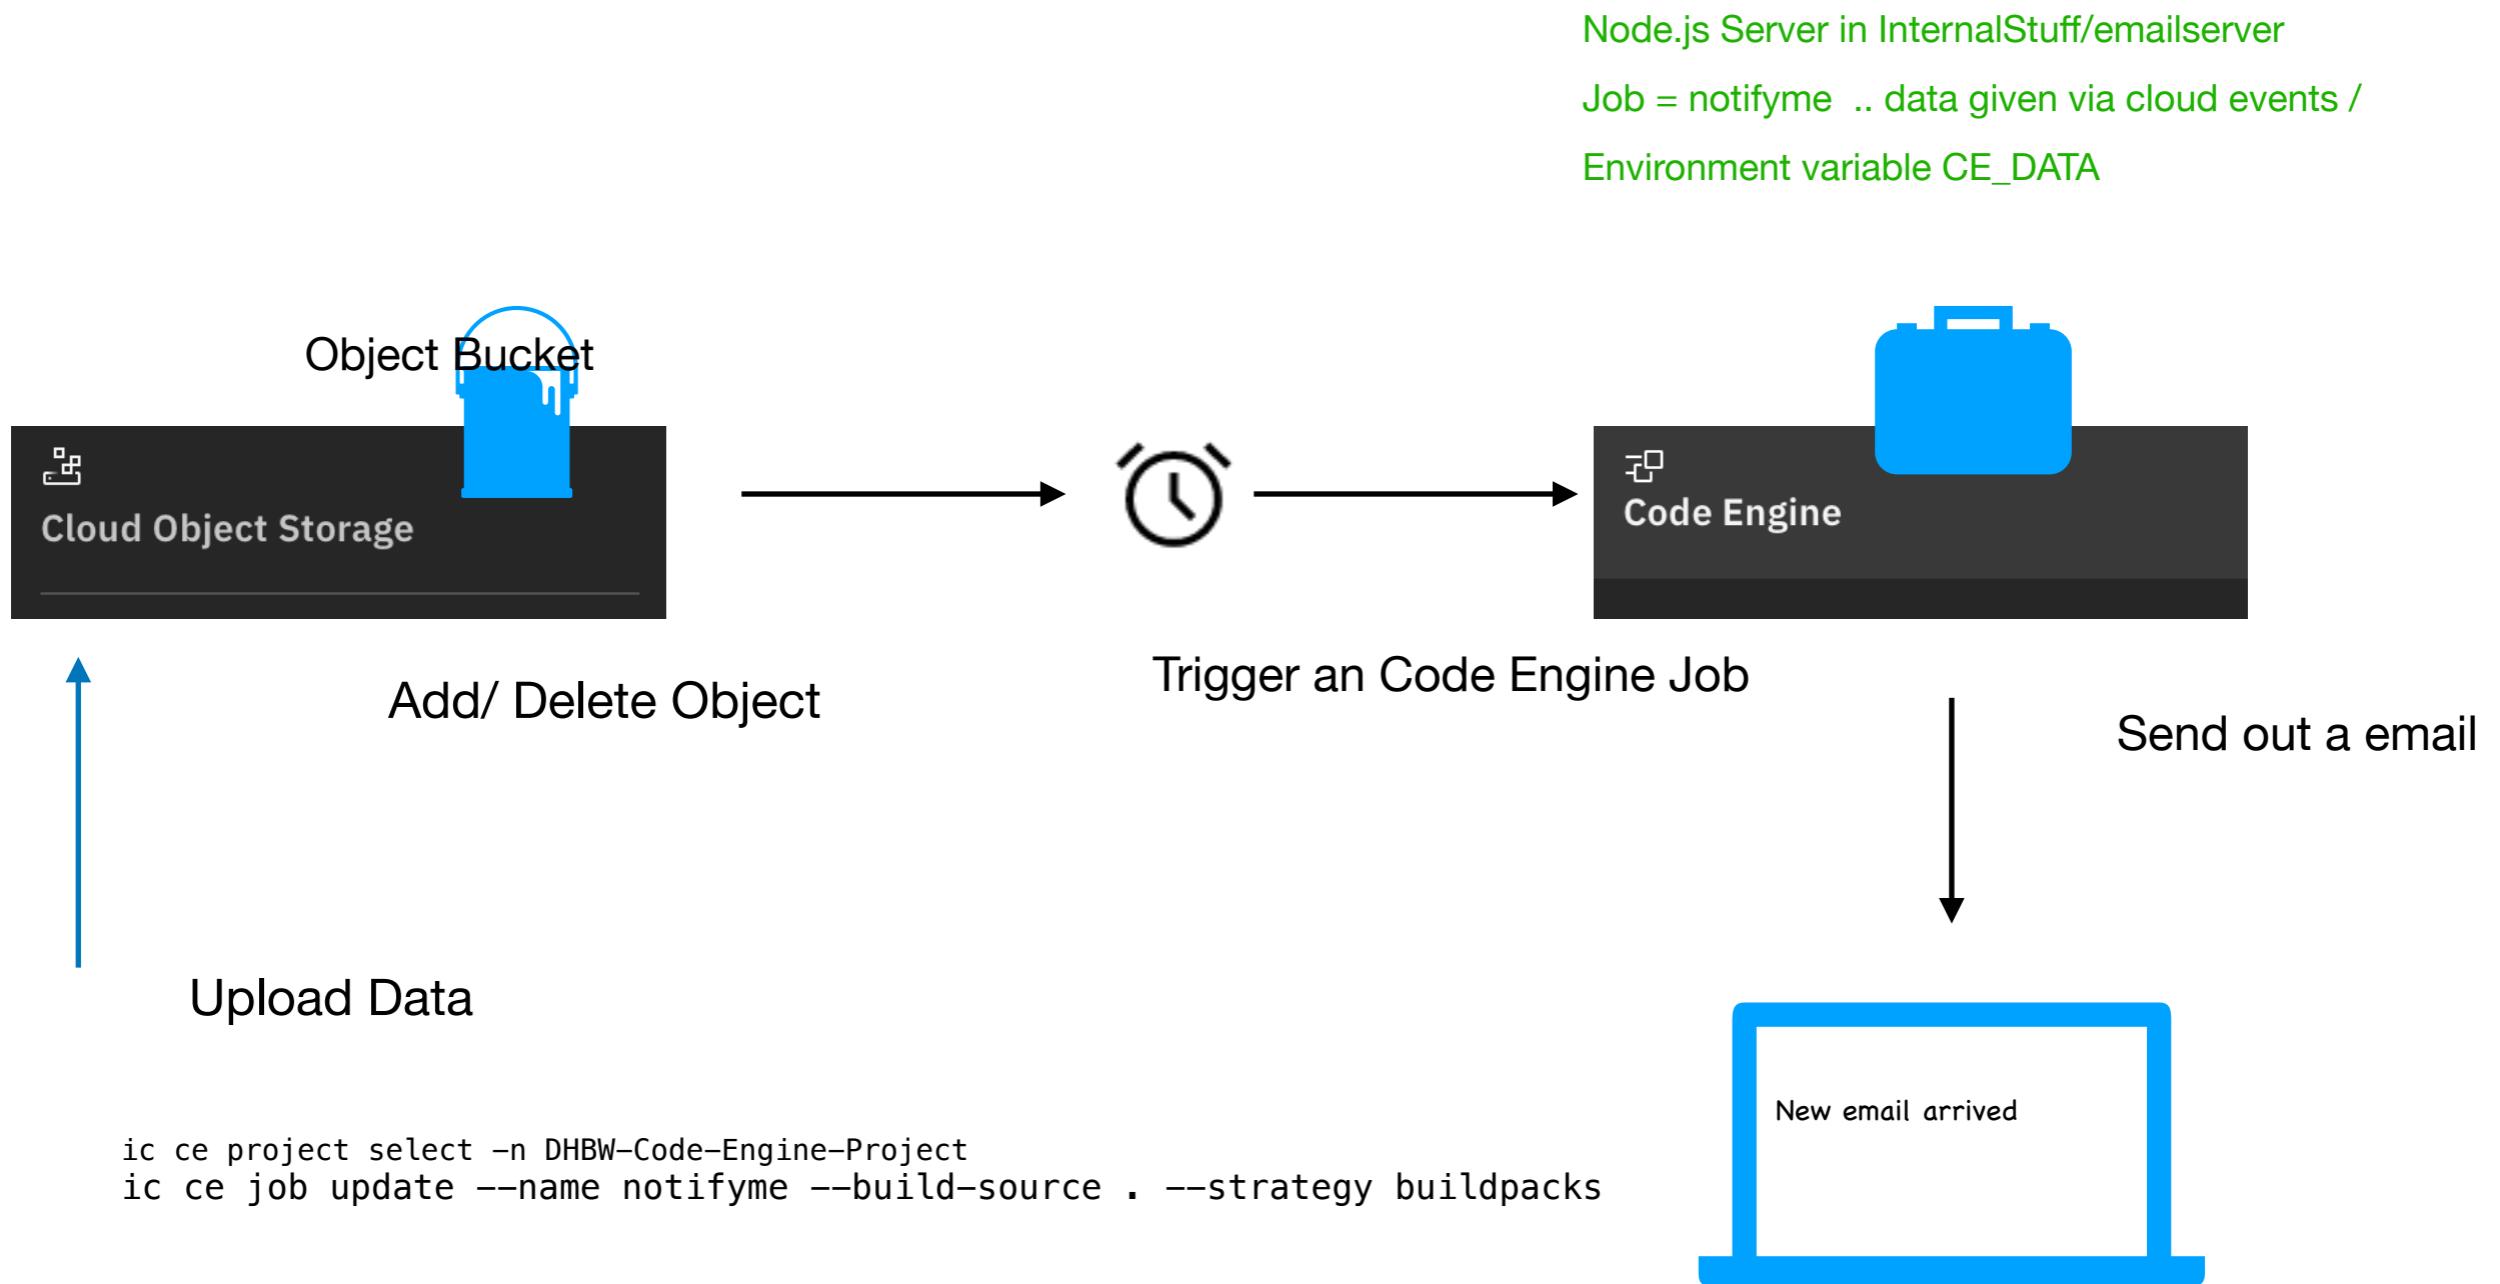
- For each incoming request a new container with the function code will run
- **Timeout** : how long is a single function allowed to run.
- **Scale down delay**: how long to wait until I give away an exited container, as maybe new requests come in soon

Demos :

- 1.Run a Job which sends out an email, of there is an update on COS DB
- 2.Trigger a Function which inverse the text given
- 3.Trigger a Function the speak the text given (Cloud Service)
- 4.Run our (almost) famous tinyApp as Serverless App
- 5.Demo a GIT (Webhook) to build an new image for our App

Example : Event Driven Invocation of Jobs in Code Engine

Demo : an Update of an COS bucket generates a email



Demo commands to be used in InternalStuff/CodeEngine/emailserver commands.txt

A basic Code Engine Function Example based on OpenFaaS

<https://docs.openfaas.com/languages/node/>

Demo : Function reverse is triggered by a HTTP Get and takes the ?text=data query string and reverse the characters. The result is shown back to the caller (browser)

<https://reverse.yk9dj48yz2s.eu-de.codeengine.appdomain.cloud?text=Geht es nun wirklich besser>



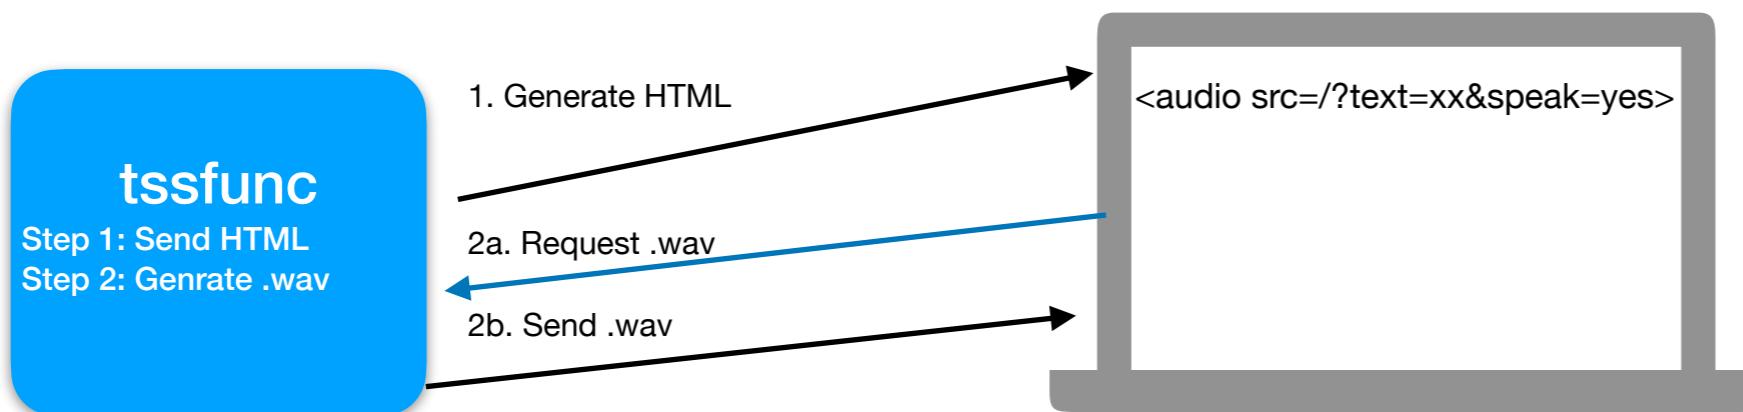
The screenshot shows a code editor window titled "Node.js 18". Inside the editor, there is a file named "main.js" with the following content:

```
1 function main(args) {
2   const body = {
3     result: (args.text == undefined) ? "Bitte mit dem richtigen Parameter" : args.text.split("").reverse().join(),
4   };
5
6   console.log(`Return body: ${JSON.stringify(body, null, 2)})`;
7
8   return {
9     statusCode: 200,
10    headers: {
11      'Content-Type': 'application/json',
12    },
13    body,
14  };
15}
16
17 module.exports.main = main;
18
```

A more sophisticated Code Engine Function based on OpenFaaS

Demo :

- Function sends a HTML document including an audio tag with an URL which again call this function to synthesise the audio using a IBM Cloud Service.
- As this function is asynchronous in nature the function is a asynchronous function
- As this function is using additional node.js packages, the function must be rebuilt generating a small image (with additional modules)
 - `ic ce fn create --name ttsfunc --runtime nodejs-20 --build-source .`
- The credentials to a Cloud Services are securely provided as Environment key/values (Kubernetes Secret File) and never externalised

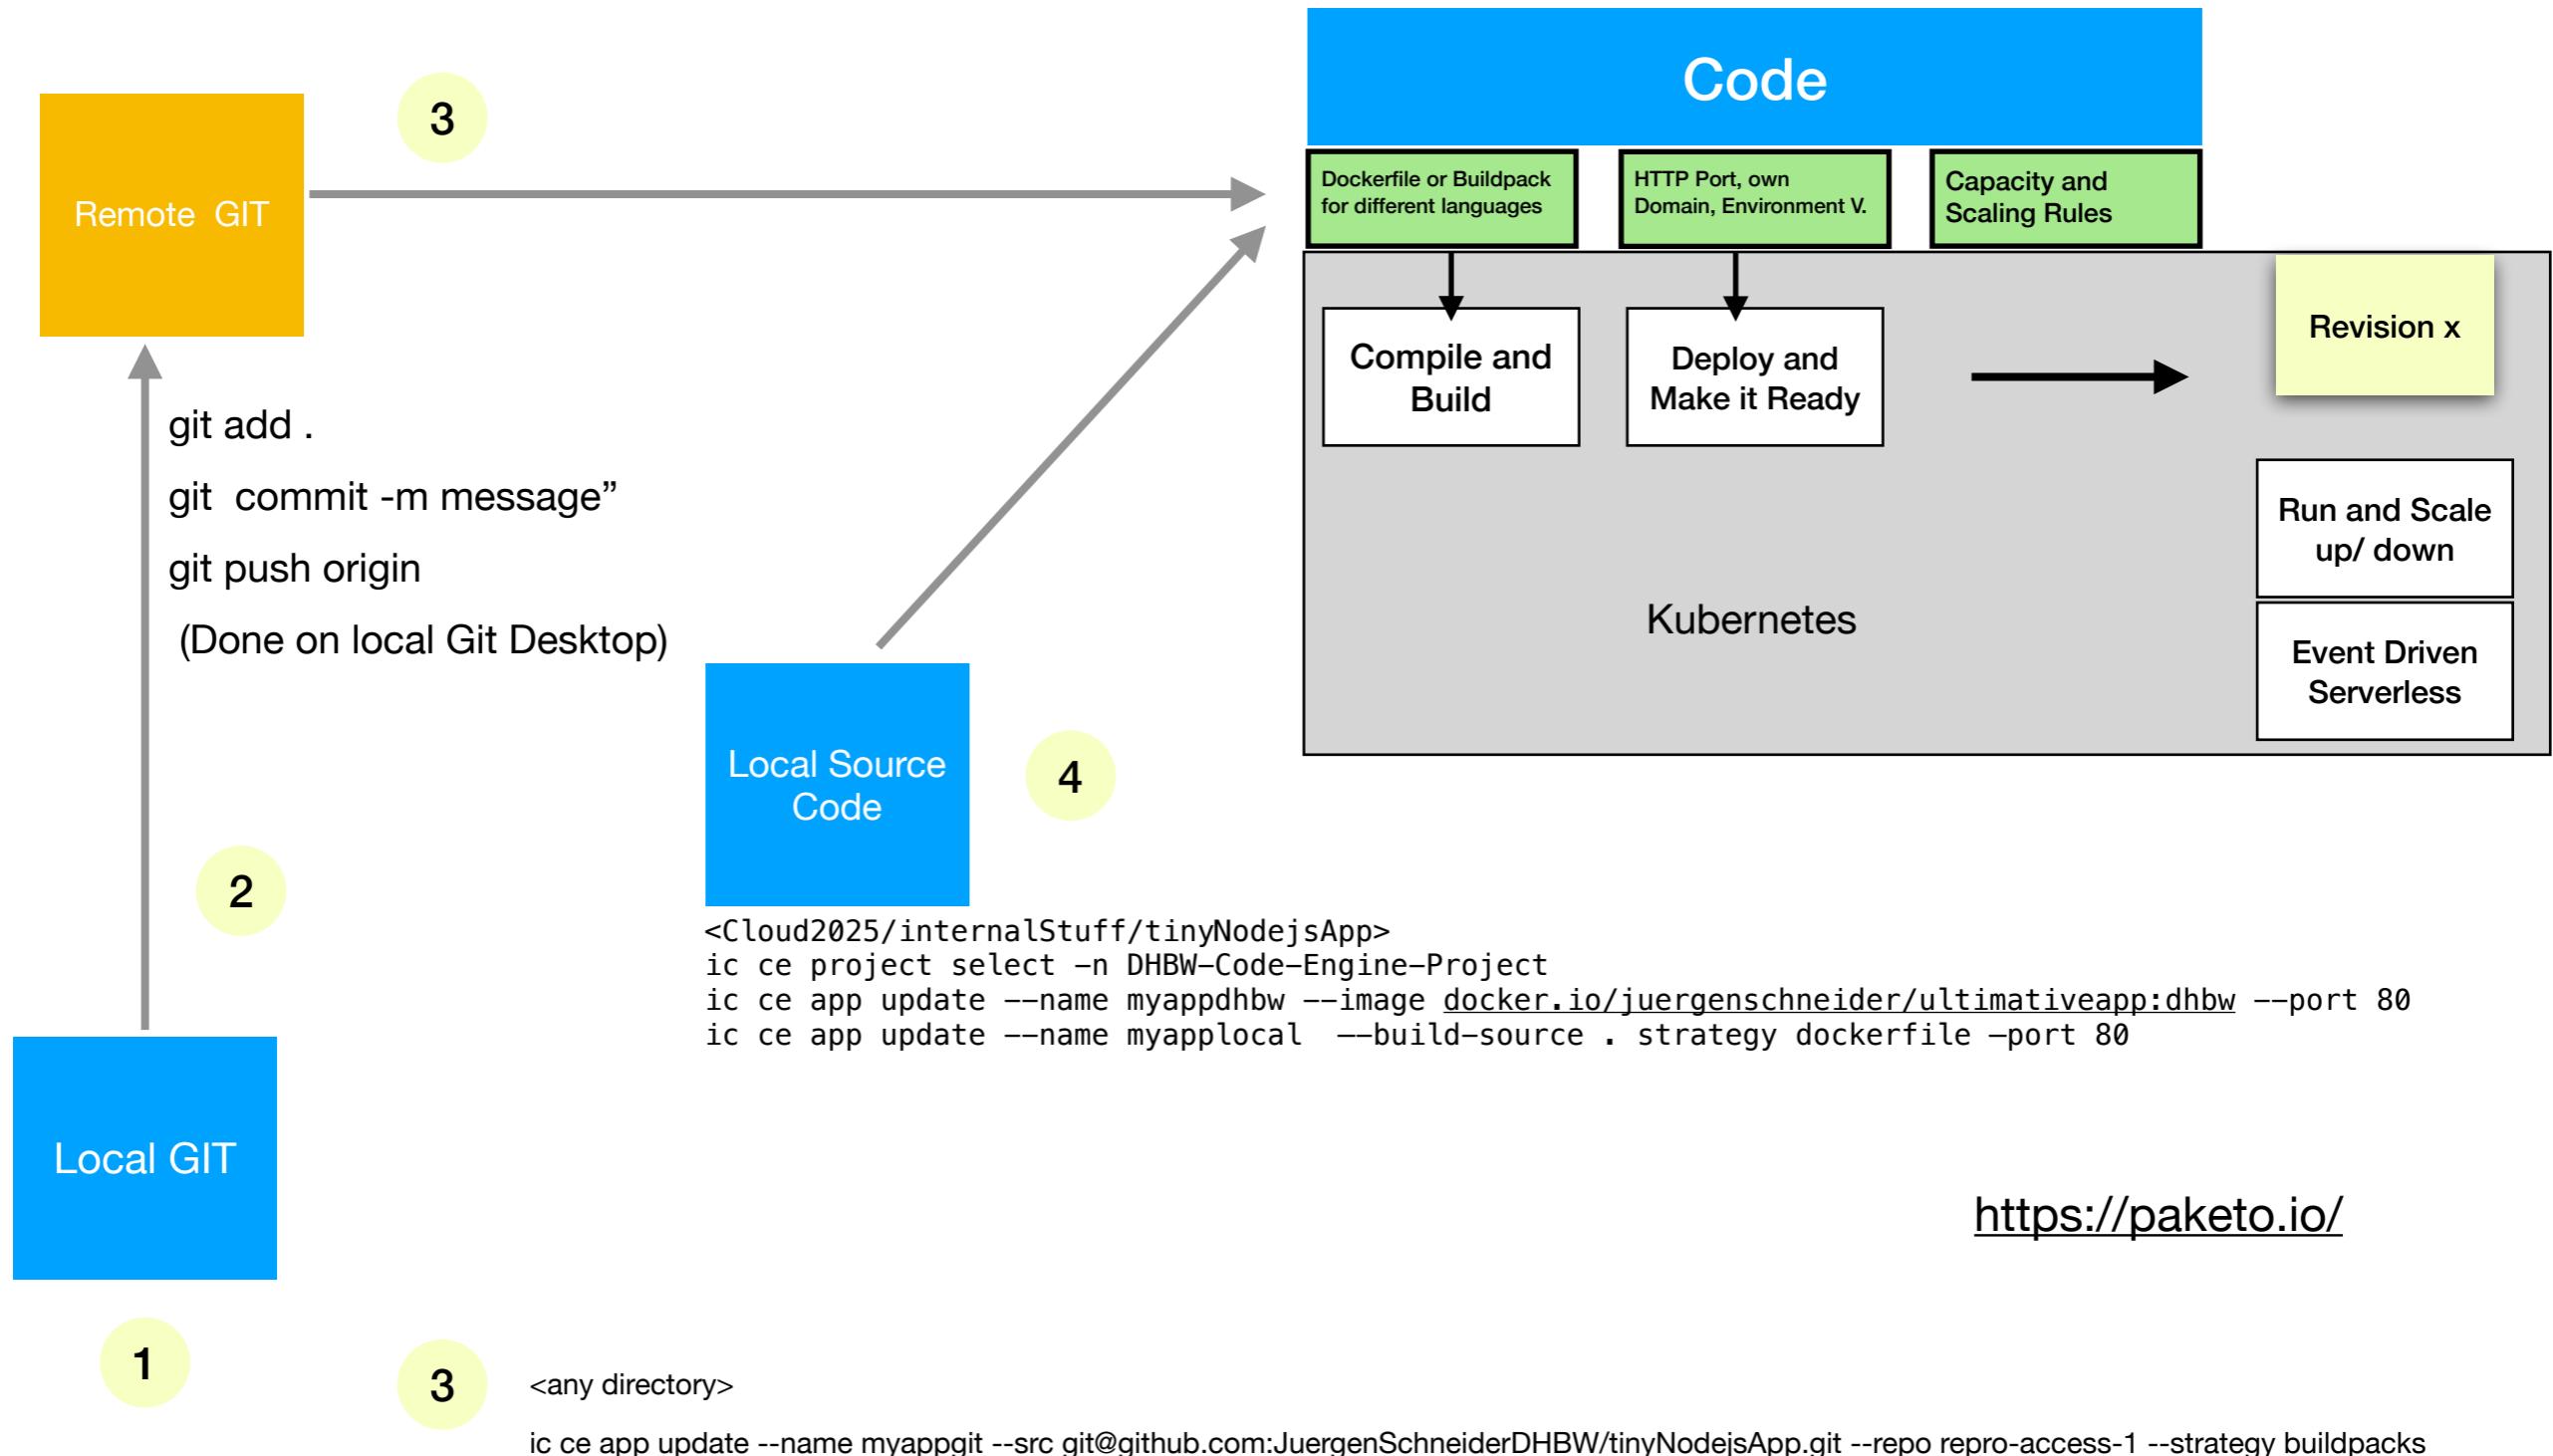


Show Code Engine GUI with Service Binding, Environment Variables

<https://ttsfunc.yk9dj48yz2s.eu-de.codeengine.appdomain.cloud?text=Geht+es+nun+wirklich+besser>

→ in Chrome oder Firefox.. Safari audio format leads to errors

Building an App either using a local build or as part of a Git Push



Summary on Workload Type

<https://cloud.ibm.com/docs/codeengine?topic=codeengine-cefunctions>

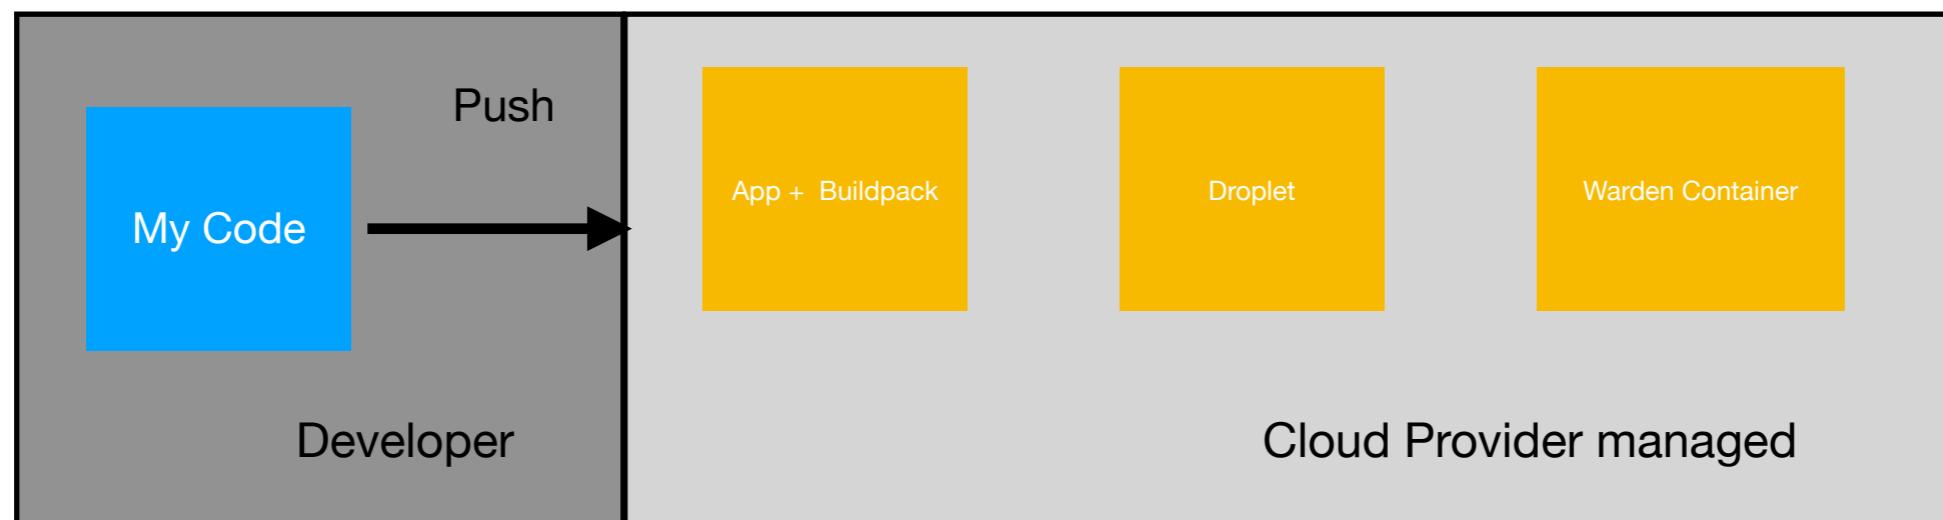
How do functions compare to apps and jobs?

Characteristic	App	Job	Function
Execution time (duration)	Long-running (10 minutes per request)	Long-running (up to 24 hours)	Short-running (2 minutes or less)
Startup latency	Medium	Scheduled start	Low
Termination	Run-continuously	Run-to-completion	Run-to-completion
Invocation	On request or permanently running	Scheduled	On request, instant
Programming Model	Container-based build and execution	Container-based build and execution	Language-specific source code files and dependency metadata
Parallelism	Parallel execution, flexible	Low to medium parallel execution	High parallel execution
Scale-out	Based on number of requests	Based on job workload definition	Based on events or direct invocations
Optimized for	Long running, highly complex workload and on-demand scale-out	Scheduled or planned workloads with high resource demands	Startup time and rapid scale-out

Collapse

Cloud Foundry

- History
 - Started in 2009 with a small team in VMware
 - First announced April 2011
 - April 2012 Usage of BOSH <https://www.bosh.io/docs/> as a development toolchain for large distributed systems to deploy Cloud Foundry PaaS
 - April 2013 Pivotal was founded by EMC and VMware to become the owner of Cloud Foundry
 - April 2014 Open Source governed.. with a broad coverage of important players (Platinum Level Cisco, Dell/EMC, IBM, SAP, Pivotal, Suse, VMware, Gold Level Alibaba, Google, Swisscom, VW, Ford, Microsoft)
 - <https://www.cloudfoundry.org/>
 - **Cloud Foundry lost the game against Kubernetes and their huge ecosystem to provide exactly this capabilities**
 - <https://krishnan.medium.com/lessons-from-the-demise-of-cloudfoundry-ed4f77a08a73>
- Here is the code, please build an executable and run it
 - CLI **cf push** (manifest.yml).
 - Build and Compile
 - **App + Buildpack** = Droplet (a sort of a image)
 - Run
 - **Droplet + CF Runtime** = App running in a container

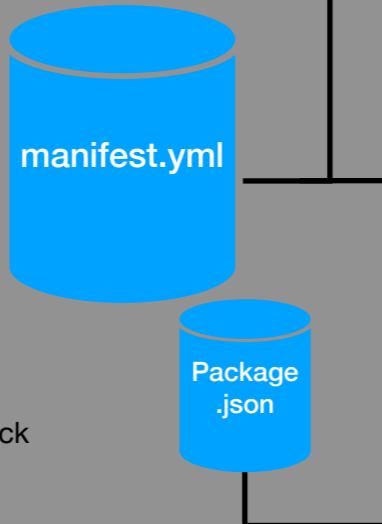


Cloud Foundry Basic Deploy Flow

CLI

```
$ ibmcloud cf push [options]
```

```
---  
applications:  
- name: ..  
  disk_quota: 1024M  
  path: .  
  instances: ..  
  memory: 256M  
  buildpack: nodejs_buildpack
```



Detailed flow later...

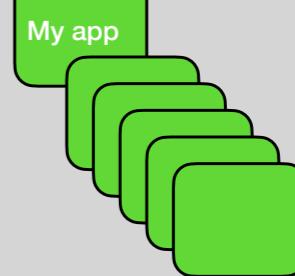
Upload Files

Download missing dependencies

buildpack

detect, build, run

CF Container



Buildpacks are basically a set of structured scripts to be used for Compile/Build and Run the Application

<https://docs.cloudfoundry.org/buildpacks/understanding-buildpacks.html>

Cloud Foundry Buildpack

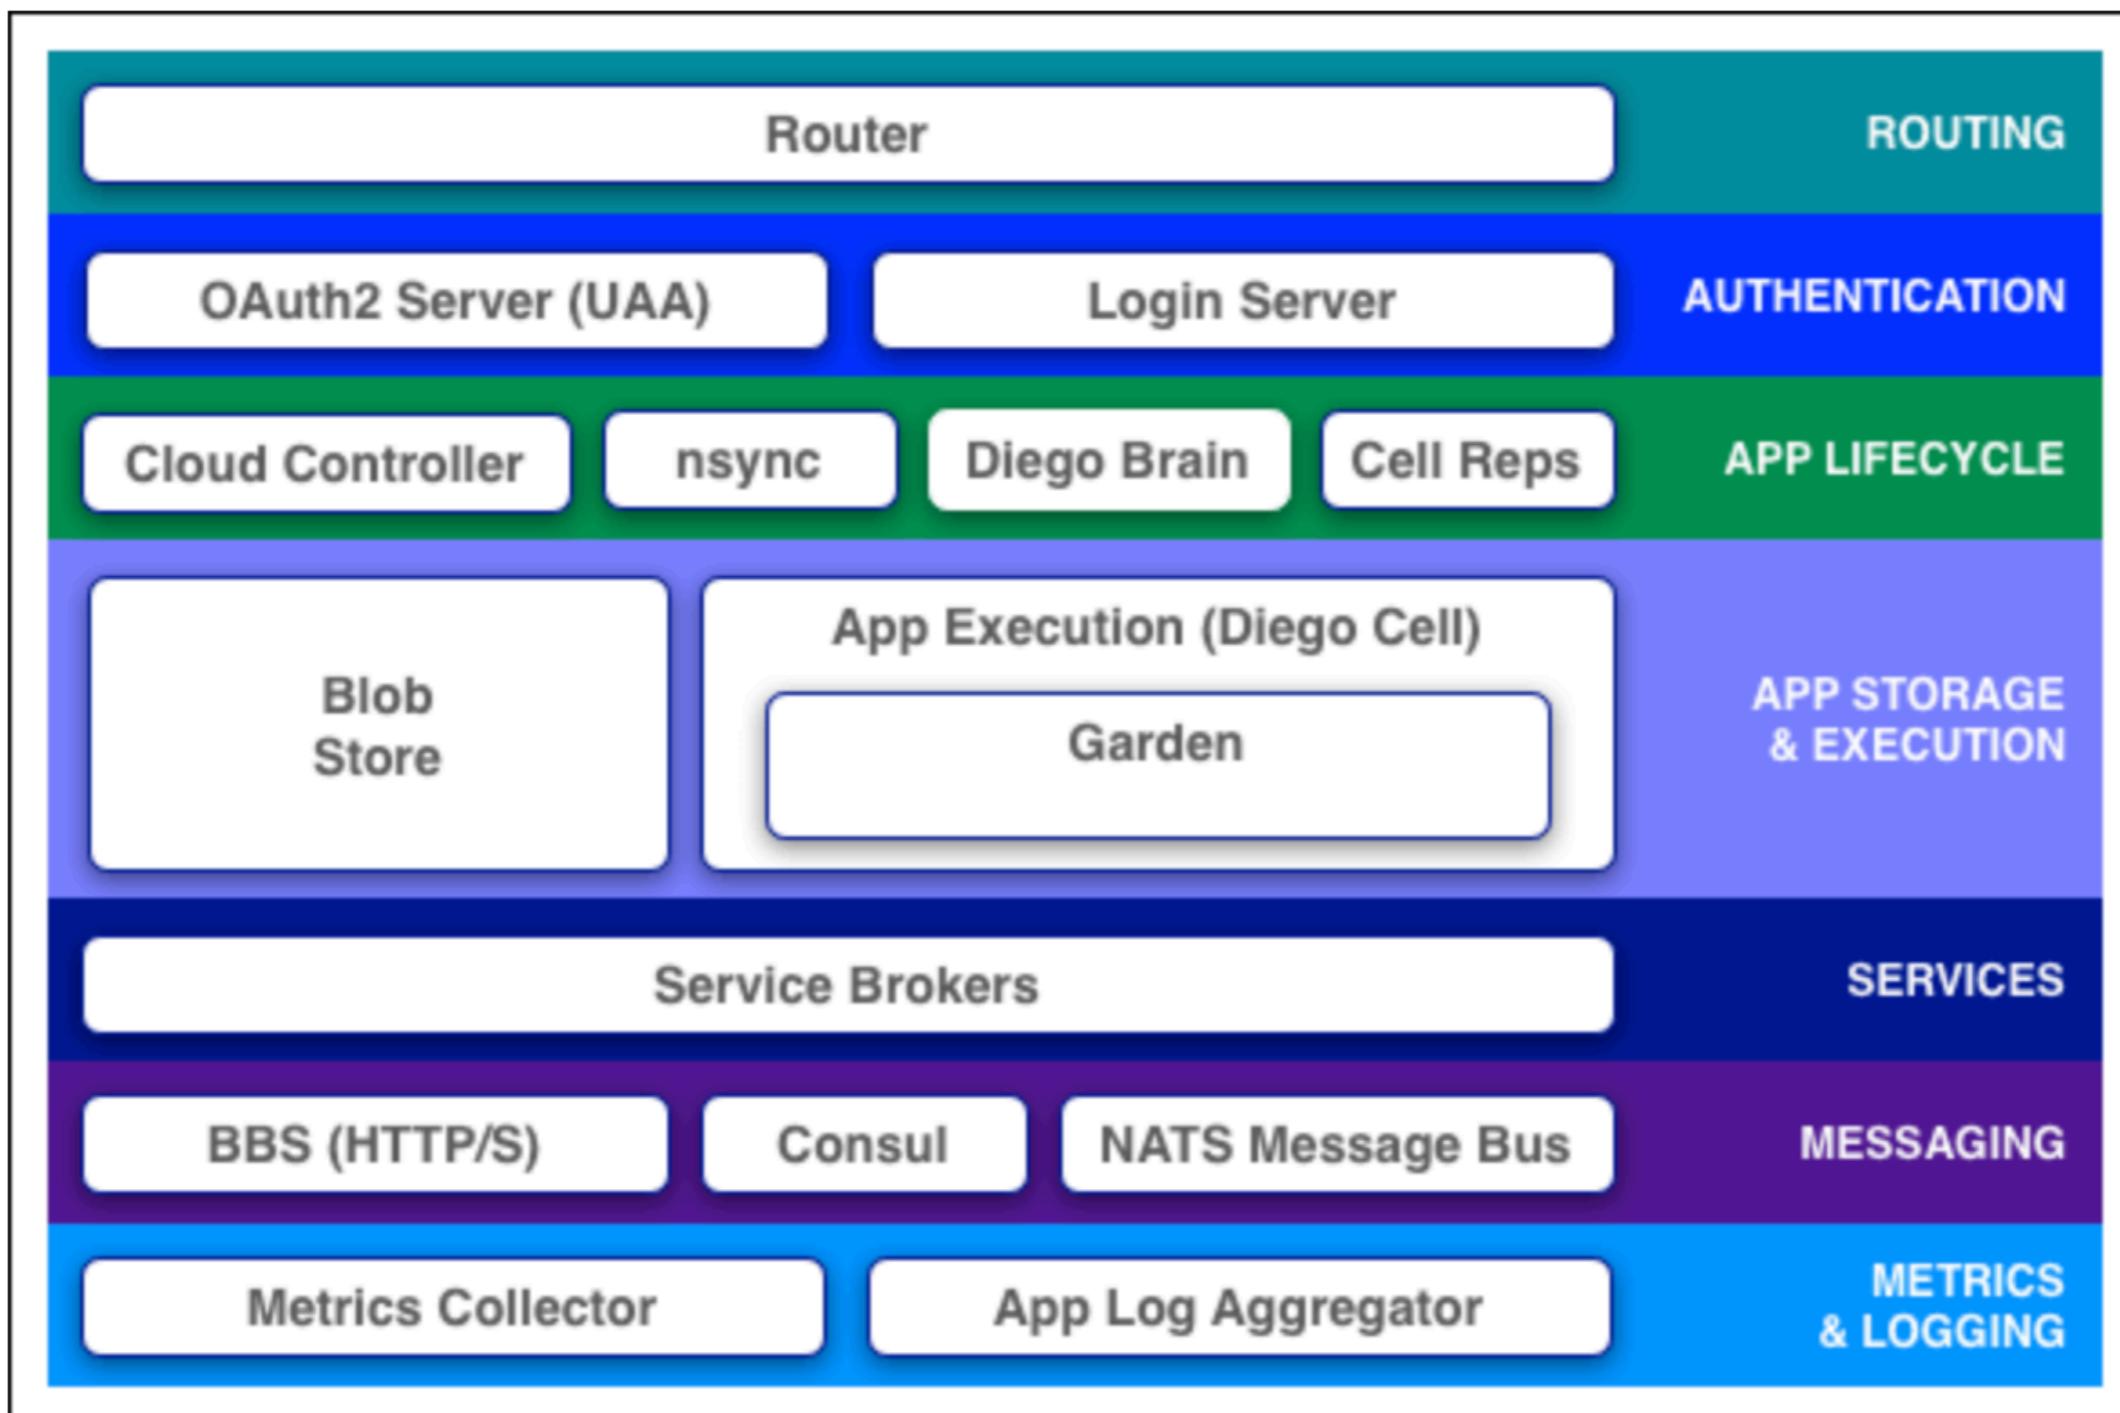
- provide the runtime support for apps (Result of a build back is a droplet (a ‘image’ ready to get deployed))
- examine the apps to determine what dependencies to download and how to configure the apps to communicate with bound services.
 - In the core it is a set of script flows as described here
 - **bin/detect** checks the code directory to figure out whether the right build pack has been chosen
 - In case of node.js it tests whether a package.json exists
 - **bin/supply** provides all the dependencies
 - In case of node.js it executes the npm install command
 - **bin/finalize** prepares the app for launch
 - In case of node.js it basically setup some environment variable and the final directory
 - **bin/release** provides feedback metadata to Cloud Foundry indicating how the app should be executed.
 - In case of node.js it tells Cloud Foundry to start node.js with NPM start
- Ready to use build packs are available for
 - Ruby, Java, Go, .NET, PHP, Node.js, Python, Staticfile (a NGINX Web Server for static Web Pages (CSS, HTML, JS))

Cloud Foundry Manifest.YML

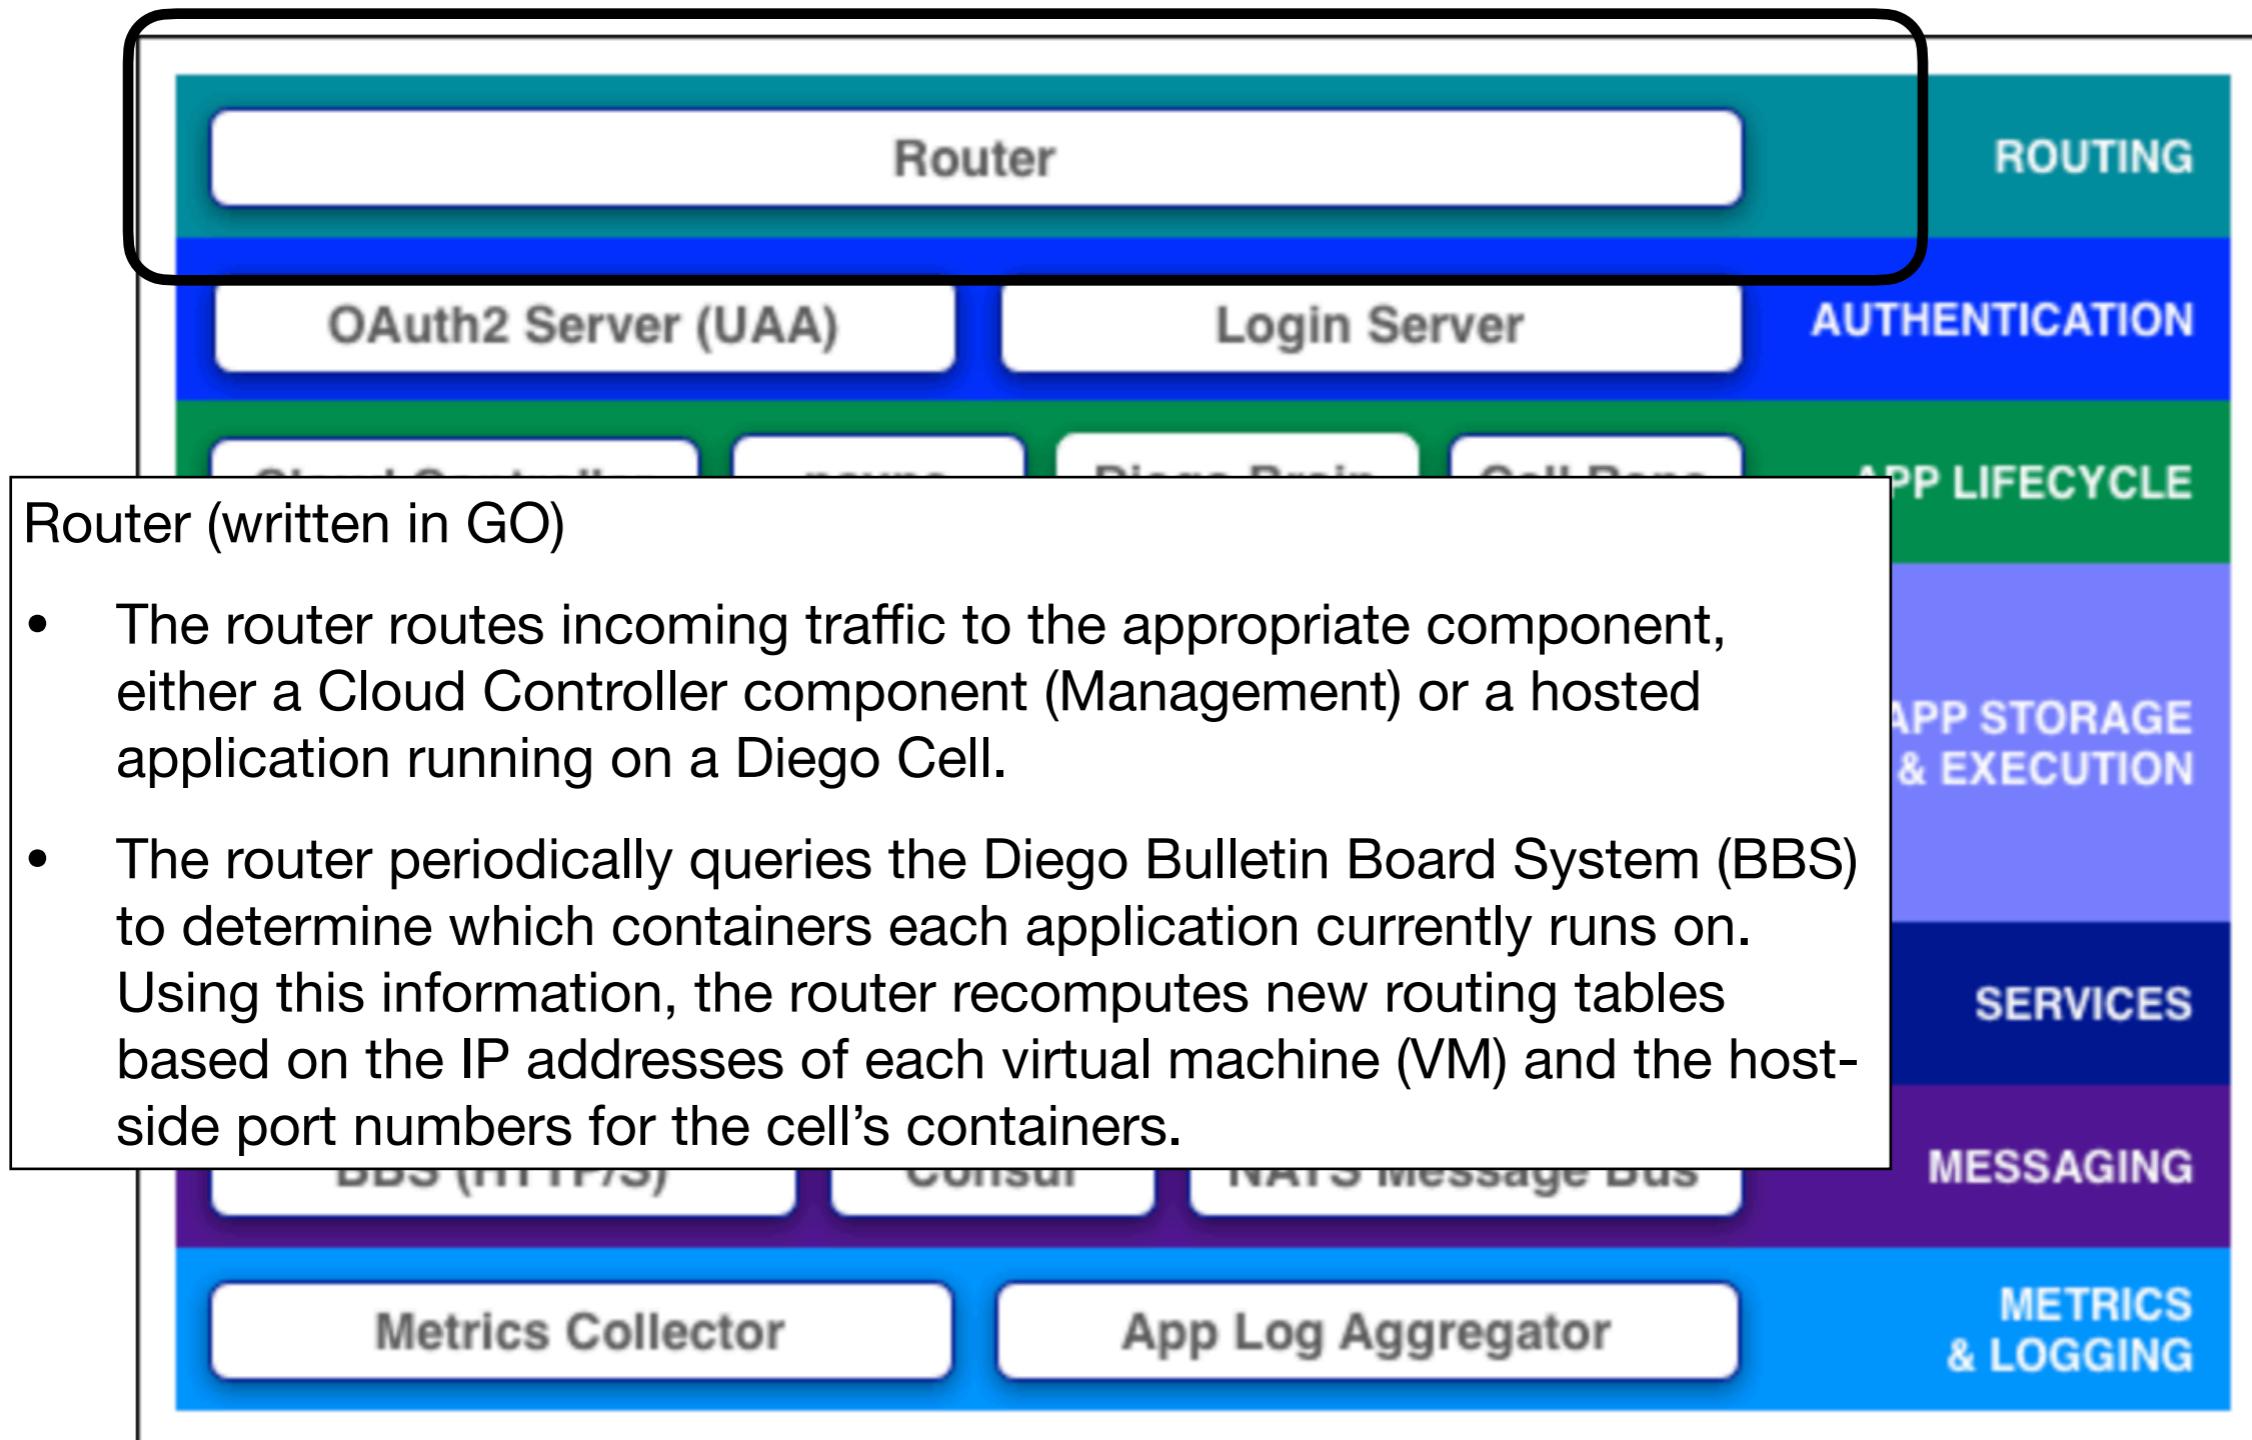
- Basically all the options you can apply to the CF Push command
 - Application
 - Name of the app (which becomes the subdomain in the dns entry)
 - Memory/CPU, Storage Quota, # of instances
 - Name of the buildpack
 - Start Command (if needed, in node.js this is typical part of the package.json)
 - Environment Variables
 - Directory Path if not just ./
 - Sidecars
 - Services used
 -
 - CF Push -f manifest

```
---  
applications:  
- name: my-app  
  memory: 512M  
  instances: 2
```

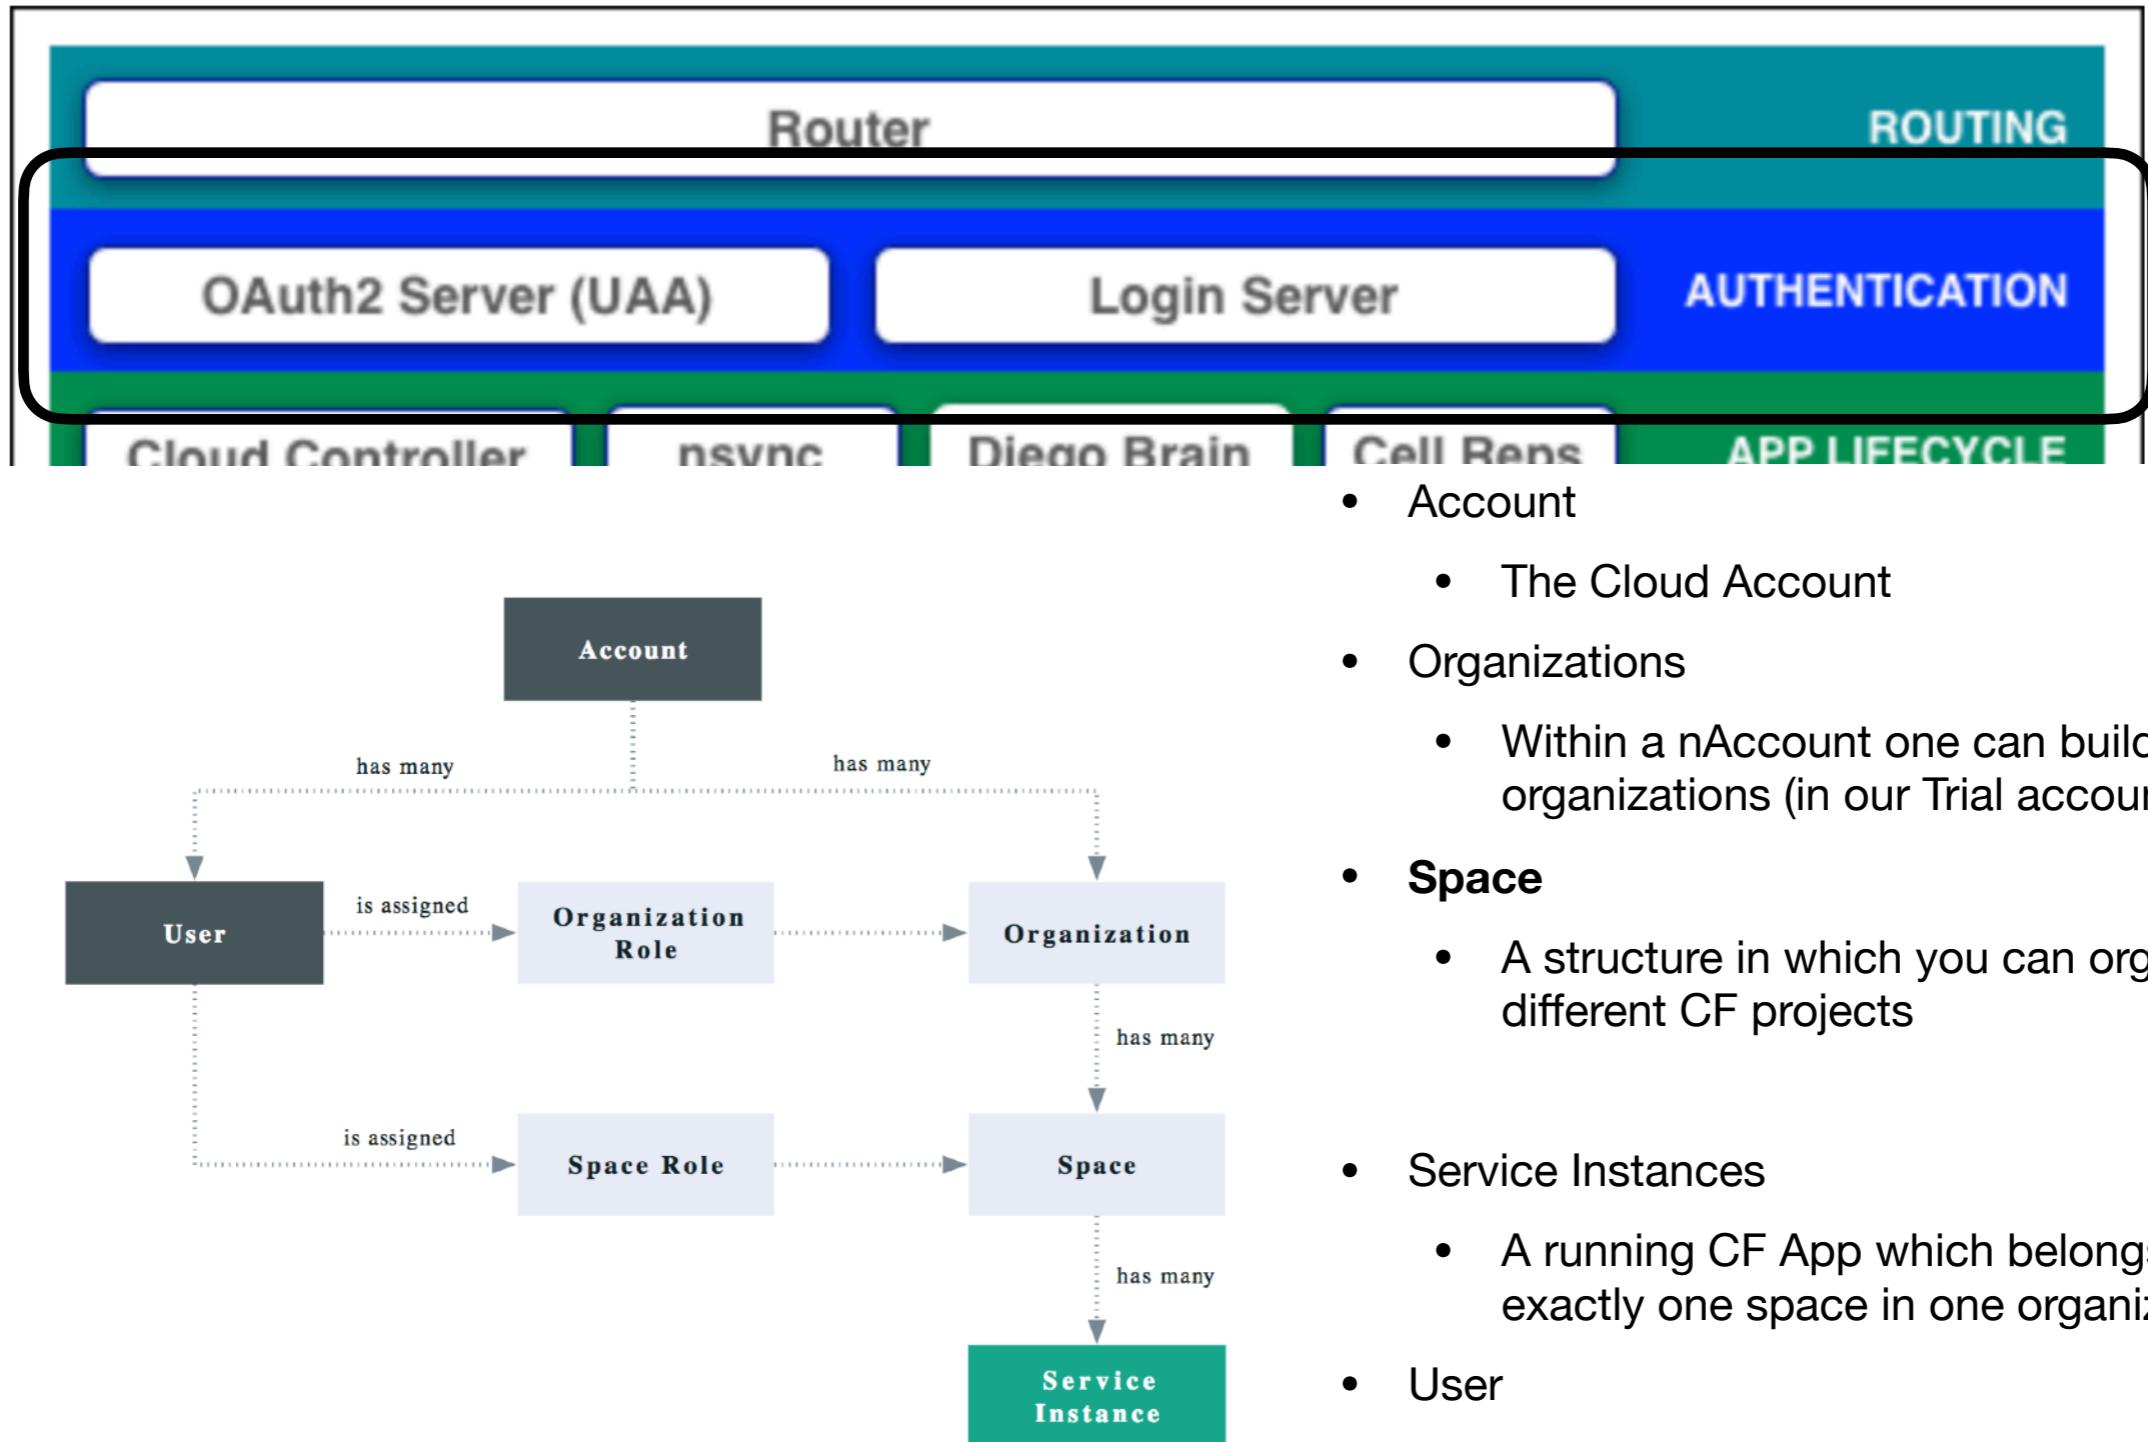
Cloud Foundry Architecture Component Overview



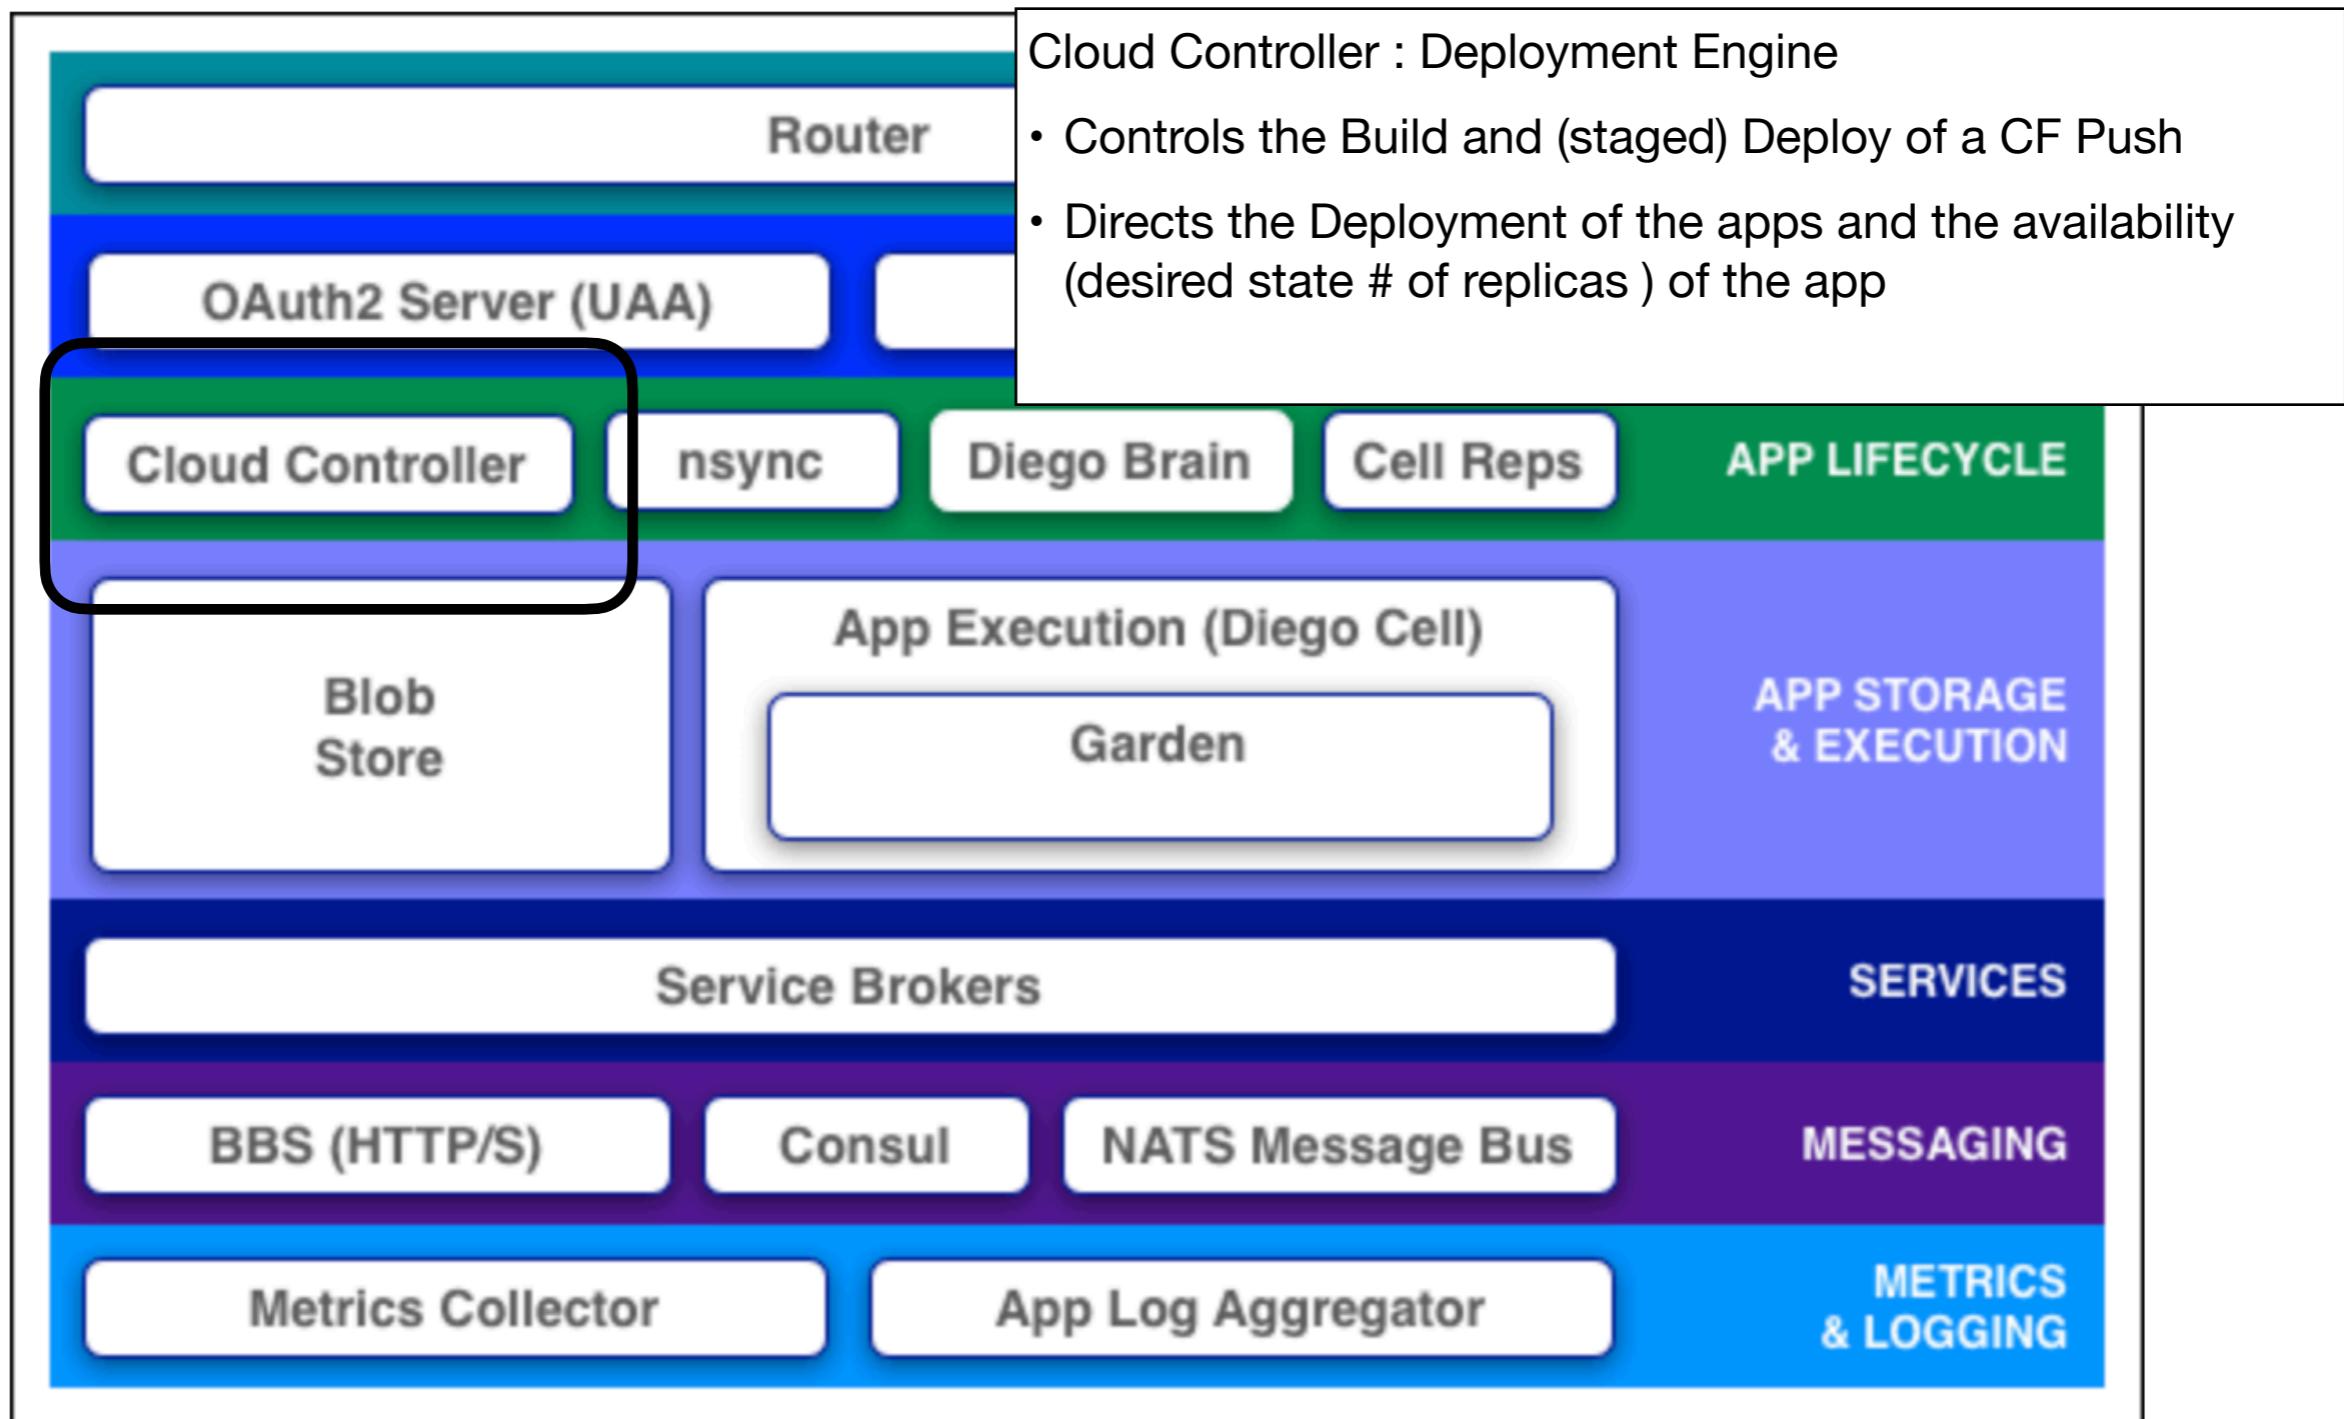
Cloud Foundry Architecture Component Overview



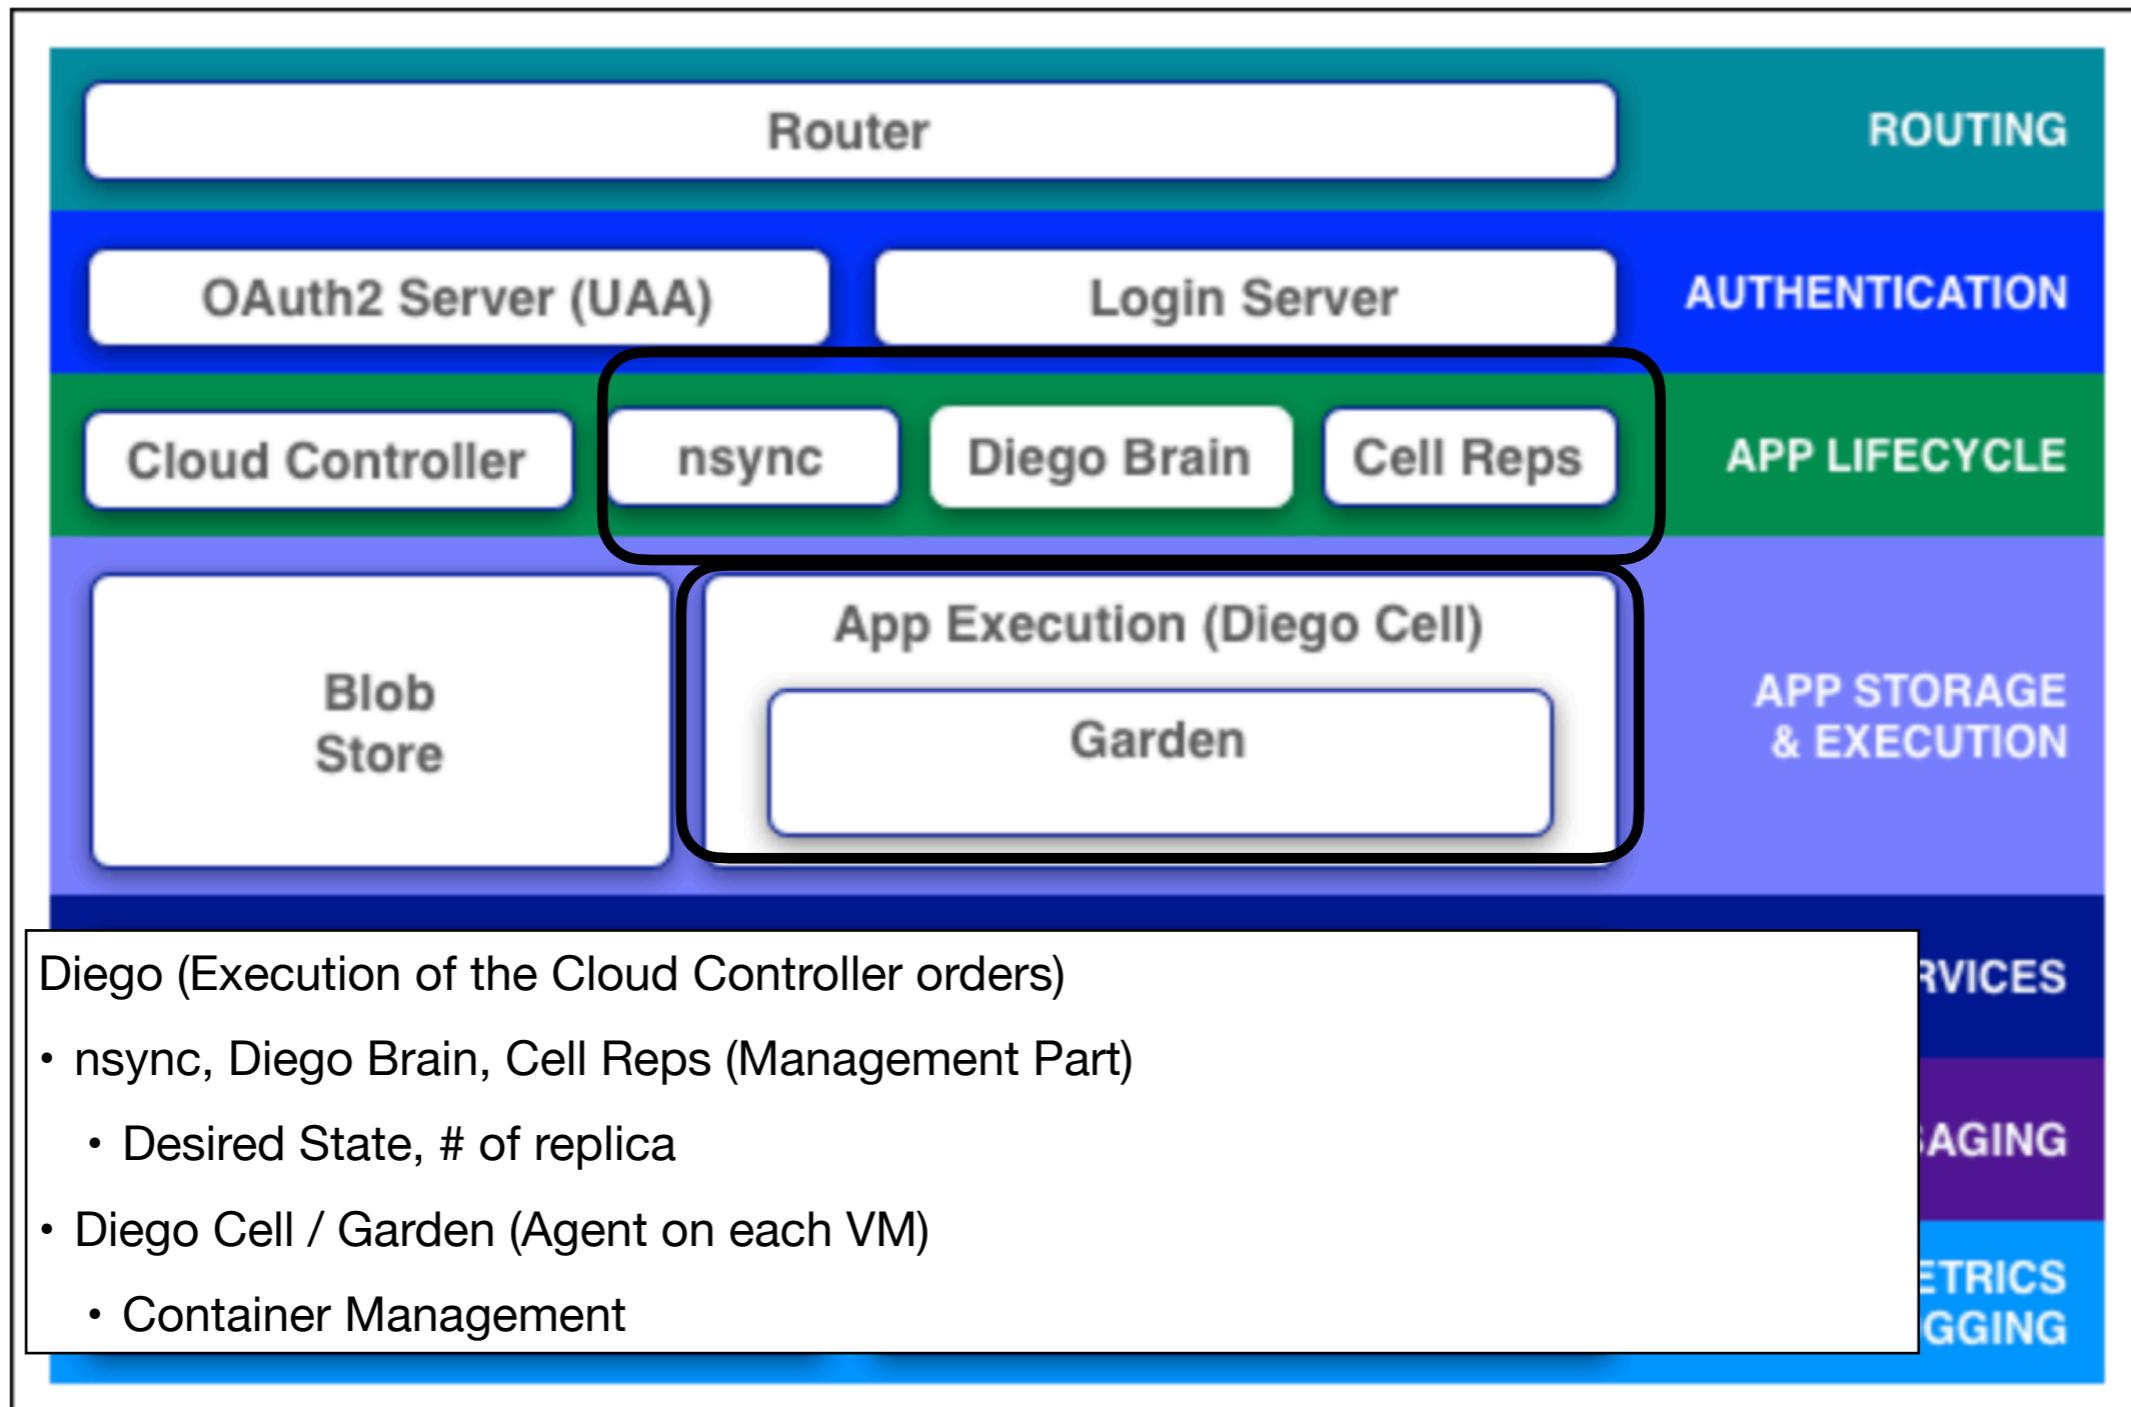
Cloud Foundry Architecture Component Overview



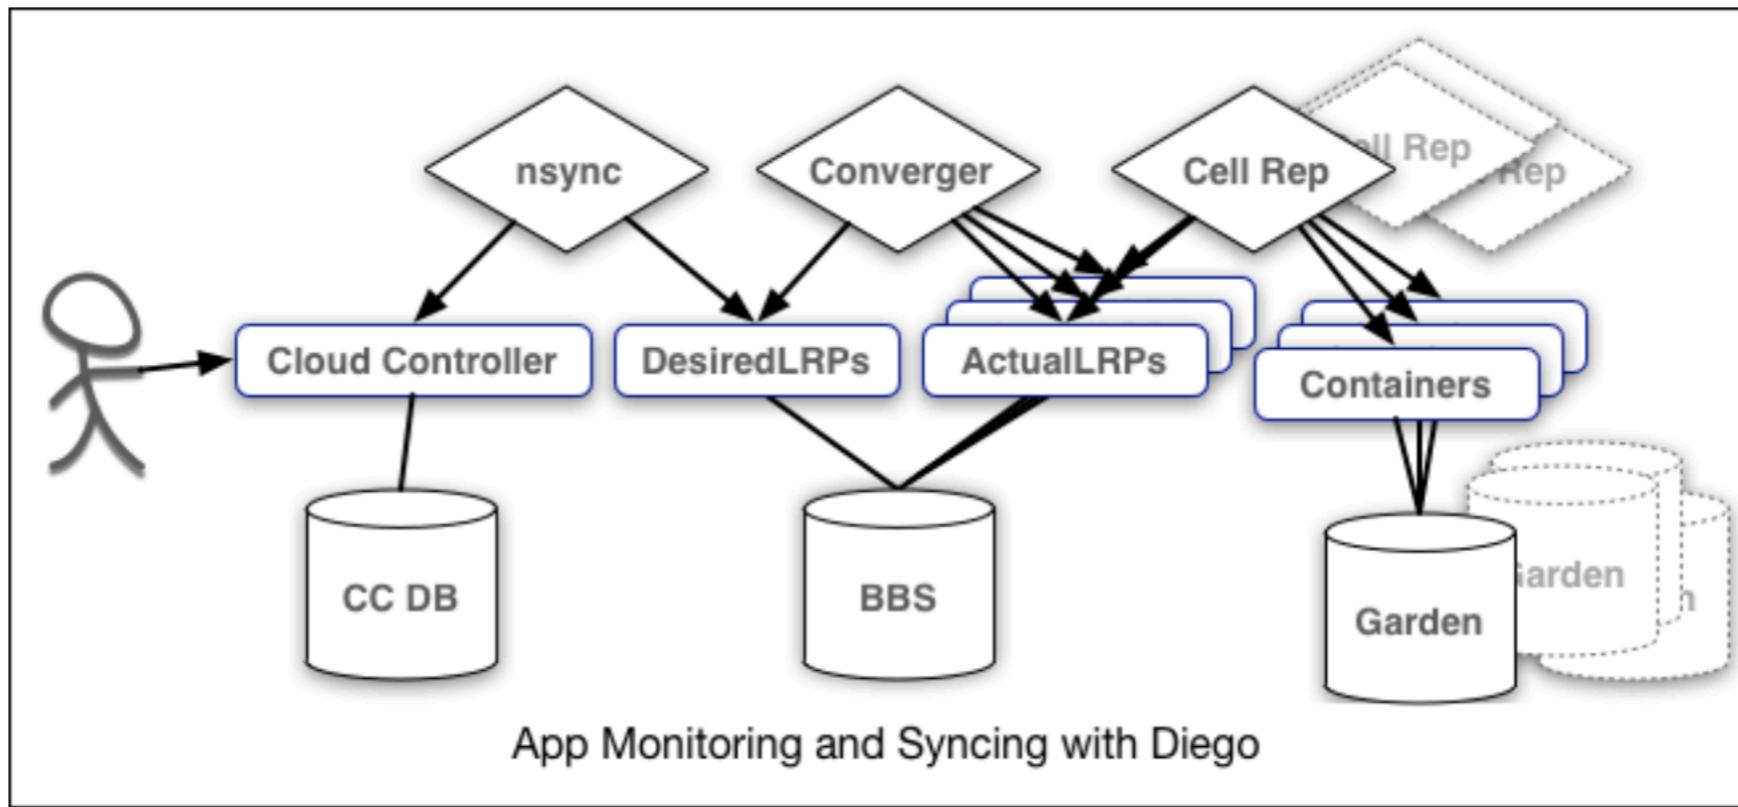
Cloud Foundry Architecture Component Overview



Cloud Foundry Architecture Component Overview

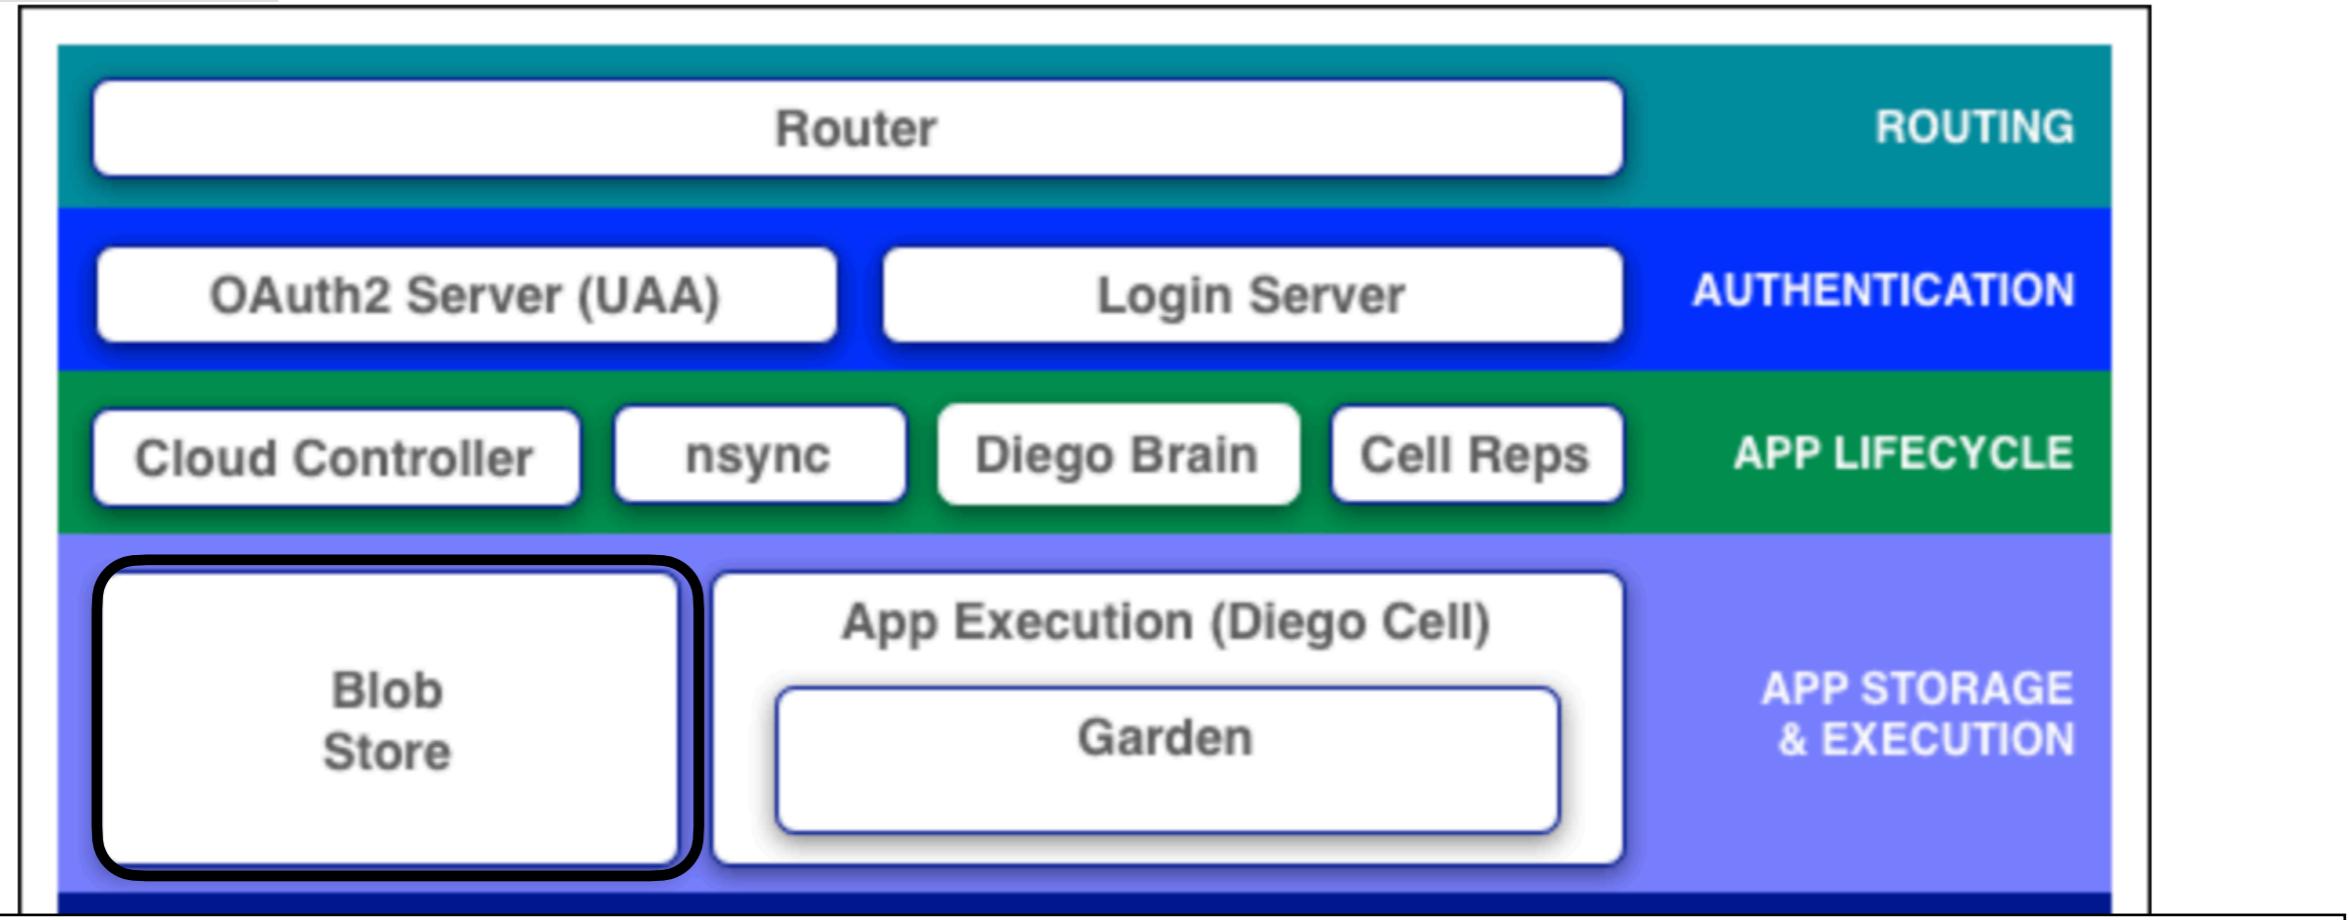


Desired State Management



- Cloud Controller directs the system wrt. # of active instances (LRP = Long running Process, normally a App instance)
- Converger (part of BBS) compares desired and actual # and either kills or launches new containers via the Diego cell) (BBS = Bulletin Board System is basically a data store for frequently updated and disposable data such as cell and application status, unallocated work, and heartbeat messages. The BBS stores data in MySQL)
- Cell Rep (one per VM) monitors the health of the containers
- CC DB BBS are persistent stores
- **Garden is the new container management backend it supports Linux Containers, runC (Open Container Initiative), Windows and Docker**
- Communications among the components via HTTP Messages

Cloud Foundry Architecture Component Overview



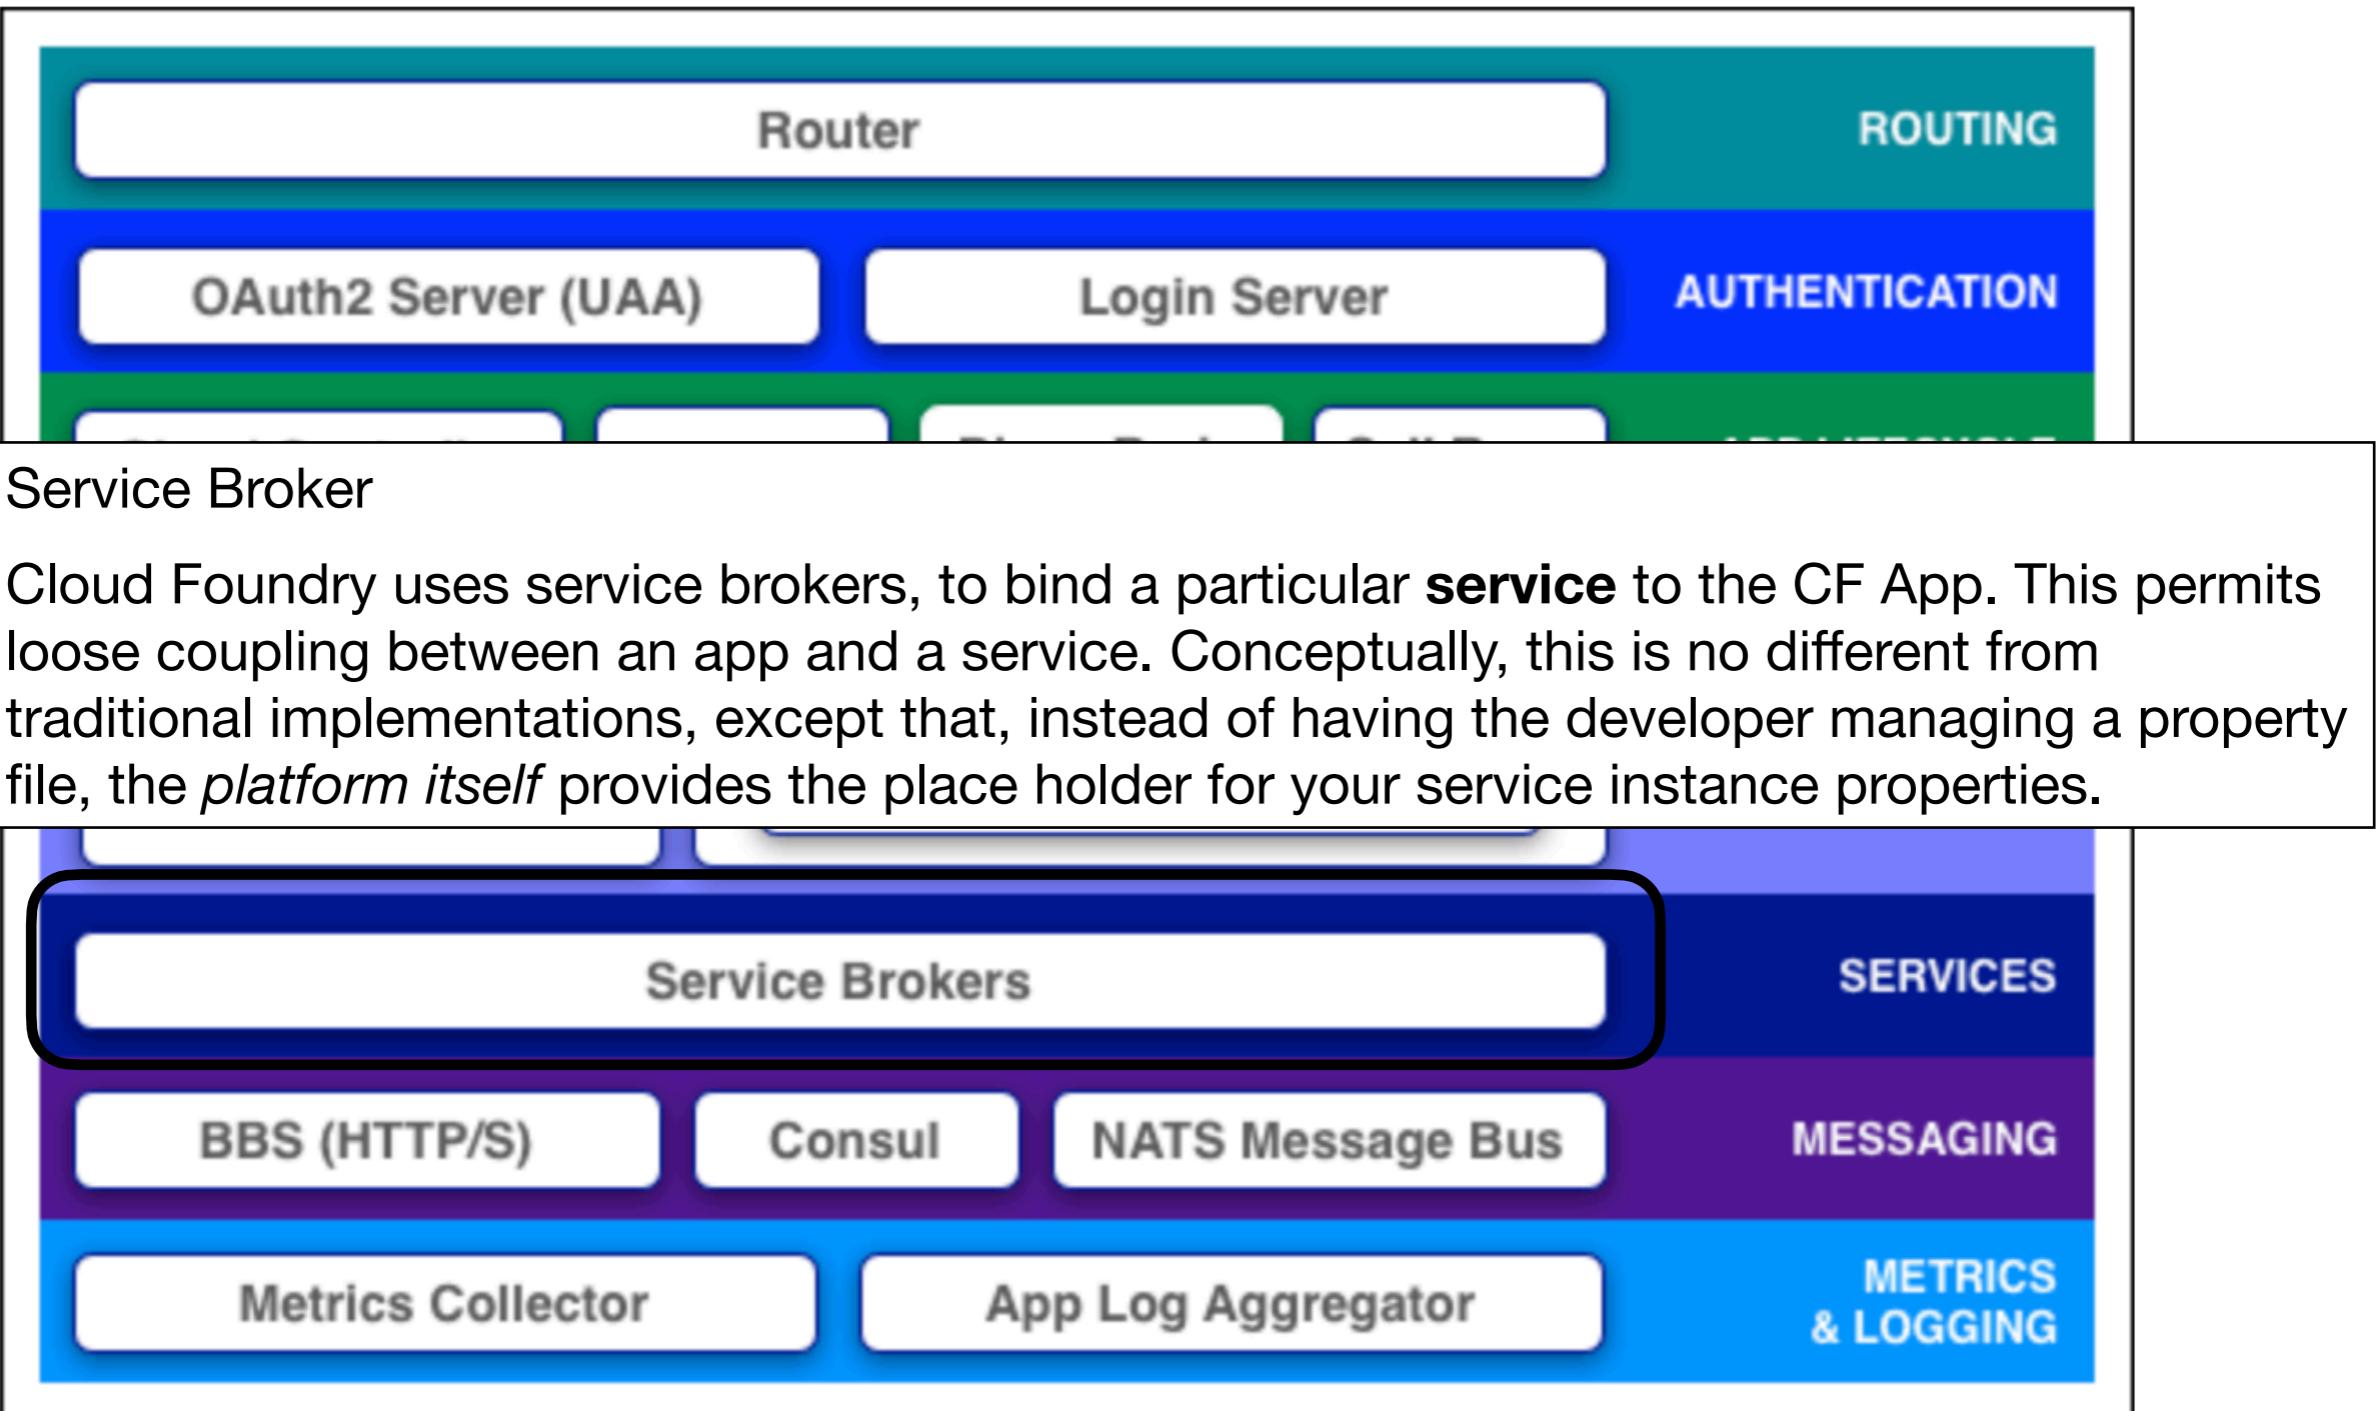
The blobstore is a repository for large binary files, the blobstore contains:

- Application code packages
- Buildpacks
- Droplets

You can configure the blobstore as either an internal server or an external S3 or S3-compatible endpoint.

- blob ? Distributed cloud applications often need to handle large data elements, also referred to as binary large objects (blob). Examples are virtual server images managed in an [Elastic Infrastructure](#), pictures, or videos.

Cloud Foundry Architecture Component Overview



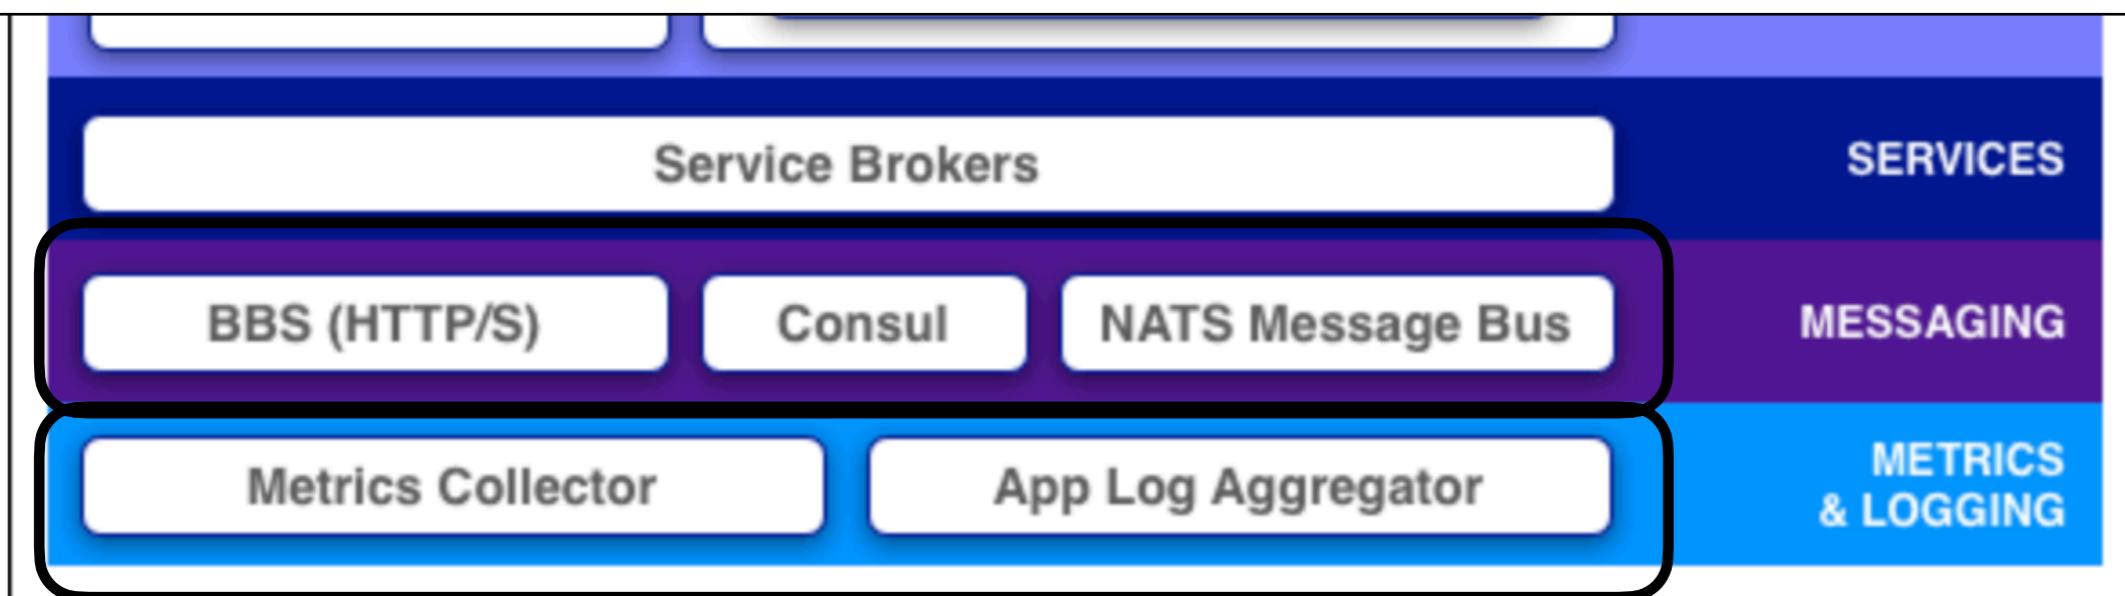
Cloud Foundry Architecture Component Overview



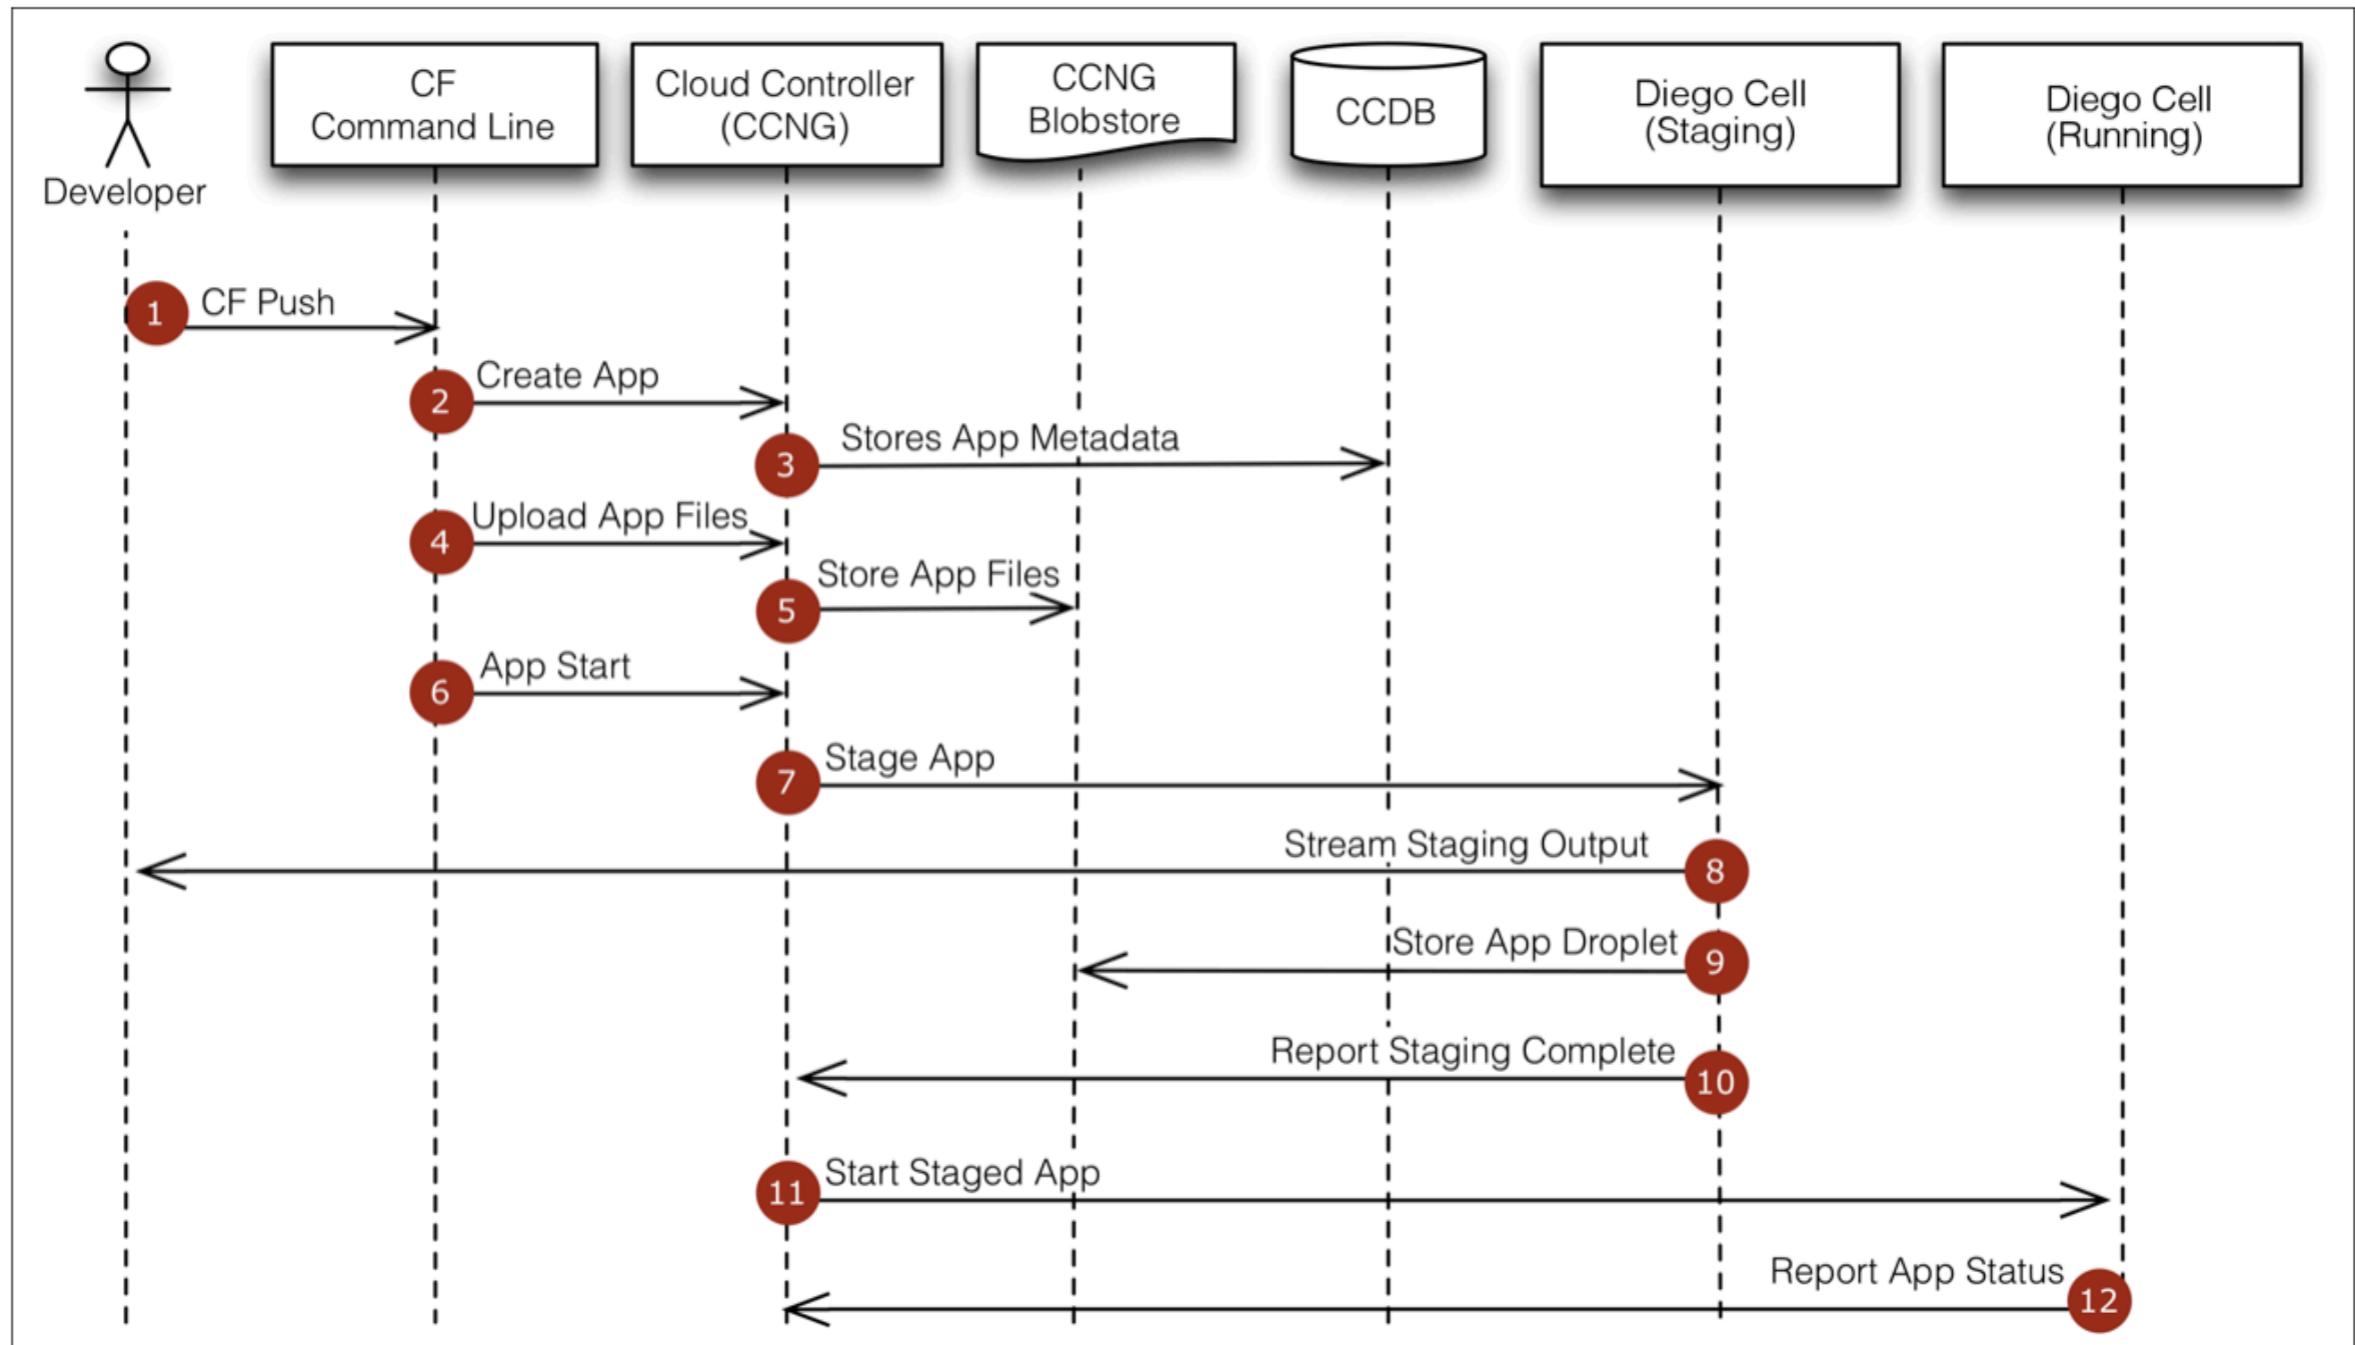
Messaging

Cloud Foundry component VMs communicate with each other internally through HTTP and HTTPS protocols, sharing temporary messages and data stored in two places.

The Loggregator (log aggregator) system streams application logs (more than one instance) to developers. The metrics collector gathers metrics and statistics from the components.



The Code Deploy Flow



Cloud Foundry Manifest.YML

- Cloud Foundry does support docker images as well..
 - Finally CF push the Docker Image

```
[Juergens-MBP:tippen juergenschneider$ bx cf push tippendocker --docker-image=juergenschneider/tippendocker:latest --no-manifest
Invoking 'cf push tippendocker --docker-image=juergenschneider/tippendocker:latest --no-manifest'...

Pushing app tippendocker to org DHBW Cloud Computing / space DHBW Exercises as jschnei1@lehre.dhbw-stuttgart.de...
Getting app info...
Creating app with these attributes...
+ name:          tippendocker
+ docker image:  juergenschneider/tippendocker:latest
  routes:
+   tippendocker.eu-gb.mybluemix.net

Creating app tippendocker...
Mapping routes...

Staging app and tracing logs...
Cell 818ffad7-d39c-4424-90d2-dfa76f97df99 successfully created container for instance 5db9716a-c836-4787-9ecd-7a72324de45a
  Staging...
  Staging process started ...
  Staging process finished
  Exit status 0
  Staging Complete
Cell 818ffad7-d39c-4424-90d2-dfa76f97df99 stopping instance 5db9716a-c836-4787-9ecd-7a72324de45a
Cell 818ffad7-d39c-4424-90d2-dfa76f97df99 destroying container for instance 5db9716a-c836-4787-9ecd-7a72324de45a

Waiting for app to start...

name:          tippendocker
requested state: started
instances:      1/1
usage:         1G x 1 instances
routes:        tippendocker.eu-gb.mybluemix.net
last uploaded: Thu 25 Oct 15:01:53 CEST 2018
stack:          cflinuxfs2
docker image:  juergenschneider/tippendocker:latest
start command: node myWebServer.js 8080

  state  since            cpu    memory   disk    details
#0  running  2018-10-25T13:03:05Z  0.0%  0 of 1G  0 of 1G
```

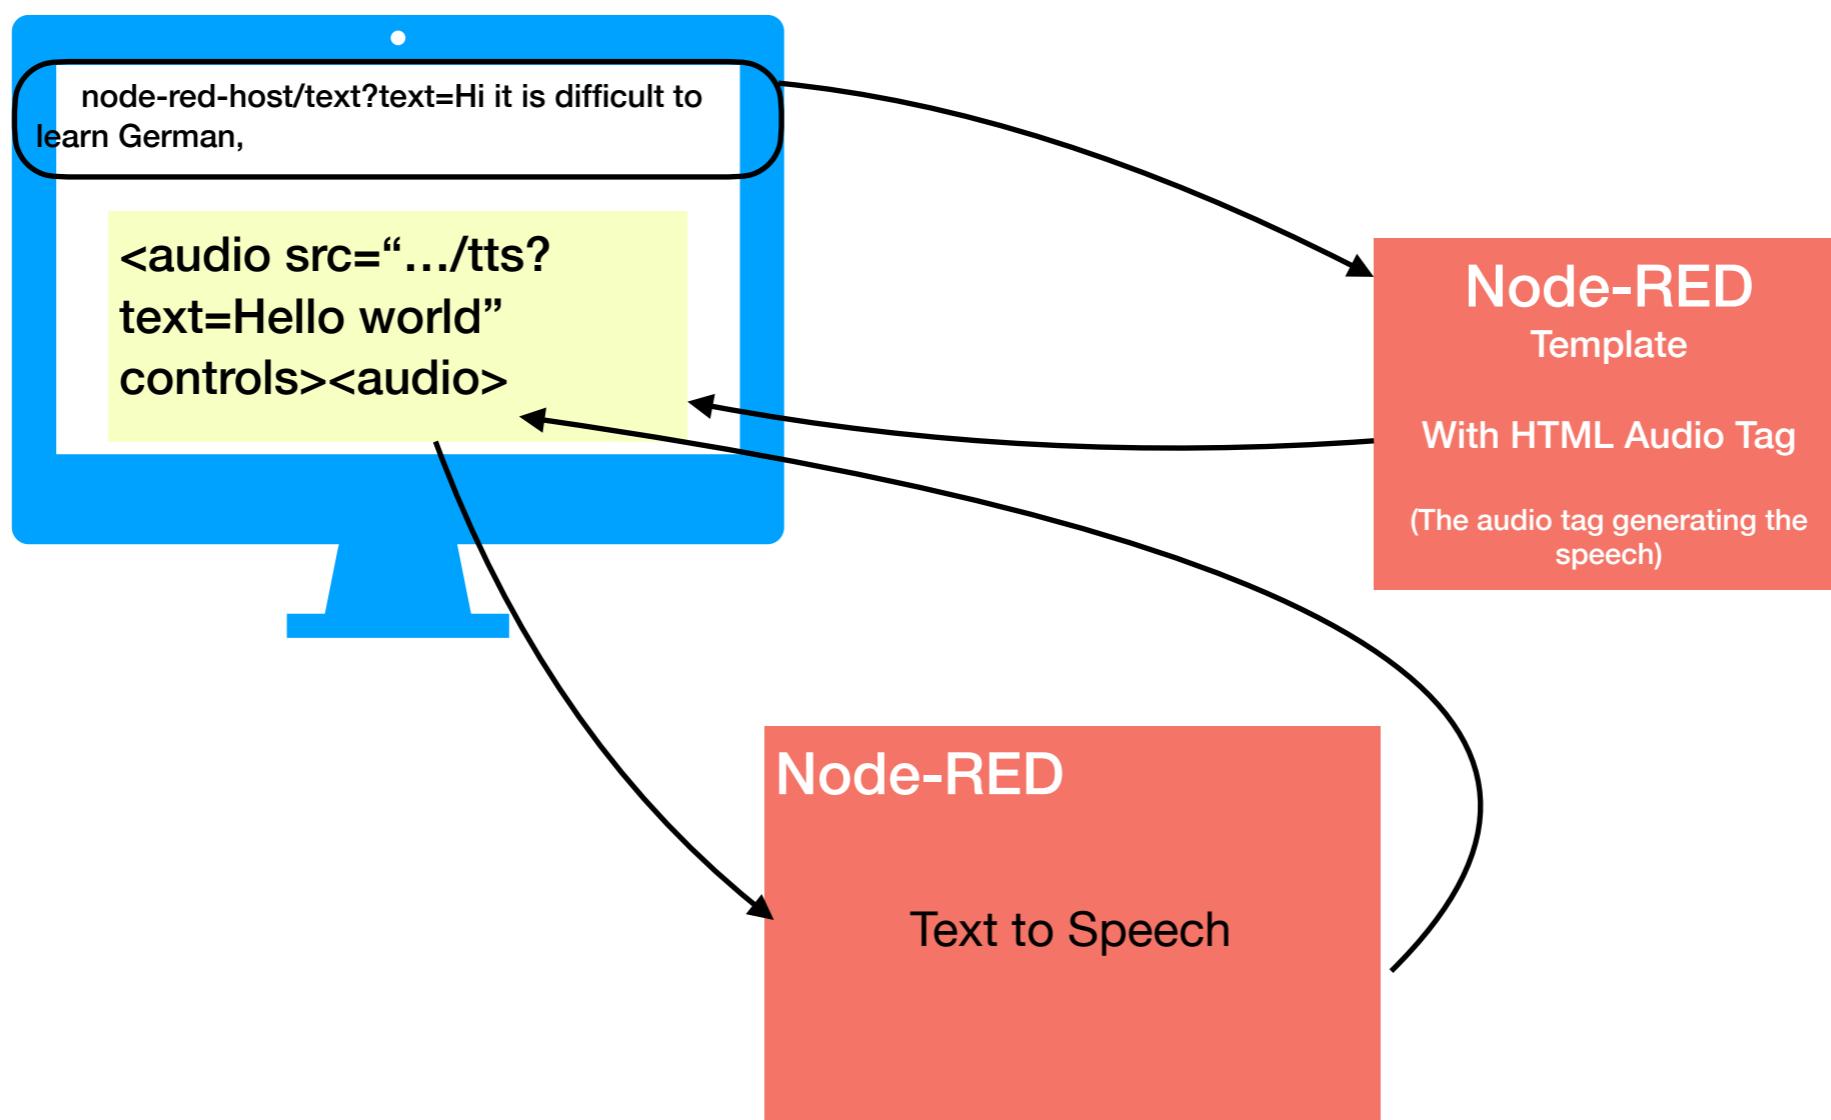
Approaching SaaS Writing Apps using AI Services

Node-Red Action Flow

- Often Apps are a flow of SaaS Services.
- Either you write them by yourself or use an action Modeller like Node-Red. It easily enables you to “**configure**” a Web Server (instead of writing your own Node.js).
- Node-Red is a node.js implementation with those modelling capabilities
- See Demo..
- A Node-RED Instance can be initiated
 - via Terraform in the bwCloud
 - Lösungen/bwCloud/nodered
- Admin. <http://<bwcloud-ip:1880>
- Services typically as http get with query string
 - Examples of flows in /internal/node-red

Node-Red Transaction Flow

- Implement a simple Text to Speech use case via Node-Red
 - Flow Example in /internal/node-red
 - The HTML element which produces the the audio file must be changed to the new bwCloud Instance
 - The url is /text?text= and the text must be in English
 - There are two flow first ist to receive the the text to generate the HTML audio which it self will invoke the TTS Service



IaaS++(CaaS)

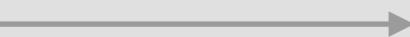
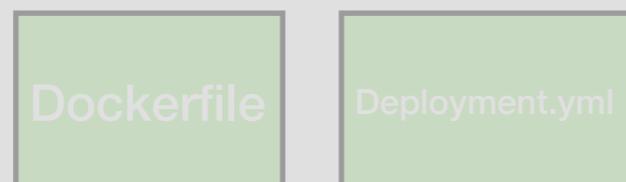
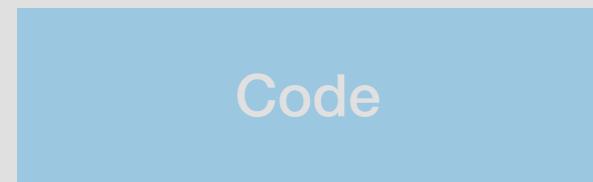
PaaS

PaaS++(FaaS)

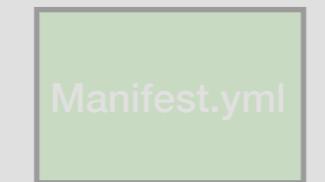
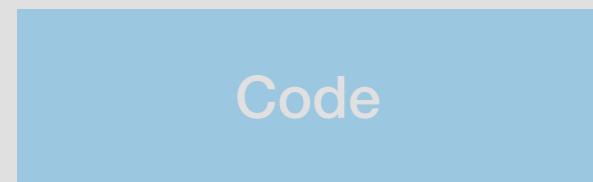
Dev / Ops (User) provides it

Somebody else manage it

Kubernetes + Toolchain



Cloud Foundry



OpenWhisk / Lambda / OpenFaaS / IBM Code Engine



Code



Event Driven Processing - Focus on Function

- Small and short program functions which are triggered by an event, execute that logic for that event and disappear.
 - Examples are:
 - Execute app logic in response to database triggers
 - Execute app logic in response to sensor data
 - Execute app logic in response of a HTTP API call
 - Execute app logic in response to scheduled tasks
 - For many of these scenarios, event-driven programming models are a better fit for **serverless architectures that can provide temporary resources on demand for intermittent tasks.**
- Cost
 - Applications deployed using a **serverless** architecture rely on a platform that provides the resources needed **on demand, at that very moment**. This obviously means that business logic execution can be mapped **to actual compute time used** (in milliseconds) rather than memory reserved for anticipated usage (by gigabyte per hour).
 - Good if those actions are ‘short’ (e.g. less than a minute) and the events happen infrequent or not equally distributed over time.
- Serverless is a bit misleading, better is to say the code (= function + MiniConfig) is in the cloud but the server (=Container) is built on demand (triggered by an event)

Some possible Scenarios

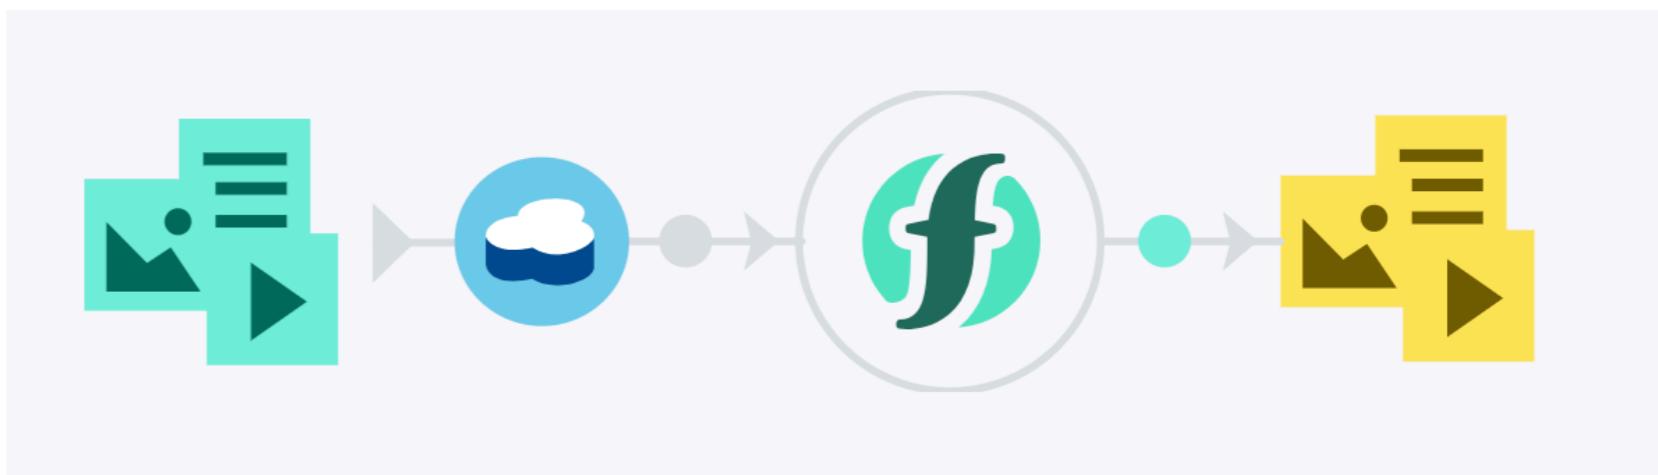
Serverless Backends

Expose application logic by implementing serverless microservices. Simply map your functions to well-defined API endpoints any client can call by making use of Web Actions or API Gateway integration.



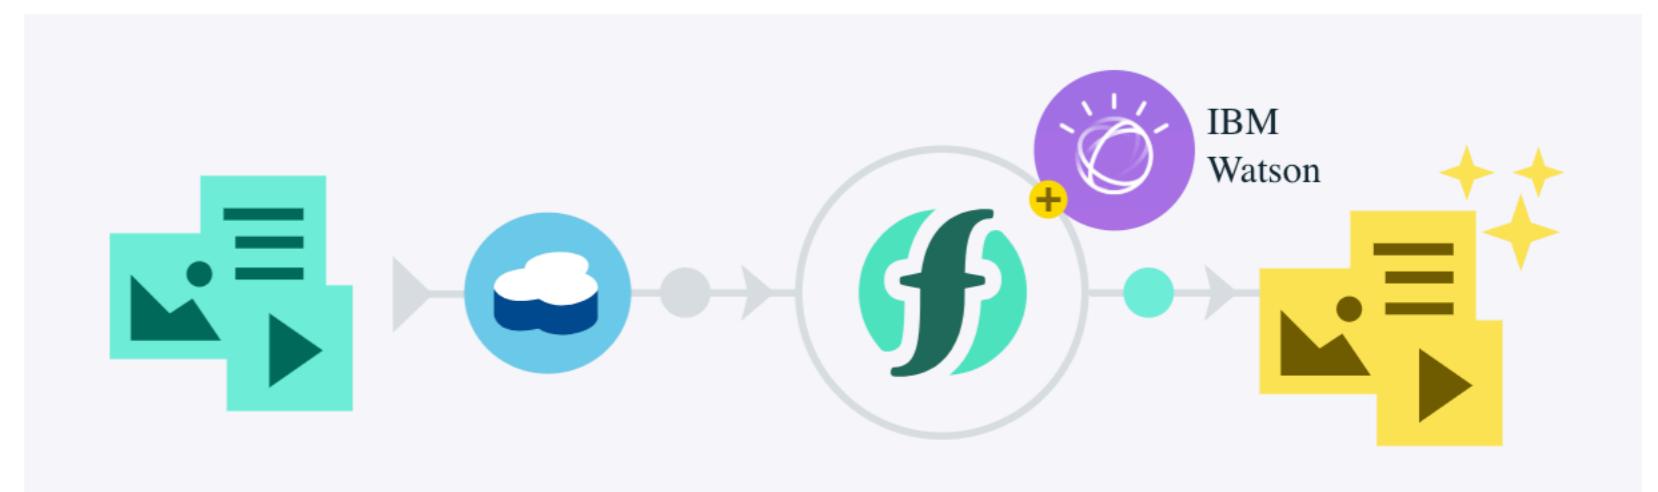
Data Processing

Execute code whenever data is updated in your datastore. Easily automate processes like audio normalization, image rotation, sharpening, noise reduction, thumbnail generation, or video transcoding.



Cognitive Data Processing

Analyze data as soon as it becomes available. Let your function make use of powerful cognitive services like translation or image recognition to detect objects or people appearing in images or videos.



Serverless Runtimes or FaaS current Market

- AWS Lambda available since 2016
- Google Cloud Functions (as part of the Google App engine) since 2016
- Microsoft Azure Functions since 2016
- IBM Functions (on top of Open Source OpenWhisk) since 2017
- Oracle Cloud Platform
- Project Riff developed by Pivotal on top of Kubernetes

The screenshot shows the Apache OpenWhisk website. At the top, there's a blue header bar with the Apache OpenWhisk logo, navigation links for Documentation, Community, and Downloads, and social media icons for GitHub, Slack, Twitter, LinkedIn, YouTube, and a blog.

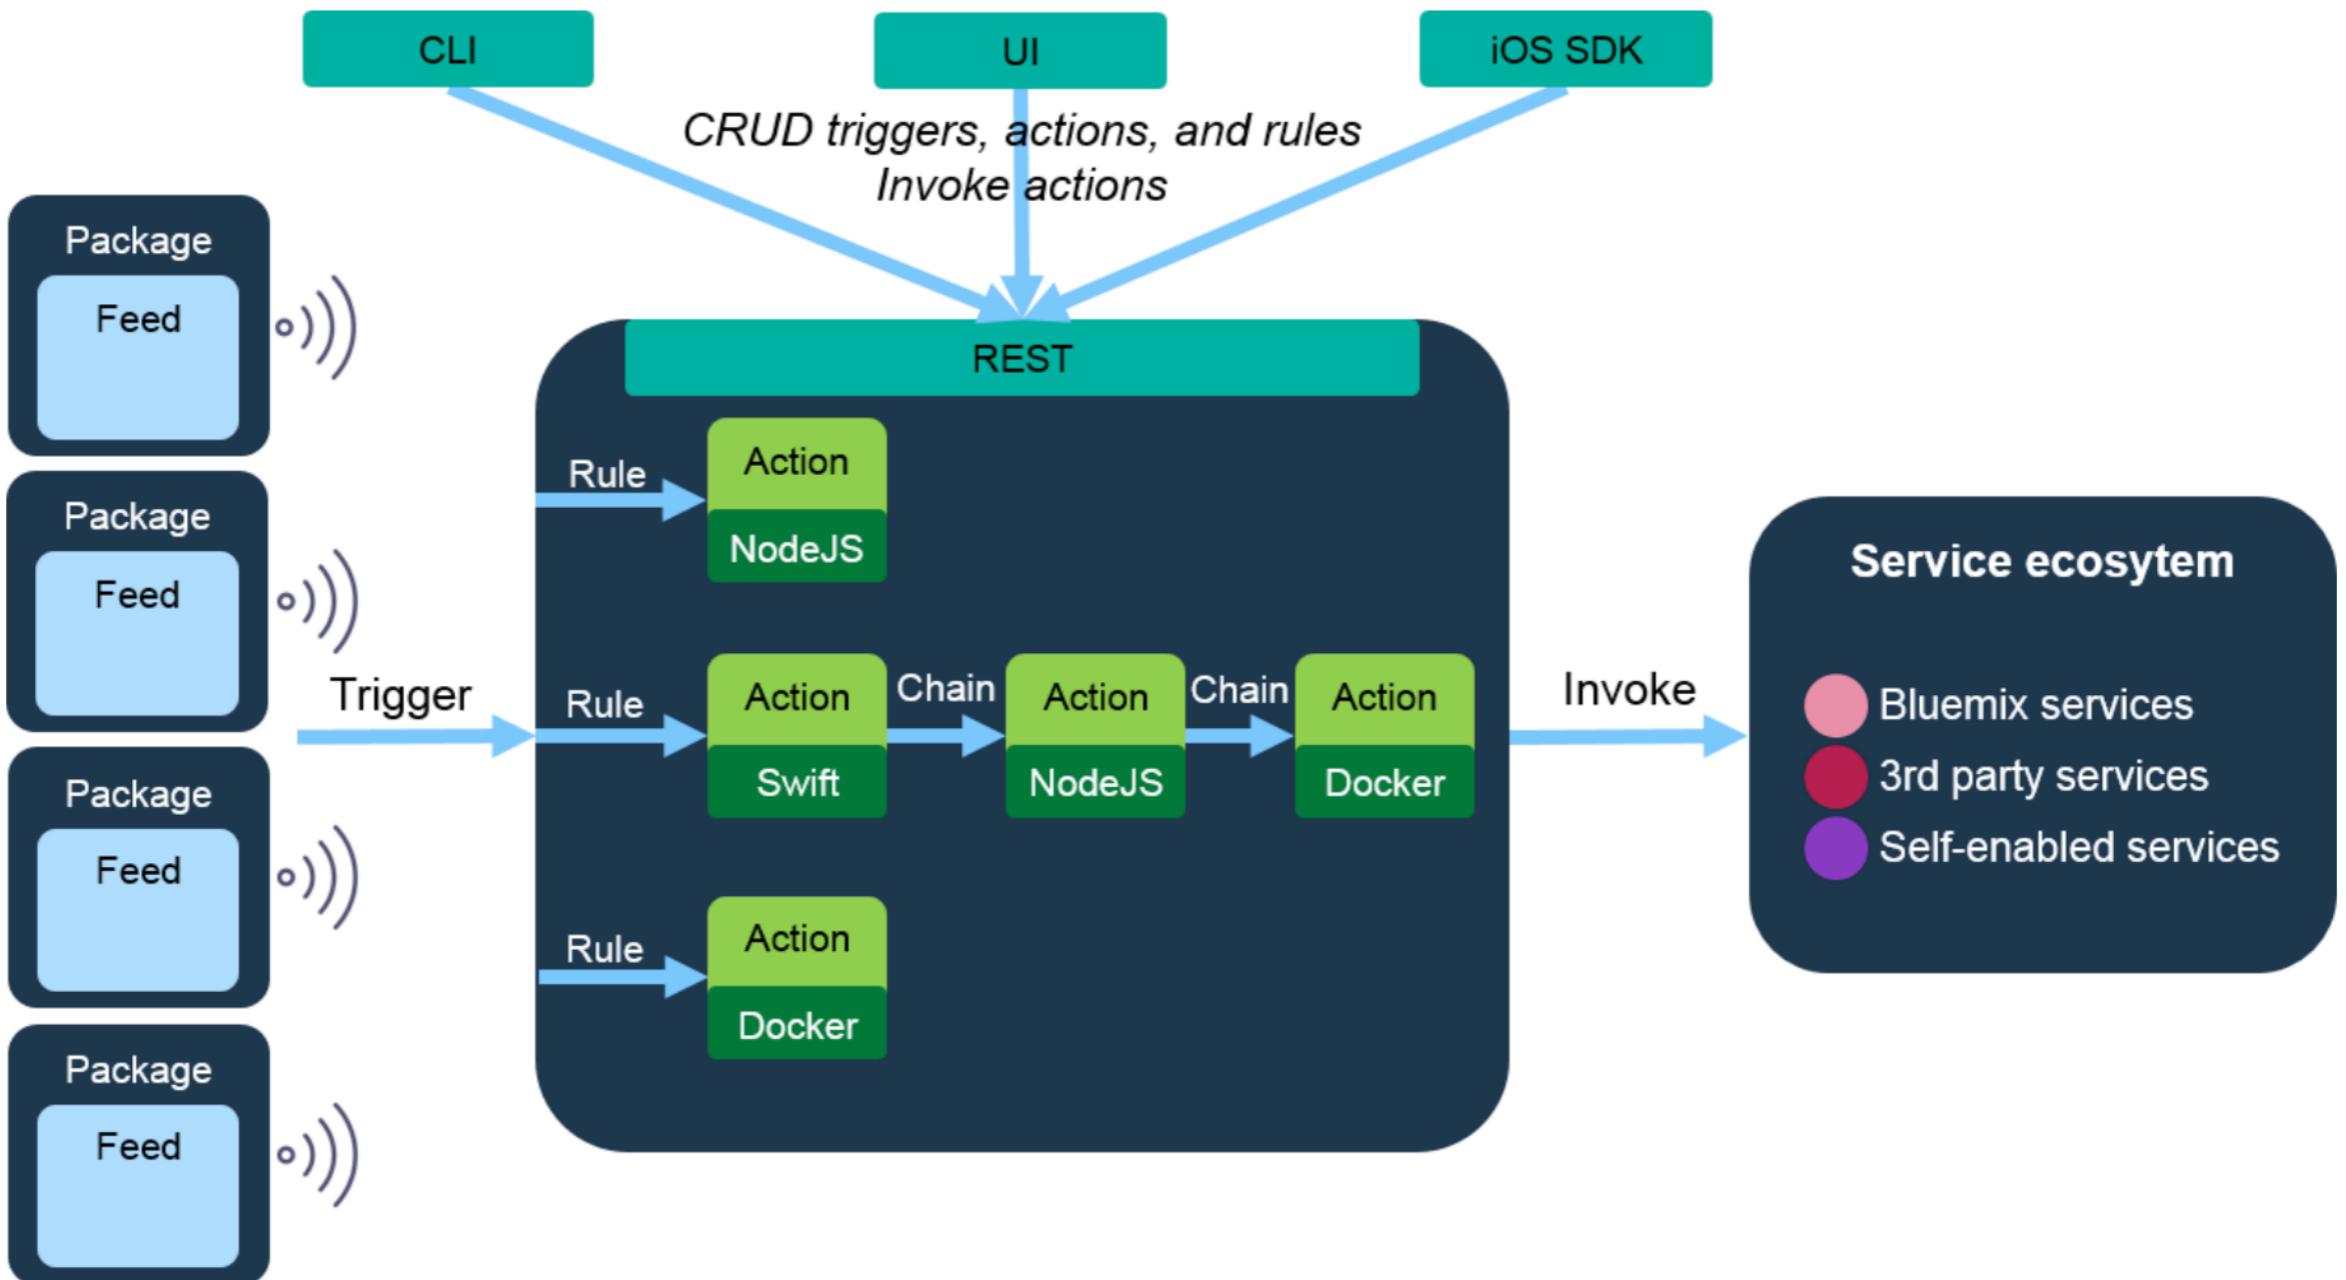
On the left, a sidebar menu includes links for Community, Media, Mailing Lists, Project Wiki, Events, and Supporters. The 'Supporters' link is currently selected.

The main content area features a section titled 'Supporters' with a sub-instruction: 'The following companies and organizations have acknowledged support of the Apache OpenWhisk project as contributors or users of the technology.' Below this, there are two rows of logos from various companies:

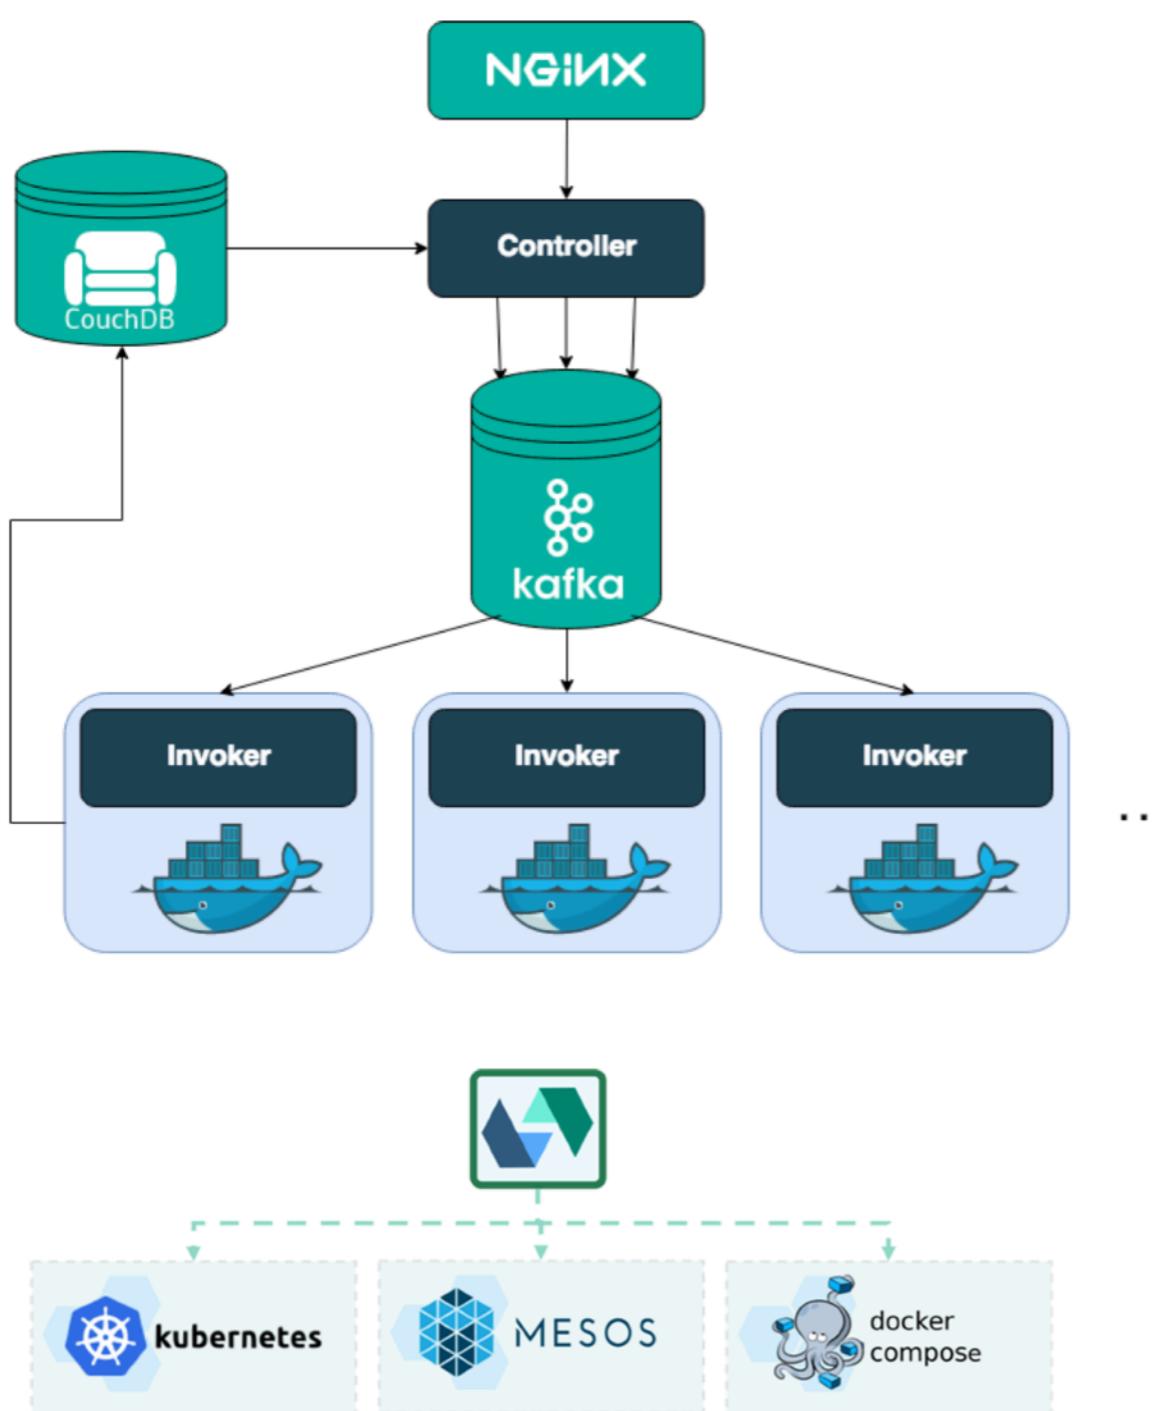
- Row 1: Adobe, ADVISOR CONNECT, ALTOROS™
- Row 2: articoolo, BIGVU POCKET TV STUDIO, CLOUDSTATION, GreenQ Making Garbage Trucks Smarter
- Row 3: IBM, KONG, MAGENTIQ SEE BEYOND THE VISIBLE, NAVER
- Row 4: NeuroApplied, nimbella Your Cloud. Beautiful., redhat., SiteSpirit
- Row 5: WSO2

Apache OpenWhisk

<https://openwhisk.apache.org/community.html>



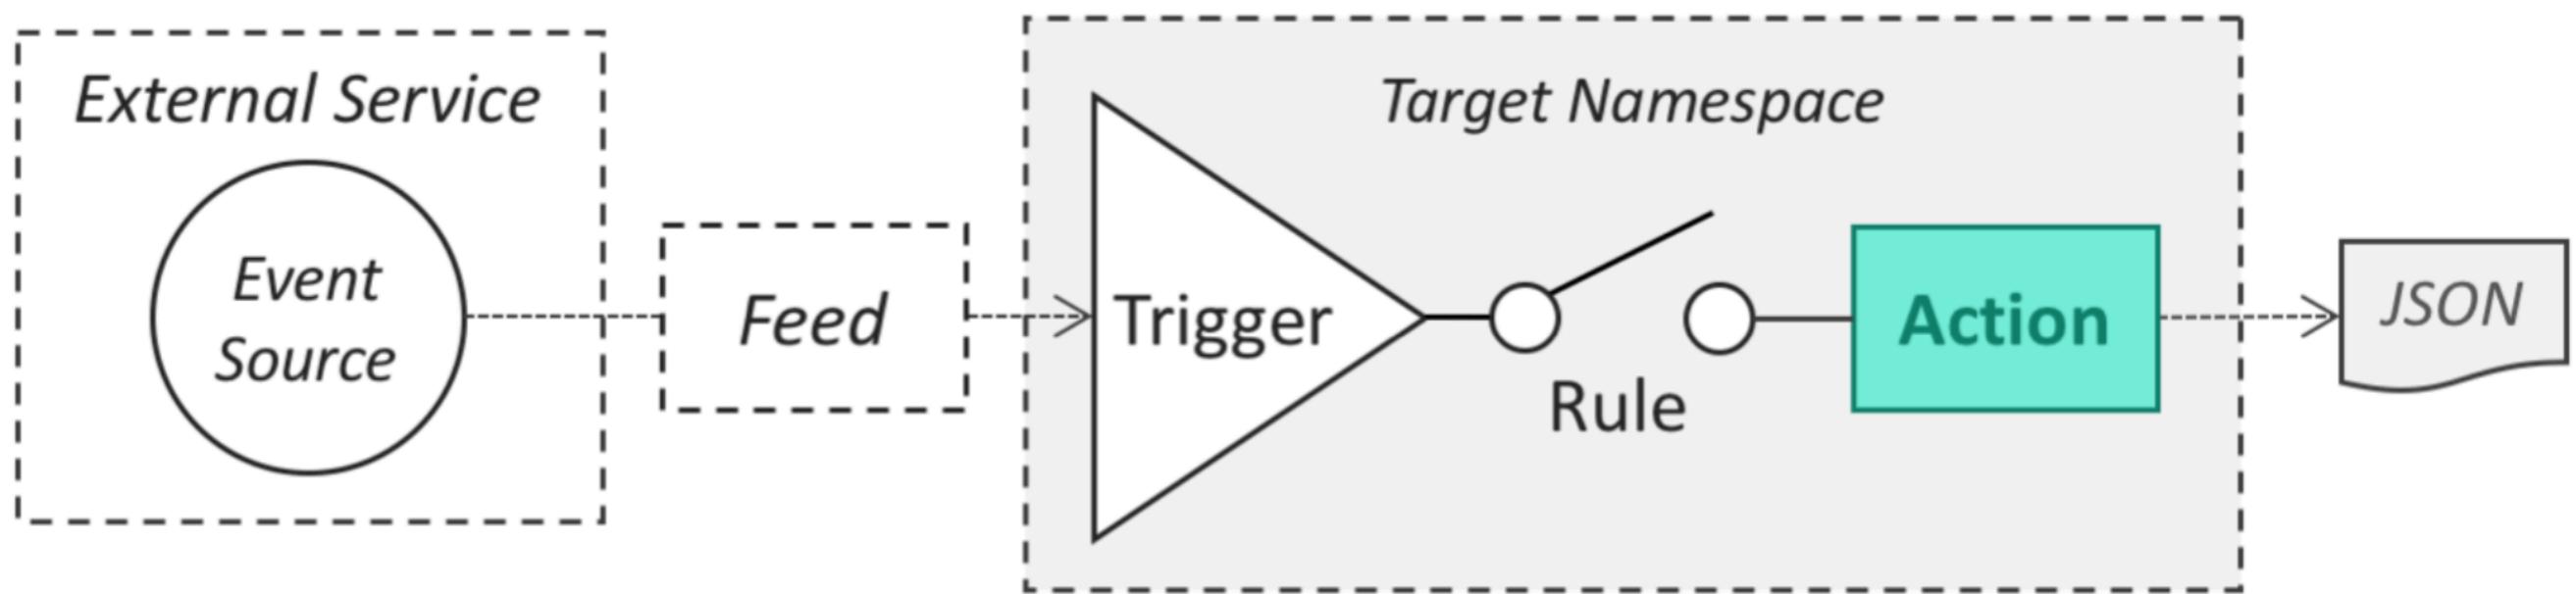
OpenWhisk High Level Flow



- NGINX (HTTP reverse Proxy Server)
 - It is mainly used for SSL termination and forwarding appropriate HTTP calls to the next component.
- Controller
 - Rest Processing
 - Authentication and Authorization
 - writes request to couch DB
 - Using Kafka message streaming
 - Assigns work among the invokers (invokers are in a service registry)
 - Returns the http caller (with an id)
 - (http invoker can query the results later with an additional HTTP Get id)
- Couch DB (Data Services)
 - Credentials
 - Packages (feeds/actions/rules)
 - Results of actions
- Consul
 - Key value store
- Invoker (Rest Executor)
 - Subscribes to requests from the controller.
 - Invoke the action
 - Each action in known docker
 - Put action results into Couch DB
- Kafka (Pub/Sub Streaming)
- A lot more....

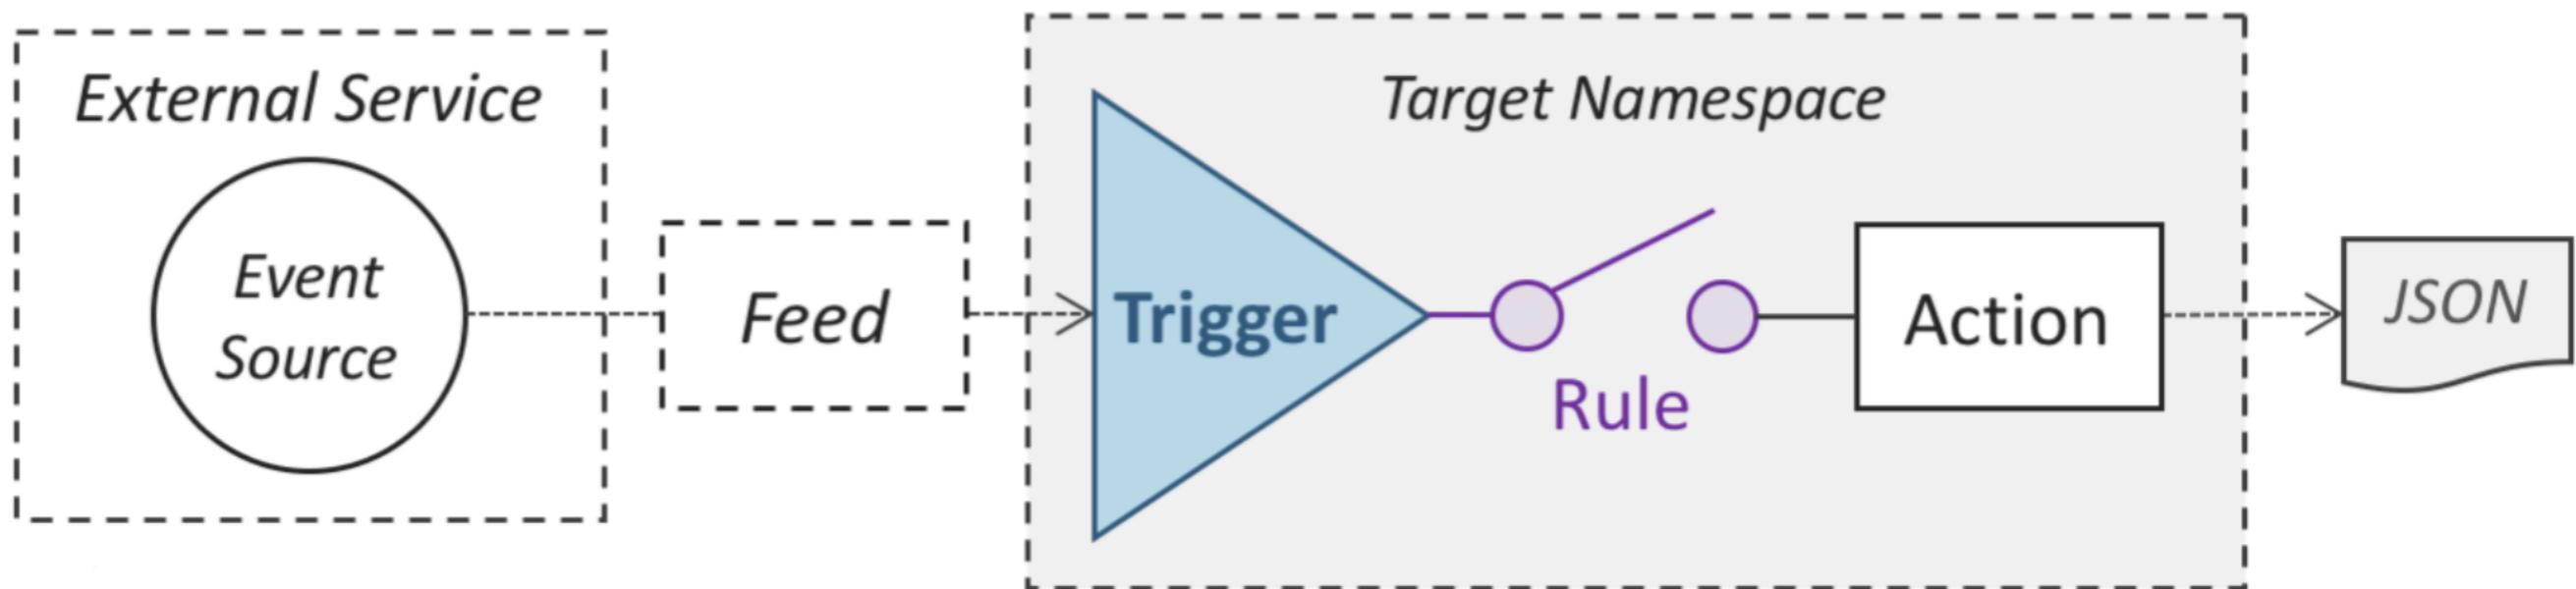
<https://github.com/apache/openwhisk/blob/master/docs/about.md#how-openWhisk-works>

Terminology in OpenWhisk



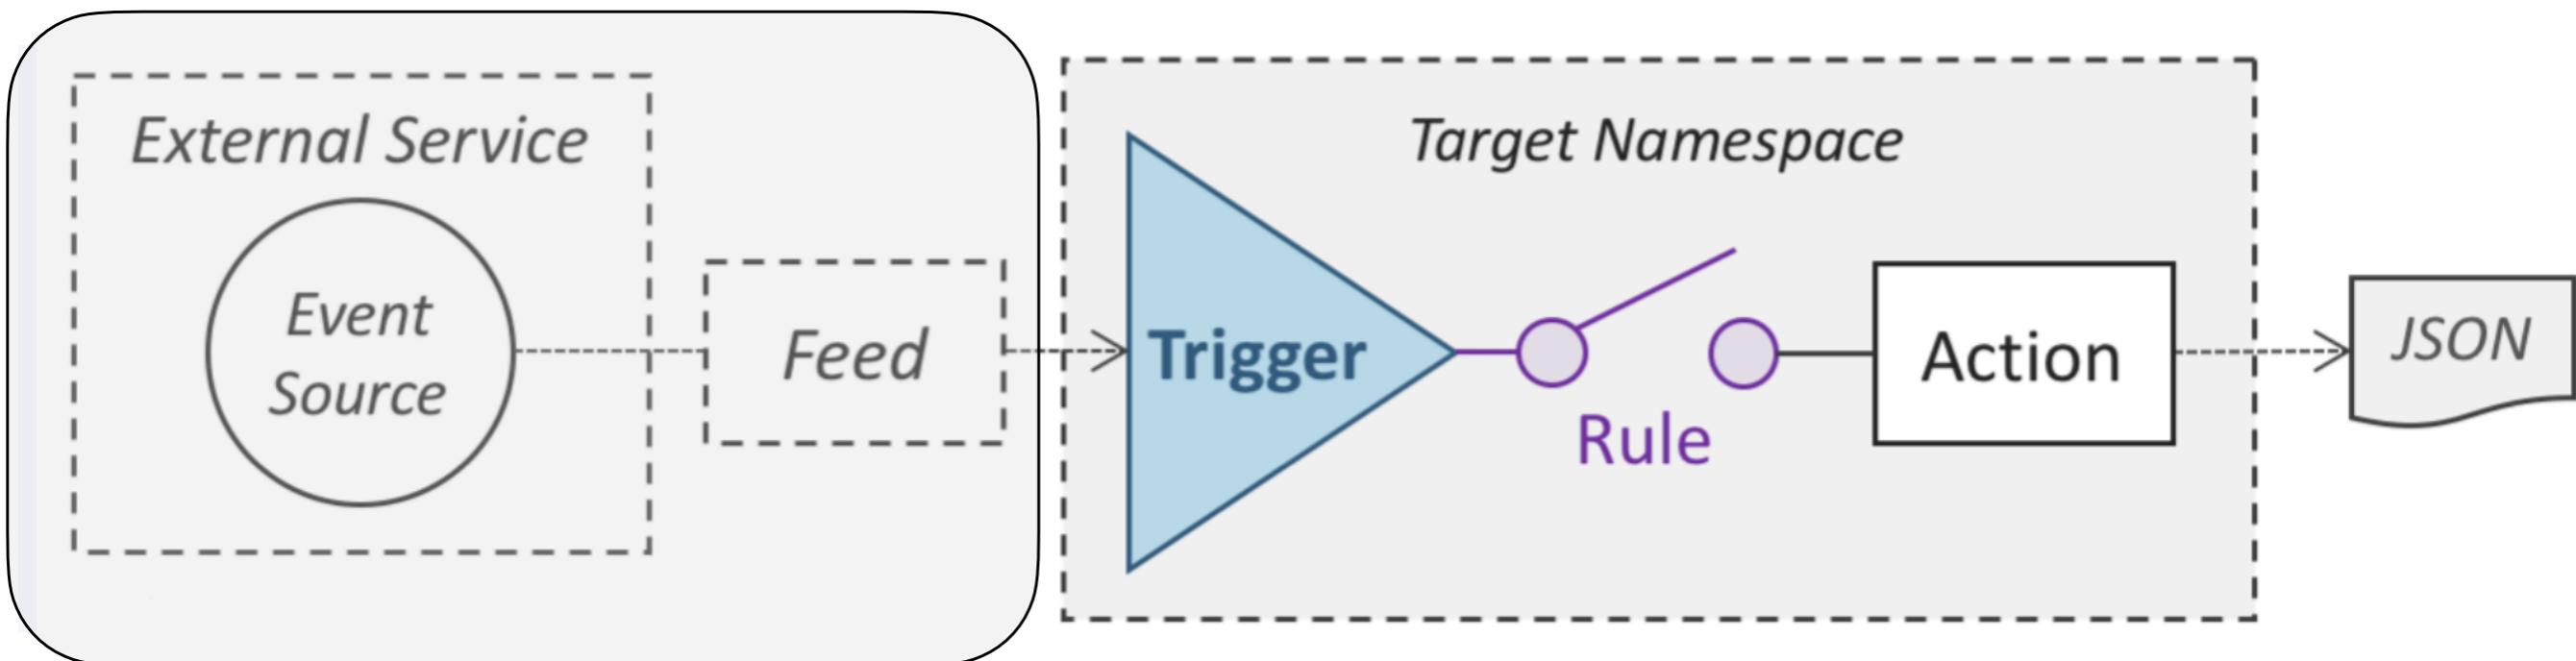
- Action
 - is a piece of code that performs one specific task. An action can be written in the languages described or a custom binary code embedded in a Docker container. **You provide your action to Cloud Functions either as source code or a Docker image.** (Limit is 40 MB)
 - An action performs work when it is directly invoked by using the Cloud Functions API, CLI or automatically respond to events (from an Cloud services) using a trigger.
- Sequence
 - A sequence is a chain of actions, invoked in order, where the output of one action is passed as input to the next action. This allows you to combine existing actions together for quick and easy re-use. A sequence can then be invoked just like an action, through a REST API or automatically in response to events.

Terminology in OpenWhisk (IBM Cloud Functions)



- Trigger
 - A trigger is a **declaration that you want to react to a certain type of event**
 - It describes the specific event (and the given parameters (key/value)) which should happen in order to fire the action
 - E.g. In **couch-db ABC** a document was added to **database XYX**
- Rule
 - Every time the trigger fires, the rule uses the trigger event as input and invokes the associated action.
 - With the appropriate set of rules, it's possible for a single trigger to invoke multiple actions, or for an action to be invoked from multiple triggers.

Terminology in OpenWhisk (IBM Cloud Functions)



- Feed
 - is a convenient way to configure an external event source to fire trigger events (e.g. web hooks). Firing an event is mostly a REST call
- Event Sources
 - These are services that generate events that often indicate changes in data or carry data themselves. Some examples of common Event Sources include:
 - Messages arriving on Message Queues
 - Changes in Databases
 - Changes in Document Stores
 - Website or Web Application interactions
 - Service APIs being invoked
 - IoT Frameworks forwarding device sensor data.

Workflows using Lambda

<https://aws.amazon.com/tutorials/orchestrate-microservices-with-message-queues-on-step-functions/>

The screenshot shows a tutorial page for AWS Step Functions. At the top, there's a navigation bar with links for Products, Solutions, Pricing, Documentation, Learn, Partner Network, AWS Marketplace, Customer Enablement, Events, Explore More, and a search icon. Below the navigation is a secondary row with links for AWS Tutorials Directory, Getting Started Resource Center, Developer Center, IT Pro Center, Architecture Center, Tools & SDKs, and More Resources.

Orchestrate Queue-based Microservices with AWS Step Functions and Amazon SQS

TUTORIAL

Overview

In this tutorial, you will learn how to use AWS Step Functions and Amazon Simple Queue Service (Amazon SQS) to design and run a serverless workflow that orchestrates a message queue-based microservice. Step Functions is a serverless orchestration service that lets you easily coordinate multiple AWS services into flexible workflows that are easy to debug and change. Amazon SQS is the AWS service that allows application components to communicate in the cloud.

This tutorial will simulate inventory verification requests from incoming orders in an e-commerce application as part of an order-processing workflow. Step Functions will send inventory verification requests to a queue on Amazon SQS. An AWS Lambda function will act as your inventory microservice that uses a queue to buffer requests. When it retrieves a request, it will check inventory and then return the result to Step Functions. When a task in Step Functions is configured this way, it is called a *callback pattern*. Callback patterns allow you to integrate asynchronous tasks in your workflow, such as the inventory verification microservice of this tutorial.

✓ AWS experience	Intermediate
⌚ Minimum time to complete	10 minutes
\$ Cost to complete	Free Tier eligible
ℹ Requires	<ul style="list-style-type: none">AWS account with administrator-level access*Recommended browser: The latest version of Chrome or Firefox

*Accounts created within the past 24

PaaS on top of Kubernetes Recap

- Eine immer gestellte Frage ist, welche technische Details notwendig sind um einem Software Entwickler oder einer Software Entwicklerin alle notwendigen PaaS Optionen zu geben, die er oder sie braucht um die eigentlich Entwicklungsaufgabe effizient zu erledigen. Der andere Teil wird von einer Admin Role übernommen. Nenne zwei Beispiele für Konfigurationsoptionen die typischerweise von der Admin Rolle und zwei die typischerweise von der Entwicklung beschrieben werden sollte.
- Positioniere Serverless zu Function as Service)
- Knative unterscheidet zwischen Applications, Jobs und Funktionen. Welche unterschiedliche Scaling Profile haben diese Workloads. (Einfache Beschreibung)
- Buildpacks (welches nun auch in einer PaaS Integration mit Kubernetes zum Einsatz kommen) ist ein Satz von Scripts die in 4 Phasen (detect, build, prepare and run) den zur Verfügung gestellten Code in der PaaS Cloud zum Laufen bringt. Welche Phasen gibt es und was tue sie (grob am Beispiel von Node.js.)
- Knative definiert den Begriff Revision. Was steckt hinter einer Revision
- Unter dem Begriff Serverless versteht man, daß ready to run code (=images) nur im Bedarfsfalle aktiviert wird.
- Der große Vorteil (den die Cloud Provider postulieren) ist, daß nur bei Ausführung des Codes (aktiver Container) auch abgerechnet wird.
- Der Nachteil ist die Performance wenn in gewissen Situationen erst der Container gestartet werden muss und die App sich ggfls. noch initialisiert .. das ist spürbar. Der Gegensatz ist die Anwendung hat einen Server ist aktiv und wartet auf eine Anfrage.
- Neben der Knative und OpenFass Implementierung gibt es noch andere Implementierung von FaaS. Nenne 2 zwei. (Openwisk, AWS mit Lambda und Azure Functions)
- Typische Events für Funktionen sind DB Updates/Changes, IOT Data Arrivals (Kafka), Cron Triggers. HTTP Web Service Endpoints.

Referenzen

- Cloud Foundry
 - Steffen Uhlig, IBM. <https://ws.uhlig.it>
 - <https://docs.cloudfoundry.org/concepts/architecture/>
 - <https://cloudacademy.com/blog/cloud-foundry-components/>
 - Vorlesungsscript, Prof. Dr. Ulrike Steffens, Hochschule für Angewandte Wissenschaften Hamburg
 - <https://docs.cloudfoundry.org/buildpacks/understand-buildpacks.html>
 -
- OpenWisk
 - <https://developer.ibm.com/opentech/2016/09/06/what-makes-serverless-attractive/>
 - <https://github.com/apache/incubator-openwhisk/blob/master/docs/about.md>
 -